

Code Sample*

Anonymous

This document contains excerpts from the replication code of a paper titled “Inequality in a Potential Equalizer of Opportunity: Variation in the Accessibility and Net Benefits of English-Medium Instruction in India.” Excerpts are followed by their corresponding outputs when appropriate. The project may have been updated since this file was generated; the most up-to-date versions of this document, the full paper, and its replication package are available online, and I would be happy to provide links for those interested. This PDF can be generated by running `make samples` or `scripts/run_full_build.sh` from the repo root.

Historical Geography and Measurement

Reconstructing District Boundary Transfers Between 1991 and 2001

A historical Census table reports how each 1991 district’s population was apportioned among 2001 districts after boundary changes. A blank 1991 population cell identifies a candidate continuation row in the headerless Census table: transfer fields distinguish an allocation record from a row that only continues a printed district name. Structural repair runs before ordinary fill-down because those blank cells carry layout information; continuation rows that contain transfer shares remain observations after the full 1991 district name and population are restored. The reader requires each 1991 district’s population shares across 2001 successor districts to sum to 100% within 0.05 percentage points.

File: `R/io/read_inputs.R`

```
# Repair physical row wraps in the historical district transfer table  
#  
# A blank 1991 population field marks a candidate continuation row. Transfer  
# fields distinguish a genuine allocation record from a row that only continues
```

*GPT models 5.5 and 5.6 were used for troubleshooting, reviewing results, and generating certain portions of this code from my detailed outlines and pseudocode, after which I reviewed and edited their responses extensively. All ideas, analytical work, and errors are my own.

```

# a printed name. For split source names, lowercase continuations remove a
# hyphen inserted by a line break; uppercase continuations preserve a hyphen
# that belongs to the district name.
repair_district_carveout_wrapped_rows <- function(x) {
  out <- safe_df(x)
  if (!nrow(out)) return(out)
  blank <- function(value) is.na(value) || !nzchar(trimws(as.character(value)))
  drop <- rep(FALSE, nrow(out))
  for (i in seq_len(nrow(out))) {
    if (i == 1L || !blank(out$pop_1991[[i]])) next
    previous <- i - 1L
    source_continuation <- !blank(out$district_1991[[i]])
    target_continuation <- !blank(out$district_2001[[i]])
    has_transfer <- !blank(out$pct_01in91[[i]]) || !blank(out$pct_91in01[[i]])

    if (source_continuation) {
      previous_name <- plain_chr(out$district_1991[[previous]])
      if (!length(previous_name) || !grepl("-", trimws(previous_name[[1L]]))) {
        stop("Unexpected wrapped row in the district boundary table.", call. = FALSE)
      }
      continuation <- plain_chr(out$district_1991[[i]])[[1L]]
      combined <- join_district_carveout_wrapped_label(previous_name[[1L]], continuation)
      out$district_1991[[previous]] <- combined
      if (has_transfer) {
        out$district_1991[[i]] <- combined
        out$pop_1991[[i]] <- out$pop_1991[[previous]]
      }
    }

    if (!has_transfer && (source_continuation || target_continuation)) {
      if (target_continuation) {
        if (blank(out$district_2001[[previous]])) {
          stop("Unexpected continuation of a 2001 district name in the district boundary table.", call. =
            FALSE)
        }
        previous_target <- trimws(as.character(out$district_2001[[previous]]))
        continuation <- trimws(as.character(out$district_2001[[i]]))
        out$district_2001[[previous]] <- join_district_carveout_wrapped_label(
          previous_target, continuation
        )
      }
    }
  }
}

```

```

      drop[[i]] <- TRUE
    }
  }
  out[!drop, , drop = FALSE]
}

district_carveout_rounding_tolerance_pp <- function() 0.05

# Validate population shares in the 1991 to 2001 district transfer table
#
validate_district_carveout_shares <- function(
  x,
  rounding_tolerance_pp = district_carveout_rounding_tolerance_pp()) {
  out <- safe_df(x)
  if (!nrow(out)) return(invisible(TRUE))
  share_cols <- c("pct_01in91", "pct_91in01")
  for (field in share_cols) {
    values <- num(out[[field]])
    bad <- is.finite(values) & (values < 0 | values > 100)
    if (any(bad)) {
      stop("District transfer shares must lie in [0, 100].", call. = FALSE)
    }
  }
  source_share <- num(out$pct_01in91)
  if (any(!is.finite(source_share))) {
    stop("Shares of each 1991 district assigned to 2001 districts must be complete.", call. = FALSE)
  }
  # Group transfers by the 1991 district name and population printed in the
# table.
  source_key <- paste(out$district_1991, out$pop_1991, sep = "__")
  source_sum <- vapply(split(source_share, source_key), sum, numeric(1))
  if (any(abs(source_sum - 100) > rounding_tolerance_pp)) {
    stop("Shares of each 1991 district assigned to 2001 districts must sum to 100 within rounding
      ↪ tolerance.", call. = FALSE)
  }
  invisible(TRUE)
}

read_district_carveouts <- function(path) {
  out <- utils::read.csv(
    path,

```

```

    header = FALSE,
    col.names = c("district_1991", "pop_1991", "district_2001", "pct_01in91", "pct_91in01"),
    stringsAsFactors = FALSE,
    na.strings = c("", "NA"),
    check.names = FALSE
  )
  # Repair layout blanks before filling downward because the blanks themselves
  # identify candidate continuation rows.
  out <- repair_district_carveout_wrapped_rows(out)
  out$district_1991 <- fill_down_missing(out$district_1991)
  out$pop_1991 <- num(gsub(",", "", fill_down_missing(out$pop_1991), fixed = TRUE))
  out$pct_01in91 <- num(out$pct_01in91)
  out$pct_91in01 <- num(out$pct_91in01)
  validate_district_carveout_shares(out)
  out
}

```

Generating District Match Candidates for Review

District names vary across survey waves and reference years, so standardized names rank candidates for subsequent district lineage review. Within each state, the code averages Jaro–Winkler, normalized Damerau–Levenshtein, trigram cosine, and token Jaccard similarities. The preferred reference year depends on survey wave: 2001 is preferred for the 2007–08 NSS and 2011 for later waves. The review flag also requires the survey district and reference district to rank each other first, compatible directional terms, Jaro–Winkler ≥ 0.90 , normalized Damerau–Levenshtein ≥ 0.70 , trigram similarity ≥ 0.55 , and a score margin ≥ 0.05 over the runner-up. Token Jaccard contributes to the composite score without a separate cutoff. The resulting candidates then enter the reviewed district lineage process.

File: R/districts/lineage_matching.R

```

score_match_candidates <- function(source_roster, reference_units, excluded_source_ids = character()) {
  need_pkg("stringdist", "district-lineage candidate scores")
  source_roster <- safe_df(source_roster)
  reference_units <- safe_df(reference_units)
  if (!"raw_district" %in% names(source_roster)) source_roster$raw_district <- source_roster$district_std
  source_open <- source_roster[!source_roster$source_row_id %in% excluded_source_ids, , drop = FALSE]
  if (!nrow(source_open)) return(empty_source_candidates())

  pairs <- merge(

```

```

    source_open, reference_units,
    by = "state_std", all = FALSE, sort = FALSE,
    suffixes = c("_source", "_candidate")
  )
  if (!nrow(pairs)) return(empty_source_candidates())
  source_name <- pairs$district_std_source
  candidate_name <- pairs$district_std_candidate
  # p = 0.1 applies Winkler's prefix adjustment rather than plain Jaro distance.
  pairs$jw <- 1 - stringdist::stringdist(
    source_name, candidate_name, method = "jw", p = 0.1
  )
  dl_distance <- stringdist::stringdist(source_name, candidate_name, method = "dl")
  denom <- pmax(nchar(source_name), nchar(candidate_name), 1L)
  pairs$dl <- pmax(0, 1 - dl_distance / denom)
  pairs$trigram <- stringdist::stringsim(source_name, candidate_name, method = "cosine", q = 3)
  pairs$token <- token_jaccard_similarity(source_name, candidate_name)
  pairs$score <- rowMeans(pairs[c("jw", "dl", "trigram", "token")], na.rm = TRUE)

  ranked <- safe_bind_rows(lapply(split(seq_len(nrow(pairs))), pairs$source_row_id, function(i) {
    x <- pairs[i, , drop = FALSE]
    # Prefer the reference year closest to the survey wave.
    preference <- vintage_preference(x$wave[[1]])
    vintage_rank <- match(x$reference_vintage, preference, nomatch = length(preference) + 1L)
    x <- x[order(-x$score, -x$jw, vintage_rank, x$unit_id), , drop = FALSE]
    x$rank <- match(x$district_std_candidate, unique(x$district_std_candidate))
    distinct_scores <- x$score[!duplicated(x$district_std_candidate)]
    second <- if (length(distinct_scores) >= 2L) distinct_scores[[2]] else -Inf
    x$margin <- ifelse(x$rank == 1L, x$score - second, NA_real_)
    x$preferred_vintage <- !duplicated(x$district_std_candidate)
    x
  })

  # Require the survey district and reference district to rank each other first.
  candidate_groups <- split(seq_len(nrow(ranked)), ranked$unit_id)
  best_by_candidate <- stats::setNames(
    vapply(candidate_groups, function(i) {
      score <- ranked$score[i]
      ranked$source_row_id[i[which.max(score)]]
    }, character(1)),
    names(candidate_groups)
  )

```

```

ranked$reciprocal_nearest <- ranked$rank == 1L &
  unname(best_by_candidate[ranked$unit_id]) == ranked$source_row_id
threshold <- district_match_candidate_thresholds()
# This flag prioritizes the pair for the subsequent lineage review.
ranked$high_precision_candidate <- ranked$rank == 1L & ranked$preferred_vintage &
  ranked$reciprocal_nearest &
  directional_tokens_compatible(ranked$district_std_source, ranked$district_std_candidate) &
  ranked$jw >= threshold[["jw"]] & ranked$dl >= threshold[["dl"]] &
  ranked$trigram >= threshold[["trigram"]] & ranked$margin >= threshold[["margin"]]
ranked <- ranked[ranked$rank <= 5L, , drop = FALSE]

data.frame(
  source_row_id = ranked$source_row_id,
  wave = ranked$wave,
  source_code = ranked$source_code,
  state_std = ranked$state_std,
  source_name_raw = ranked$raw_district,
  source_name = ranked$district_std_source,
  candidate_unit = ranked$unit_id,
  candidate_name = ranked$district_std_candidate,
  candidate_source_id = ranked$source_id,
  reference_vintage = ranked$reference_vintage,
  candidate_method = "fuzzy_name_candidate",
  rank = ranked$rank,
  jw = ranked$jw,
  dl = ranked$dl,
  trigram = ranked$trigram,
  token = ranked$token,
  score = ranked$score,
  margin = ranked$margin,
  reciprocal_nearest = ranked$reciprocal_nearest,
  high_precision_candidate = ranked$high_precision_candidate,
  stringsAsFactors = FALSE
)
}

```

Reconstructing District Data from the 1991 Language Atlas

The 1991 Census Language Atlas repeats each district across eight language pages. Districts are identified by a Census state code plus a serial within the state, but some repeated rows omit the serial or contain errors introduced during text extraction. When a row itself has no serial, the code recovers one only if the same raw

row label within the same Atlas block has exactly one distinct serial across repeated pages. Before a district row can be promoted, the parser requires the observed state headings to equal the complete reviewed set of 1991 states and the state codes in the Primary Census Abstract. Promotion then requires two independent checks: the reconstructed state and district code pair must identify a unique official 1991 Primary Census Abstract district, and the Atlas population must be within 1% of that district's Census population.

File: scripts/build_language_atlas_1991.py

```
def cross_page_serial_evidence(cells: list[dict[str, object]]) -> dict[tuple[int, str], tuple[int, ...]]:
    """Collect district serials printed with repeated Language Atlas row labels.

    The Atlas prints each district on eight language pages. The district serial is
    the within-state number paired with the Census state code. If one copy omits
    that number, repeated copies can supply it only when their printed serials agree.
    """
    observed: dict[tuple[int, str], set[int]] = {}
    unique_rows = {
        (int(cell["block"]), int(cell["page"]), int(cell["row_sequence"]),
         str(cell["row_label_raw"]).strip(), str(cell["serial_candidate"]).strip())
        for cell in cells
    }
    for block, _page, _sequence, label, serial_raw in unique_rows:
        if not label or label == "INDIA":
            continue
        serial = int(serial_raw) if serial_raw else parse_leading_serial(label)
        if serial is not None:
            observed.setdefault((block, label), set()).add(serial)
    return {key: tuple(sorted(values)) for key, values in observed.items()}

def reconcile_district_serial(
    cell: dict[str, object],
    evidence: dict[tuple[int, str], tuple[int, ...]],
) -> tuple[int | None, str, str]:
    """Choose a district serial from the row itself or agreeing repeated copies."""
    label = str(cell["row_label_raw"]).strip()
    serial_column = int(cell["serial_candidate"]) if str(cell["serial_candidate"]).strip() else None
    serial_label = parse_leading_serial(label)
    repeated = evidence.get((int(cell["block"]), label), ())
    repeated_text = ";".join(map(str, repeated))
```

```

if serial_column is not None and serial_label is not None and serial_column != serial_label:
    return None, "direct_conflict", repeated_text
direct = serial_column if serial_column is not None else serial_label
direct_source = "serial_column" if serial_column is not None else "label_prefix"

if direct is not None:
    # Prefer a serial printed on the current row. Use repeated copies of the
    # same label only when the current row has no serial because
    # continuation pages contain occasional text extraction errors.
    return direct, direct_source, repeated_text
if len(repeated) == 1:
    return repeated[0], "cross_page_exact_label", repeated_text
if len(repeated) > 1:
    return None, "cross_page_conflict", repeated_text
return None, "missing", repeated_text

def build_district_population_validation(
    cells: list[dict[str, object]],
    state_crosswalk: dict[str, dict[str, str]],
    pca91: dict[tuple[str, str], int],
) -> list[dict[str, object]]:
    rows = merge_population_continuations(cells, state_crosswalk)
    serial_evidence = cross_page_serial_evidence(cells)
    seen_state_codes: set[str] = set()
    current_state: dict[str, str] | None = None
    candidates: list[dict[str, object]] = []
    for cell in rows:
        label = str(cell["row_label_raw"]).strip()
        if label == "INDIA":
            continue
        mapped_state = state_crosswalk.get(label)
        if mapped_state is not None and mapped_state["state_code_1991"] not in seen_state_codes:
            current_state = mapped_state
            seen_state_codes.add(mapped_state["state_code_1991"])
            continue

        serial, serial_source, cross_page_serials = reconcile_district_serial(cell, serial_evidence)

        state_code = current_state["state_code_1991"] if current_state else ""

```

```

state_name = current_state["state_name_1991"] if current_state else ""
district_code = f"{serial:02d}" if serial is not None else ""
pca_population = pca91.get((state_code, district_code)) if state_code and district_code else None
# The extractor stores any parsed numeric cell in one candidate field;
# on the population page that field is the district population.
atlas_population = (
    int(cell["speaker_count_candidate"])
    if str(cell["speaker_count_candidate"]).strip()
    else None
)
if atlas_population is not None and pca_population is not None and pca_population > 0:
    abs_diff = abs(atlas_population - pca_population)
    rel_diff = abs_diff / pca_population
else:
    abs_diff = None
    rel_diff = None
candidates.append({
    "block": cell["block"],
    "page": cell["page"],
    "row_sequence": cell["row_sequence"],
    "row_sequence_end": cell["row_sequence_end"],
    "row_merge_status": cell["row_merge_status"],
    "row_label_raw": label,
    "state_code_1991": state_code,
    "state_name_1991": state_name,
    "district_code_1991": district_code,
    "serial_source": serial_source,
    "cross_page_serials": cross_page_serials,
    "atlas_population_raw": cell["raw_value"],
    "atlas_population_candidate": "" if atlas_population is None else atlas_population,
    "pca91_population": "" if pca_population is None else pca_population,
    "population_abs_diff": "" if abs_diff is None else abs_diff,
    "population_relative_diff": "" if rel_diff is None else round(rel_diff, 8),
    "identity_status": "",
    "population_status": "",
    "promotion_status": "review_required",
})

expected_state_codes = {row["state_code_1991"] for row in state_crosswalk.values()}
if seen_state_codes != expected_state_codes:
    missing = sorted(expected_state_codes - seen_state_codes)

```

```

        extra = sorted(seen_state_codes - expected_state_codes)
        raise ValueError(
            "Language Atlas state headings do not match the expected 1991 states; "
            f"missing={missing}; extra={extra}"
        )
pca_state_codes = {state_code for state_code, _ in pca91}
if pca_state_codes != expected_state_codes:
    raise ValueError("Language Atlas state names and 1991 Primary Census Abstract state codes disagree")

key_counts: dict[tuple[str, str], int] = {}
for row in candidates:
    key = (str(row["state_code_1991"]), str(row["district_code_1991"]))
    if all(key):
        key_counts[key] = key_counts.get(key, 0) + 1

for row in candidates:
    key = (str(row["state_code_1991"]), str(row["district_code_1991"]))
    if not row["state_code_1991"]:
        identity_status = "missing_state_context"
    elif row["serial_source"] in {"direct_conflict", "cross_page_conflict"}:
        identity_status = "serial_conflict"
    elif not row["district_code_1991"]:
        identity_status = "missing_district_serial"
    elif key_counts.get(key, 0) > 1:
        identity_status = "duplicate_district_code"
    elif key not in pca91:
        identity_status = "missing_pca_district"
    else:
        identity_status = "candidate_official_code"
    row["identity_status"] = identity_status

    if row["atlas_population_candidate"] == "":
        population_status = "atlas_population_unparsed"
    elif row["pca91_population"] == "":
        population_status = "pca_population_missing"
    elif float(row["population_relative_diff"]) <= POPULATION_RELATIVE_TOLERANCE:
        population_status = "within_1pct"
    else:
        population_status = "over_1pct"
    row["population_status"] = population_status
    if identity_status == "candidate_official_code" and population_status == "within_1pct":

```

```

    row["promotion_status"] = "population_validated_candidate"
  }
  return candidates
}

```

Bounding Historical Linguistic Distance with Unresolved Population Counts

Some 1991 Language Atlas population remains unresolved after language name reconciliation. The code carries that population into explicit bounds on the speaker-weighted distance measure. It separates accepted speakers whose distance is positive from known groups excluded from that mean (zero distance and native English speakers), then allows unresolved population either to leave the mean unchanged or to take any value on the positive distance support observed in the reviewed preferred mapping. The preferred historical measure requires at least 99% accepted speaker coverage and a worst case bound no wider than 0.5 Shastry degree.

File: R/measures/build_linguistic_distance_iv.R

```

historical_linguistic_population_coherent <- function(speakers, population) {
  speakers <- num(speakers)
  if (!is.numeric(population) || length(population) != 1L ||
      !is.finite(population) || population <= 0) {
    stop("Historical linguistic bounds require a positive district population.", call. = FALSE)
  }
  valid_count <- is.finite(speakers) & speakers >= 0
  valid_total <- sum(speakers[valid_count], na.rm = TRUE)
  isTRUE(valid_total <= population)
}

historical_linguistic_distance_bounds <- function(
  speakers, degree, language, population, nonzero_degree_range = c(1, 5)) {
  speakers <- num(speakers)
  degree <- num(degree)
  language <- normalize_language_label(language)
  population_coherent <- historical_linguistic_population_coherent(speakers, population)
  valid_count <- is.finite(speakers) & speakers >= 0
  english <- valid_count & language == "English"
  known_nonzero <- valid_count & is.finite(degree) & degree > 0 & !english
  known_zero <- valid_count & is.finite(degree) & degree == 0 & !english
  # These resolved speakers are excluded from the positive distance mean.
  known_excluded <- known_zero | english
}

```

```

nonzero_speakers <- sum(speakers[known_nonzero], na.rm = TRUE)
numerator <- sum(speakers[known_nonzero] * degree[known_nonzero], na.rm = TRUE)
known_excluded_speakers <- sum(speakers[known_excluded], na.rm = TRUE)
resolved_speakers <- nonzero_speakers + known_excluded_speakers
unresolved_mass <- max(0, population - resolved_speakers)
point <- if (nonzero_speakers > 0) numerator / nonzero_speakers else NA_real_

nonzero_degree_range <- sort(unique(num(nonzero_degree_range)))
if (length(nonzero_degree_range) < 2L || any(!is.finite(nonzero_degree_range)) ||
    any(nonzero_degree_range <= 0)) {
  stop("Historical linguistic distance bounds require a finite positive degree range.", call. = FALSE)
}
degree_min <- min(nonzero_degree_range)
degree_max <- max(nonzero_degree_range)
if (!population_coherent) {
  point <- lower <- upper <- width <- unresolved_mass <- NA_real_
} else if (!is.finite(point) || !is.finite(degree_min) || !is.finite(degree_max)) {
  lower <- upper <- width <- NA_real_
} else if (unresolved_mass <= 0) {
  lower <- upper <- point
  width <- 0
} else {
  lower_if_nonzero <- (numerator + unresolved_mass * degree_min) /
    (nonzero_speakers + unresolved_mass)
  upper_if_nonzero <- (numerator + unresolved_mass * degree_max) /
    (nonzero_speakers + unresolved_mass)
  # Unresolved population may consist of Hindi speakers, English speakers, or
  # language counts without an accepted distance value; in those cases the
  # positive distance mean is unchanged.
  lower <- min(point, lower_if_nonzero)
  upper <- max(point, upper_if_nonzero)
  width <- upper - lower
}

list(
  point = point,
  lower = lower,
  upper = upper,
  width = width,
  nonzero_speakers = nonzero_speakers,
  known_excluded_speakers = known_excluded_speakers,

```

```

resolved_speakers = resolved_speakers,
unresolved_mass_upper_bound = unresolved_mass,
population_coherent = population_coherent,
degree_min = degree_min,
degree_max = degree_max
)
}

```

Output: `outputs/figures/appendix/appendix_historical_language_persistence.pdf`

Rendering code: `R/output/appendix_data_construction_exhibits.R` and `R/output/save_figures.R`




Figure 2: Persistence of speaker-weighted linguistic distance from Hindi from 1991 to 2001. Each point is a 2001 Census district with a reconstructed measure in both years; the dashed line is the 45-degree line. Sources: Census of India 1991 Language Atlas (Office of the Registrar General and Census Commissioner, India 2004) and Census of India 2001 language tables (Office of the Registrar General and Census Commissioner of India 2001).

Independent Schooling Measures and Survey Estimation

Reconstructing English-Medium Enrollment from Archived DISE Reports

District Information System for Education (DISE) report cards from 2008–09 through 2014–15 provide an independent validation of the NSS English-medium measure. Reviewed pages are converted with Poppler’s `pdftotext -layout`, which preserves the physical text layout as far as possible. For tables with languages in rows, the parser requires a nearby **Total** header and takes the final numeric entry on each English or Hindi row. For tables with languages in columns, it locates the language headings and reads the nearest numeric entry on an explicit **Total**

or **Grand Total** row. The full script reconstructs each registered district and year from its reviewed page and compares the result with the tracked count table. When several reports cover the same district and year, the recorded report priority must select exactly one.

File: scripts/build_dise_report_language_enrollment.py

```
LANGUAGES = ("english", "hindi")
NUMBER = re.compile(r"(?![\w.])?(?:\d{1,3}(?:,\d{3})+|\d+)(?![\w.])")
MEDIUM_HEADING = re.compile(r"medium\s+of\s+instruction", re.IGNORECASE)
# Keep both searches local to the table region covered by the supported layouts.
MEDIUM_BLOCK_LINES = 28
TOTAL_HEADER_LOOKBACK = 6

def medium_block(text: str, radius: int = MEDIUM_BLOCK_LINES) -> list[str]:
    lines = text.splitlines()
    hit = next((i for i, line in enumerate(lines) if MEDIUM_HEADING.search(line)), None)
    if hit is None:
        raise ValueError("page has no 'medium of instruction' heading")
    return lines[hit : min(len(lines), hit + radius)]

def _has_total_header(lines: list[str], row_index: int) -> bool:
    start = max(0, row_index - TOTAL_HEADER_LOOKBACK)
    return any(re.search(r"\btotal\b", line, re.IGNORECASE) for line in lines[start:row_index])

def parse_languages_as_rows(lines: list[str]) -> dict[str, int | None] | None:
    values: dict[str, int | None] = {language: None for language in LANGUAGES}
    matched = False
    for i, line in enumerate(lines):
        stripped = line.strip()
        language = next(
            (name for name in LANGUAGES if re.match(rf"^{name}\b", stripped, re.IGNORECASE)),
            None,
        )
        if language is None:
            continue
        numbers = [parse_int(m.group(0)) for m in NUMBER.finditer(stripped)]
        if not numbers:
            continue
```

```

        if not _has_total_header(lines, i):
            raise ValueError(
                f"{language} row has numeric fields but no nearby Total header"
            )
        values[language] = numbers[-1]
        matched = True
    return values if matched else None

def _column_positions(header: str) -> dict[str, int]:
    # pdftotext -layout retains page spacing closely enough here for character
    # positions to identify the English and Hindi columns.
    positions = {}
    lower = header.lower()
    for language in LANGUAGES:
        pos = lower.find(language)
        if pos >= 0:
            positions[language] = pos
    return positions

def _nearest_number_at_column(line: str, position: int) -> int | None:
    matches = list(NUMBER.finditer(line))
    if not matches:
        return None
    # Match the language heading to the nearest numeric field on the Total row.
    nearest = min(matches, key=lambda m: abs(m.start() - position))
    return parse_int(nearest.group(0))

def parse_languages_as_columns(lines: list[str]) -> dict[str, int | None] | None:
    header_index = None
    positions: dict[str, int] = {}
    for i, line in enumerate(lines):
        candidate = _column_positions(line)
        if candidate:
            header_index = i
            positions = candidate
            break
    if header_index is None:
        return None

```

```

total_line = next(
    (
        line
        for line in lines[header_index + 1 :]
        if re.match(r"\s*(?:grand\s+)?total\b", line, re.IGNORECASE)
    ),
    None,
)

if total_line is None:
    raise ValueError("language columns found, but no Total or Grand Total row was found")

return {
    language: _nearest_number_at_column(total_line, position)
    for language, position in positions.items()
} | {language: None for language in LANGUAGES if language not in positions}

def parse_language_counts(text: str) -> dict[str, int | None]:
    lines = medium_block(text)
    parsed = parse_languages_as_rows(lines)
    if parsed is None:
        parsed = parse_languages_as_columns(lines)
    if parsed is None or parsed.get("english") is None:
        raise ValueError("could not resolve English enrollment from medium-of-instruction table")
    return parsed

def preferred_rows(rows: Iterable[dict[str, str]]) -> list[dict[str, str]]:
    grouped: dict[tuple[str, str, str], list[dict[str, str]]] = defaultdict(list)
    for row in rows:
        key = (
            row["academic_year"].strip(),
            row["state_report"].strip(),
            row["district_report"].strip(),
        )
        grouped[key].append(row)

    selected = []
    for key, group in grouped.items():
        priorities = [int(row["report_priority"]) for row in group]

```

```

        best = min(priorities)
        winners = [row for row in group if int(row["report_priority"]) == best]
        if len(winners) != 1:
            raise ValueError(f"multiple reports share the highest priority for {key}")
        selected.append(winners[0])
    return sorted(
        selected,
        key=lambda row: (
            row["academic_year"],
            row["state_report"],
            row["district_report"],
        ),
    )

def rebuild(
    metadata: Path,
    registry: Path,
    archive_root: Path,
    pdftotext: str,
) -> list[dict[str, str]]:
    report_map = registered_reports(registry)
    rebuilt = []
    for row in preferred_rows(read_csv(metadata)):
        source = row["source_pdf"].strip()
        if source not in report_map:
            raise ValueError(f"report metadata references an unregistered PDF: {source}")
        pdf = archive_root / report_map[source]
        if not pdf.is_file():
            raise FileNotFoundError(f"missing registered DISE report PDF: {pdf}")
        page = int(row["source_page"])
        counts = parse_language_counts(extract_pdf_page(pdf, page, pdftotext))
        out = dict(row)
        out["english_enrollment"] = str(counts["english"])
        out["hindi_enrollment"] = "" if counts["hindi"] is None else str(counts["hindi"])
        rebuilt.append(out)
    return rebuilt

```

Constructing the NSS Survey Design

The 2007–08 National Sample Survey (NSS) education survey samples households within first-stage units such as villages or urban blocks. This function defines each sampling stratum by state, rural/urban sector, stratum, and substratum, supplies the household multiplier as the sampling weight, and creates the `survey::svydesign()` used for the enrollment probit. In the `survey` package, `nest = TRUE` relabels cluster identifiers to enforce nesting within strata. Final estimation requires all six design fields before fitting the survey-weighted probit.

File: `R/selection/estimate_selection_probit.R`

```
# List NSS fields needed to reproduce the survey sample design
#
selection_survey_design_variables <- function() {
  c("FSU_SL_NO", "weight", "STATE", "SECTOR", "STRATUM", "SUB_STRATUM_NO")
}

# Construct the NSS design used for weighted enrollment estimation
#
build_survey_design_selection <- function(selection_df, require_all = FALSE) {
  # Final estimation calls this with require_all = TRUE.
  required <- selection_survey_design_variables()
  missing <- setdiff(required, names(selection_df))
  if (isTRUE(require_all) && length(missing)) {
    stop(
      "NSS survey design is missing required fields: ",
      paste(missing, collapse = ", "),
      call. = FALSE
    )
  }

  psu <- first_col(selection_df, c("FSU_SL_NO", "fsu", "PSU", "psu"))
  weight <- first_col(selection_df, c("weight", "WEIGHT", "Multiplier", "multiplier"))
  strata_cols <- intersect(
    setdiff(required, c("FSU_SL_NO", "weight")),
    names(selection_df)
  )
  if (is.null(psu) || is.null(weight) || !length(strata_cols)) {
    if (isTRUE(require_all)) {
      stop("The required NSS survey design could not be constructed.", call. = FALSE)
    }
  }
}
```

```

    return(NULL)
  }
  # NSS strata are defined by the joint state, rural/urban sector, stratum, and
  # substratum identifiers used in the survey files.
  selection_df$.survey_strata <- interaction(selection_df[strata_cols], drop = TRUE)
  # FSU identifiers can repeat across NSS strata; nest = TRUE relabels them to
  # enforce nesting within the joint sampling stratum.
  survey::svydesign(
    ids = stats::as.formula(paste0("~", psu)),
    strata = ~.survey_strata,
    weights = stats::as.formula(paste0("~", weight)),
    data = selection_df,
    nest = TRUE
  )
}

```

Average Marginal Effects on the Estimation Sample

The code recovers the fitted model frame and NSS sampling weights stored with the enrollment probit, verifies one finite nonnegative weight per fitted observation and positive total weight, and passes both to `marginaleffects::avg_slopes()`. The fitted survey design determines model estimation and its covariance matrix. `newdata` fixes evaluation at the fitted observations, while `wts` determines the weighted average of observation level quantities on the response scale. Continuous regressors are summarized by average response-scale slopes; for binary and categorical regressors, `marginaleffects` reports discrete comparisons.

File: R/selection/compute_average_marginal_effects.R

```

# Recover observations and sampling weights used to fit the enrollment model
#
ame_model_data_and_weights <- function(model) {
  model_data <- as.data.frame(stats::model.frame(model))

  model_weights <- if (inherits(model, "svyglm") && !is.null(model$survey.design)) {
    stats::weights(model$survey.design, type = "sampling")
  } else {
    tryCatch(stats::weights(model, type = "prior"), error = function(e) NULL)
  }

  if (is.null(model_weights)) model_weights <- rep(1, nrow(model_data))
}

```

```

model_weights <- as.numeric(model_weights)
if (length(model_weights) != nrow(model_data)) {
  stop("AME averaging weights must have one value per fitted observation.", call. = FALSE)
}
if (any(!is.finite(model_weights)) || any(model_weights < 0) || sum(model_weights) <= 0) {
  stop("AME averaging weights must be finite, nonnegative, and have positive total weight.", call. = FALSE)
}

list(data = model_data, wts = model_weights)
}

# Compute average marginal effects over the fitted NSS sample
#
compute_ames_margineffects <- function(model) {
  # Evaluate effects on the fitted model frame and use its sampling weights when
  # averaging observation level quantities on the response scale. For binary and
  # categorical regressors, margineffects reports discrete comparisons.
  amed <- ame_model_data_and_weights(model)
  margineffects::avg_slopes(
    model,
    newdata = amed$data,
    wts = amed$wts,
    vcov = TRUE,
    type = "response"
  )
}

```

Output: outputs/tables/appendix/appendix_selection_ame.tex

Rendering code: R/output/appendix_selection_exhibits.R and R/output/save_tables.R

Table 7: Average Marginal Effects for School Enrollment

Term	Enrolled (1 = yes)
Age (years)	-0.028*** (0.000)
Female (ref: Male)	0.044*** (0.003)
Household size	-0.004*** (0.001)
Religion: Muslim (ref: Hindu)	-0.126*** (0.007)
Religion: Christian	-0.043*** (0.013)
Religion: Sikh	0.040

Table 7: Average Marginal Effects for School Enrollment (*continued*)

Term	Enrolled (1 = yes)
	(0.026)
Religion: Jain	-0.001
	(0.012)
Religion: Buddhist	0.047*
	(0.024)
Religion: Zoroastrian	0.101***
	(0.034)
Religion: Other	0.252***
	(0.003)
Social group: Scheduled Tribe (ref: Other)	-0.102***
	(0.008)
Social group: Scheduled Caste	-0.075***
	(0.006)
Social group: Other Backward Class	-0.029***
	(0.005)
Urban (ref: Rural)	-0.003
	(0.005)
Distance 1–2km (ref: <1km)	-0.077***
	(0.026)
Distance 2–3km	-0.001
	(0.010)
Distance 3–5km	0.004
	(0.066)
Distance > 5km	0.036
	(0.035)
Male household education: Literate, no school (ref: Illiterate)	0.097***
	(0.017)
Male household education: Literate, school < primary	0.127***
	(0.006)
Male household education: Primary	0.142***
	(0.005)
Male household education: Upper primary	0.195***
	(0.006)
Male household education: Secondary	0.212***
	(0.007)
Male household education: Higher secondary	0.247***
	(0.008)
Male household education: Postsecondary+	0.245***
	(0.009)

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Entries are average changes in predicted enrollment probability. Continuous covariates are average slopes; categorical covariates are average discrete comparisons against the stated reference category, evaluated over the observed covariate distribution. Standard errors use the NSS sampling design, including household weights, primary sampling units, and strata formed from state, rural/urban sector, stratum, and substratum identifiers. The model also controls for district averages of schooling costs and benefits among enrolled children; the table displays marginal effects for child and household covariates. Source: NSS 64th Round, 2007–08, Participation and Expenditure in Education.

Identification and Econometric Robustness

Assessing Weak-Instrument Strength with the Montiel Olea–Pflueger Effective F

The paper uses linguistic distance from Hindi as an instrument for district exposure to English-medium schooling. For specifications with one endogenous regressor, this function evaluates first-stage strength with the Montiel Olea–Pflueger effective F, using `momentfit::MOPtest()` with a 10% relative-bias benchmark for TSLS, 5% test size, and the conservative simplified critical value. It reconstructs the fitted regression and instrument formulas, checks the variables they require, and verifies that state clusters align with the fitted observations before constructing the MOP moment model. The table reports the effective F alongside two-stage least squares and Anderson–Rubin inference robust to weak identification.

File: `R/iv/weak_identification.R`

```
mop_effective_f <- function(
  model, analysis_data, tau = 0.10, size = 0.05) {
  unavailable <- function(reason) {
    list(
      statistic = NA_real_, critical_value = NA_real_, p.value = NA_real_,
      effective_df = NA_real_, tau = tau, size = size,
      status = "not_estimated", reason = reason
    )
  }

  if (!inherits(model, "ivreg")) {
    return(unavailable("Montiel Olea-Pflueger effective F requires an ivreg model."))
  }

  if (!requireNamespace("momentfit", quietly = TRUE)) {
    return(unavailable("Package 'momentfit' is not installed."))
  }

  if (!is.finite(tau) || tau <= 0 || !is.finite(size) || size <= 0 || size >= 1) {
    return(unavailable("MOP tau must be positive, and size must lie strictly between 0 and 1."))
  }

  endogenous <- suppressWarnings(as.integer(model$endogenous))
  if (length(endogenous) != 1L || !is.finite(endogenous)) {
    return(unavailable("Montiel Olea-Pflueger effective F requires exactly one endogenous regressor."))
  }

  regression_formula <- tryCatch(
```

```

stats::formula(model, component = "regressors"),
error = function(e) NULL
)
instrument_formula <- tryCatch(
  stats::formula(stats::delete.response(stats::terms(model, component = "instruments"))),
  error = function(e) NULL
)
if (is.null(regression_formula) || is.null(instrument_formula)) {
  return(unavailable("Could not recover the fitted regression and instrument formulas from ivreg."))
}
needed <- unique(c(all.vars(regression_formula), all.vars(instrument_formula)))
data <- iv_analysis_frame(analysis_data, needed)
missing <- setdiff(needed, names(data))
if (length(missing)) {
  return(unavailable(paste(
    "Missing variables for Montiel Olea-Pflueger effective F:",
    paste(missing, collapse = ", ")
  )))
}
moment_data <- data[needed]
# Keep MOP on the fitted IV sample; deleting additional rows for missing
# values would change the calculation.
if (any(!stats::complete.cases(moment_data))) {
  return(unavailable("Regressors or instruments used by the fitted IV model contain missing values."))
}

stored_cluster <- attr(model, "cluster_state", exact = TRUE)
cluster <- if (
  length(stored_cluster) == stats::nobs(model) &&
  !anyNA(stored_cluster) && length(unique(stored_cluster)) >= 2L
) {
  as.vector(stored_cluster)
} else {
  iv_model_cluster(model, data)
}
covariance <- if (is.null(cluster)) "MDS" else "CL"
covariance_options <- if (is.null(cluster)) {
  list(type = "HCO")
} else {
  # MOPtest() uses clustered HCO moment covariance with its finite cluster
# adjustment here. The separate first-stage Wald statistic uses CR2

```

```

# covariance and the HTZ small sample reference.
list(
  cluster = data.frame(cluster = cluster),
  type = "HCO", cadjust = TRUE, multi0 = FALSE
)
}

result <- tryCatch({
  moment_model <- momentfit::momentModel(
    regression_formula,
    instrument_formula,
    data = moment_data,
    vcov = covariance,
    vcovOptions = covariance_options
  )
  # tau specifies the desired relative-bias benchmark for TSLS; the
  # simplified test uses the more conservative critical value.
  momentfit::MOPtest(
    moment_model,
    tau = tau,
    size = size,
    estMethod = "TSLS",
    simplified = TRUE,
    print = FALSE
  )
}, error = function(e) e)
if (inherits(result, "error")) return(unavailable(conditionMessage(result)))

values <- suppressWarnings(as.numeric(result[c("Feff", "critValue", "pvalue", "Keff"))])
names(values) <- c("statistic", "critical_value", "p.value", "effective_df")
if (length(values) != 4L || any(!is.finite(values))) {
  return(unavailable("momentfit::MOPtest() did not return finite MOP results."))
}

c(
  as.list(values),
  list(tau = tau, size = size, status = "estimated", reason = NA_character_)
)
}

```

Output: outputs/tables/appendix/appendix_iv_weak_inference.tex

Table 10: Inference with Weak Instruments for Long-Run Consumption Change

	Real log consumption change, 2004-05 to 2022-23	Real log consumption change, 2004-05 to 2023-24
EMI exposure	0.117 (0.182)	0.137 (0.281)
MOP effective F	0.69	0.41
AR p -value at zero	0.029	0.081
AR 95% accepted set	$[-0.193, -0.030]^{\dagger} \cup [0.009, 0.193]^{\dagger}$	$[-0.186, 0.186]^{\dagger}$
Direct effect threshold for $\beta = 0$ (share of reduced form)	16.6%	0.0%
Observations	524	522

The table reports conventional 2SLS coefficients with state-clustered standard errors. Instrument strength is summarized by the Montiel Olea–Pflueger effective F statistic, and Anderson–Rubin tests provide inference robust to weak instruments. The 95% AR accepted sets are displayed as unions of closed intervals. [†]A component reaches an evaluated search boundary; the displayed outer endpoint is the numerical search limit. The direct effect threshold is the smallest permitted direct effect, expressed as a share of the absolute reduced form, for which the AR set includes $\beta = 0$. Sources: Census of India 2001; NSS 61st Round, 2004–05; NSS 64th Round, 2007–08; HCES 2022–23; HCES 2023–24.

Rendering code: R/output/appendix_identification_exhibits.R and R/output/save_tables.R

Constructing Falsification-Adaptive Sets for Overidentified IV Designs

Under the exclusion relaxation used here, one endogenous regressor and multiple instruments assumed relevant yield a falsification-adaptive set spanning just-identified IV estimands. Each instrument serves in turn as the sole excluded instrument while the remaining instruments enter as controls. All constituent estimates use the same observations, requiring complete data on every variable in the registered specification, and inference is clustered by state. Every registered constituent enters the FAS calculation. Conditional first-stage statistics separately summarize evidence on instrument relevance.

File: R/iv/falsification_adaptive_set.R

```
# For the exclusion relaxation used here, each relevant instrument serves in
# turn as the sole excluded instrument while the remaining instruments enter as
# controls. The FAS spans the resulting just-identified IV estimates. Every
# registered constituent enters the set calculation; conditional first-stage
# statistics separately summarize evidence on instrument relevance.
iv_fas_component_formula <- function(specification, instrument) {
  spec <- as_single_iv_specification(specification)
  excluded <- plain_chr(unlist(spec$excluded_instruments[[1L]], use.names = FALSE))
  instrument <- plain_chr(instrument)[[1L]]
  if (!instrument %in% excluded) {
```

```

    stop("Selected FAS instrument is absent from the specification's excluded instruments.", call. = FALSE)
  }
included <- unique(c(
  unlist(spec$included_language_controls[[1L]], use.names = FALSE),
  unlist(spec$controls[[1L]], use.names = FALSE),
  setdiff(excluded, instrument)
))
make_iv_formula(
  spec$outcome[[1L]], spec$treatment[[1L]], instrument,
  controls = included,
  fixed_effects = iv_fixed_effect_terms(spec$fixed_effect[[1L]])
)
}

estimate_iv_fas_component <- function(data, specification, instrument) {
  spec <- as_single_iv_specification(specification)
  excluded <- plain_chr(unlist(spec$excluded_instruments[[1L]], use.names = FALSE))
  instrument <- plain_chr(instrument)[[1L]]
  other_instruments <- setdiff(excluded, instrument)
  # Resolve transformed formula terms such as factor(region) to the underlying
  # columns used to define the common estimation sample.
  needed <- iv_specification_variables(spec)
  missing <- setdiff(needed, names(data))
  if (length(missing)) {
    stop(
      "FAS specification ", spec$specification_id[[1L]],
      " is missing columns: ", paste(missing, collapse = ", "),
      call. = FALSE
    )
  }
  x <- as.data.frame(data)
  x <- x[stats::complete.cases(x[needed]), , drop = FALSE]
  if (nrow(x) < 3L) {
    return(data.frame(
      specification_id = spec$specification_id[[1L]],
      instrument = instrument, estimate = NA_real_, std.error = NA_real_,
      p.value = NA_real_, first_stage_estimate = NA_real_,
      first_stage_std.error = NA_real_, first_stage_f = NA_real_,
      first_stage_p.value = NA_real_, n = nrow(x), status = "not_estimated",
      reason = "Insufficient complete observations.", stringsAsFactors = FALSE
    ))
  }
}

```

```

}

formula <- iv_fas_component_formula(spec, instrument)
fit <- tryCatch(
  ivreg::ivreg(formula, data = x, model = TRUE, x = TRUE, y = TRUE),
  error = function(e) e
)
if (inherits(fit, "error")) {
  return(data.frame(
    specification_id = spec$specification_id[[1L]],
    instrument = instrument, estimate = NA_real_, std.error = NA_real_,
    p.value = NA_real_, first_stage_estimate = NA_real_,
    first_stage_std.error = NA_real_, first_stage_f = NA_real_,
    first_stage_p.value = NA_real_, n = nrow(x), status = "not_estimated",
    reason = conditionMessage(fit), stringsAsFactors = FALSE
  ))
}

cluster <- iv_specification_cluster(x, spec)
inference <- iv_clustered_inference(fit, cluster)
treatment_term <- model_term_inference(
  fit, spec$treatment[[1L]], inference$vcov
)

included <- unique(c(
  unlist(spec$included_language_controls[[1L]], use.names = FALSE),
  unlist(spec$controls[[1L]], use.names = FALSE),
  other_instruments,
  iv_fixed_effect_terms(spec$fixed_effect[[1L]])
))
first_stage <- stats::lm(
  stats::reformulate(unique(c(instrument, included)), response = spec$treatment[[1L]]),
  data = x
)
first_stage_inference <- iv_clustered_inference(first_stage, cluster)
first_stage_term <- model_term_inference(
  first_stage, instrument, first_stage_inference$vcov
)
first_stage_joint <- clustered_joint_wald_test(
  first_stage, instrument, cluster, inference = first_stage_inference
)

```

```

first_stage_f <- unname(first_stage_joint[["statistic"]])

treatment_ok <- all(is.finite(
  treatment_term[c("estimate", "std.error", "p.value")]
))

first_stage_ok <- all(is.finite(
  first_stage_term[c("estimate", "std.error", "p.value")]
))

status <- if (
  identical(inference$status, "estimated") &&
  identical(first_stage_inference$status, "estimated") &&
  treatment_ok && first_stage_ok
) "estimated" else "inference_unavailable"

reasons <- unique(na.omit(c(
  if (identical(inference$status, "unavailable")) inference$reason else NA_character_,
  if (identical(first_stage_inference$status, "unavailable")) first_stage_inference$reason else
    ↪ NA_character_,
  if (identical(inference$status, "estimated") && !treatment_ok)
    "Treatment coefficient inference is nonfinite." else NA_character_,
  if (identical(first_stage_inference$status, "estimated") && !first_stage_ok)
    "Conditional first-stage coefficient inference is nonfinite." else NA_character_
)))

data.frame(
  specification_id = spec$specification_id[[1L]],
  instrument = instrument,
  estimate = treatment_term[["estimate"]],
  std.error = treatment_term[["std.error"]],
  p.value = treatment_term[["p.value"]],
  first_stage_estimate = first_stage_term[["estimate"]],
  first_stage_std.error = first_stage_term[["std.error"]],
  first_stage_f = first_stage_f,
  first_stage_p.value = unname(first_stage_joint[["p.value"]]),
  n = stats::nobs(fit),
  status = status,
  reason = if (length(reasons)) paste(reasons, collapse = "; ") else NA_character_,
  stringsAsFactors = FALSE
)
}

estimate_iv_falsification_adaptive_set_spec <- function(data, specification) {

```

```

spec <- as_single_iv_specification(specification)
excluded <- plain_chr(unlist(spec$excluded_instruments[[1L]], use.names = FALSE))
if (spec$n_endogenous[[1L]] != 1L || length(excluded) <= 1L) {
  stop("FAS requires one endogenous regressor and multiple excluded instruments.", call. = FALSE)
}
components <- safe_bind_rows(lapply(excluded, function(instrument) {
  estimate_iv_fas_component(data, spec, instrument)
}))
estimated <- components$status == "estimated" & is.finite(components$estimate)
complete <- length(estimated) == length(excluded) && all(estimated)

lower <- if (complete) min(components$estimate) else NA_real_
upper <- if (complete) max(components$estimate) else NA_real_
lower_instrument <- if (complete) components$instrument[which.min(components$estimate)] else NA_character_
upper_instrument <- if (complete) components$instrument[which.max(components$estimate)] else NA_character_
first_stage_f <- num(components$first_stage_f)

summary <- data.frame(
  specification_id = spec$specification_id[[1L]],
  adjustment_id = spec$adjustment_id[[1L]],
  construction_id = spec$construction_id[[1L]],
  n_instruments = length(excluded),
  n_components_estimated = sum(estimated),
  fas_lower = lower,
  fas_upper = upper,
  fas_width = if (complete) upper - lower else NA_real_,
  fas_contains_zero = if (complete) lower <= 0 && upper >= 0 else NA,
  lower_endpoint_instrument = lower_instrument,
  upper_endpoint_instrument = upper_instrument,
  min_conditional_first_stage_f = if (all(is.finite(first_stage_f))) min(first_stage_f) else NA_real_,
  max_conditional_first_stage_f = if (all(is.finite(first_stage_f))) max(first_stage_f) else NA_real_,
  n_conditional_first_stage_f_below_10 = sum(is.finite(first_stage_f) & first_stage_f < 10),
  any_conditional_first_stage_f_below_10 = if (all(is.finite(first_stage_f))) any(first_stage_f < 10) else
    ↪ NA,
  n = if (nrow(components) && length(unique(components$n)) == 1L) components$n[[1L]] else NA_integer_,
  status = if (complete) "estimated" else "incomplete",
  reason = if (complete) NA_character_ else "At least one registered just-identified constituent could not
    ↪ be estimated.",
  stringsAsFactors = FALSE
)
list(summary = summary, components = components)

```

```
}
```

Price Harmonization and Real Consumption

Linking Historical and Later State CPI Series onto a Common Scale

Historical CPI-RL for rural areas and state CPI-IW for urban areas are linked to the later state rural/urban CPI (CPI-R/U). For each state and rural/urban sector, the link is the median monthly ratio of CPI-R/U to the historical CPI over January 2013–December 2014, reducing sensitivity to any one overlap month. Every historical series used before January 2013 must have at least six paired months in that interval. The link summary retains the number of paired months, first and last ratios, median absolute deviation, and range. The overlap interval estimates the link; the combined monthly series uses CPI-R/U itself from January 2013 onward and verifies that the intended CPI family appears on each side of the switch.

File: R/prices/build_temporal_price_series.R

```
# Link an earlier CPI series to a later one with the median ratio across valid
# overlap months.
price_link_factor <- function(old_index, new_index) {
  old <- num(old_index)
  new <- num(new_index)
  keep <- positive_finite(old) & positive_finite(new)
  if (!any(keep)) return(NA_real_)
  stats::median(new[keep] / old[keep], na.rm = TRUE)
}

build_temporal_price_series <- function(
  price_sources,
  switch_date = as.Date("2013-01-01"),
  overlap_start = as.Date("2013-01-01"),
  overlap_end = as.Date("2014-12-01"),
  minimum_link_months = 6L,
  pre_switch_start = NULL,
  pre_switch_end = NULL) {
  switch_date <- price_boundary(switch_date)
  old <- select_pre_2013_price_series(price_sources)
  new <- select_post_2013_price_series(price_sources)
  links <- summarise_price_links(old, new, overlap_start, overlap_end)
```

```

minimum_link_months <- as.integer(minimum_link_months)

if (length(minimum_link_months) != 1L || is.na(minimum_link_months) || minimum_link_months < 1L) {
  stop("minimum_link_months must be one positive integer.", call. = FALSE)
}

old_groups <- unique(old[old$period < switch_date, c("state_code", "sector"), drop = FALSE])
linked_groups <- links[links$link_months >= minimum_link_months, c("state_code", "sector"), drop = FALSE]
linked_key <- paste(linked_groups$state_code, linked_groups$sector, sep = "\r")
missing_key <- !paste(old_groups$state_code, old_groups$sector, sep = "\r") %in% linked_key
if (any(missing_key)) {
  bad <- old_groups[missing_key, , drop = FALSE]
  stop(
    "Pre-2013 price series have fewer than the required overlap months with CPI-R/U: ",
    paste(paste(bad$state_code, bad$sector, sep = "/"), collapse = ", "),
    call. = FALSE
  )
}

old <- merge(old, links, by = c("state_code", "sector"), all.x = TRUE, sort = FALSE)
old <- old[old$period < switch_date, , drop = FALSE]
if (!is.null(pre_switch_start) || !is.null(pre_switch_end)) {
  if (is.null(pre_switch_start) || is.null(pre_switch_end)) {
    stop("Both pre_switch_start and pre_switch_end are required when trimming the old price series.",
      call. = FALSE)
  }
  pre_switch_start <- price_boundary(pre_switch_start)
  pre_switch_end <- price_boundary(pre_switch_end)
  if (pre_switch_end < pre_switch_start || pre_switch_end >= switch_date) {
    stop("The retained historical interval must be ordered and end before the CPI switch date.", call. =
      FALSE)
  }
  old <- old[old$period >= pre_switch_start & old$period <= pre_switch_end, , drop = FALSE]
}

old$index_unlinked <- old$index
# Put each historical CPI series on the CPI-R/U scale before appending
# observations from January 2013 onward.
old$index <- num(old$index) * num(old$link_factor)

new <- new[new$period >= switch_date, , drop = FALSE]
new$link_factor <- 1

```

```

new$link_months <- NA_integer_
new$link_ratio_mad <- NA_real_
new$link_ratio_min <- NA_real_
new$link_ratio_max <- NA_real_
new$index_unlinked <- new$index

columns <- c(
  "state_code", "sector", "period", "index", "index_unlinked", "price_source",
  "link_factor", "link_months", "link_ratio_mad", "link_ratio_min", "link_ratio_max"
)

index <- rbind(old[columns], new[columns])
index$year <- as.integer(format(index$period, "%Y"))
index$month <- as.integer(format(index$period, "%m"))
index <- index[order(index$state_code, index$sector, index$period), , drop = FALSE]
rownames(index) <- NULL
validate_temporal_price_chain(index, switch_date)

structure(
  list(index = index, links = links, switch_date = switch_date),
  class = "emi_temporal_price_series"
)
}

```

Constructing Comparable Rural and Urban Price Deflators Across States

The code uses each state and rural/urban Tendulkar poverty line divided by the all-India rural 2011–12 line of 816 as its spatial price relative. Production supplies a separately validated 2011–12 reference CPI table; the helper can also calculate a reference mean directly when used independently. For each month, the deflator multiplies the Tendulkar spatial relative by the CPI relative to that reference period. At the reference period, a nominal amount equal to the applicable Tendulkar poverty line therefore equals 816 on the common real-consumption scale. When a state relies on a declared donor CPI series, the donor state’s reference CPI is used as well so the temporal numerator and denominator come from the same series. The figure provides a downstream check on the resulting real-consumption measures by comparing district estimates from the 2022–23 and 2023–24 Household Consumption Expenditure Survey rounds.

File: R/prices/price_deflators.R

```

build_state_sector_deflator <- function(
  temporal_index, spatial_relatives, reference_period = NULL, reference_index = NULL) {
  idx <- safe_df(temporal_index)
  spatial <- safe_df(spatial_relatives)
  validate_price_index(idx, keys = c("state_code", "sector", "period"))
  required_spatial <- c("state_code", "sector", "spatial_price_relative")
  missing <- setdiff(required_spatial, names(spatial))
  if (length(missing)) {
    stop("Spatial price relatives are missing columns: ", paste(missing, collapse = ", "), call. = FALSE)
  }
  if (anyDuplicated(spatial[c("state_code", "sector")])) {
    stop("Spatial price relatives must have one row per state and sector.", call. = FALSE)
  }
  if (any(!positive_finite(num(spatial$spatial_price_relative)))) {
    stop("Spatial price relatives must be positive and finite.", call. = FALSE)
  }

  if (is.null(reference_index)) {
    if (is.null(reference_period)) stop("A price reference period or reference-index table is required.",
      ↪ call. = FALSE)
    reference_period <- as.Date(reference_period)
    ref <- idx[idx$period %in% reference_period, , drop = FALSE]
    if (!nrow(ref)) stop("No observations fall in the requested price reference period.", call. = FALSE)
    ref_mean <- stats::aggregate(index ~ state_code + sector, ref, function(x) mean(num(x), na.rm = TRUE))
    names(ref_mean)[names(ref_mean) == "index"] <- "reference_index"
  } else {
    ref_mean <- safe_df(reference_index)
    required_reference <- c("state_code", "sector", "reference_index")
    missing_reference <- setdiff(required_reference, names(ref_mean))
    if (length(missing_reference)) {
      stop("Reference price table is missing columns: ", paste(missing_reference, collapse = ", "), call. =
        ↪ FALSE)
    }
    if (anyDuplicated(ref_mean[c("state_code", "sector")]) ||
      ↪ any(!positive_finite(ref_mean$reference_index))) {
      stop("Reference price table must have one positive row per state and rural/urban sector.", call. =
        ↪ FALSE)
    }
    ref_mean <- ref_mean[required_reference]
  }
  out <- merge(

```

```

    idx,
    spatial,
    by = c("state_code", "sector"), all.x = TRUE, sort = FALSE
  )
  out <- merge(out, ref_mean, by = c("state_code", "sector"), all.x = TRUE, sort = FALSE)
  if ("source_state_code" %in% names(out) && any(!positive_finite(out$reference_index))) {
    # When a state and sector use a declared donor CPI series, use the donor's
    # reference period index as well so both terms come from the same CPI
    # series.
    donor_reference <- ref_mean
    names(donor_reference)[names(donor_reference) == "state_code"] <- "source_state_code"
    names(donor_reference)[names(donor_reference) == "reference_index"] <- "donor_reference_index"
    out <- merge(out, donor_reference, by = c("source_state_code", "sector"), all.x = TRUE, sort = FALSE)
    missing_reference <- !positive_finite(out$reference_index)
    out$reference_index[missing_reference] <- out$donor_reference_index[missing_reference]
    out$donor_reference_index <- NULL
  }

  # The product adjusts for the rural/urban price level in each state and for
  # price change from the reference period to the survey month.
  out$temporal_price_relative <- num(out$index) / num(out$reference_index)
  out$price_deflator <- num(out$spatial_price_relative) * num(out$temporal_price_relative)
  if (any(!positive_finite(out$price_deflator))) {
    stop("Price deflators must be positive and finite for every retained state, sector, and month.", call. =
      ↪ FALSE)
  }

  out <- out[order(out$state_code, out$sector, out$period), , drop = FALSE]
  rownames(out) <- NULL
  out
}

```

Output: outputs/figures/appendix/appendix_consumption_hces_consistency.pdf

Rendering code: R/output/appendix_data_construction_exhibits.R and R/output/save_figures.R




Figure 8: Agreement between district real-consumption estimates in the 2022-23 and 2023-24 HCES rounds. Each panel uses districts retained in both rounds under the consumption sample criteria; the dashed line is the 45-degree line. Sources: National Sample Survey Office, HCES 2022-23 and HCES 2023-24 (National Sample Survey Office 2024, 2025).

Testing CPI Links Against Known Answers

A small rural example has monthly ratios of the later CPI to the historical CPI equal to 2, 2, and 3, so the median link is 2 while the final overlap ratio is 3. The urban example has a constant ratio of 0.5. The test verifies the median link, the first and last overlap ratios, and the number of paired months.

File: tests/testthat/test-prices-and-controls.R

```
test_that("CPI link factors equal median overlap ratios", {
  months <- as.Date(c("2013-01-01", "2013-02-01", "2013-03-01"))
  # Rural monthly ratios are 2, 2, and 3, so the median link is 2.
  old <- rbind(
    data.frame(state_code = "A", sector = "rural", period = months, index = c(100, 110, 120)),
    data.frame(state_code = "A", sector = "urban", period = months, index = c(200, 220, 240))
  )
  new <- rbind(
    data.frame(state_code = "A", sector = "rural", period = months, index = c(200, 220, 360)),
    data.frame(state_code = "A", sector = "urban", period = months, index = c(100, 110, 120))
  )
  links <- summarise_price_links(old, new, months[1], months[3])
  expect_equal(links$link_factor, c(2, 0.5))
  expect_equal(links$link_factor_first, c(2, 0.5))
  expect_equal(links$link_factor_last, c(3, 0.5))
})
```

```
expect_equal(links$link_months, c(3L, 3L))
})
```

Spatial Presentation on a Common Geography

Preserving All 2001 Census Districts in Published Maps

Published district maps begin with the full reviewed 2001 Census district geometry. Analysis values are joined by district identifier, so every district remains in the figure and missing estimates appear as the explicit “No data” category. State codes are derived from the 2001 Census district identifiers and checked against any state codes supplied with the analysis values. The completed district table is then passed to the shared geographic figure builder.

File: R/output/save_figures.R

```
complete_public_map_geometry <- function(district_panel, map_geometry) {
  if (!has_sf_geometry(map_geometry)) {
    stop("District maps require the complete 2001 Census district geometry.", call. = FALSE)
  }
  key <- "target_unit_2001"
  if (!key %in% names(map_geometry) && "unit_id" %in% names(map_geometry)) {
    map_geometry[[key]] <- map_geometry$unit_id
  }
  if (!key %in% names(map_geometry) || !key %in% names(district_panel)) {
    stop("Map geometry and analysis data must contain a 2001 Census district key.", call. = FALSE)
  }
  if (anyDuplicated(map_geometry[[key]])) {
    stop("Complete 2001 Census map geometry must contain one row per district.", call. = FALSE)
  }
  attributes <- if (inherits(district_panel, "sf")) sf::st_drop_geometry(district_panel) else
  ↪ safe_df(district_panel)
  if (anyDuplicated(attributes[[key]])) {
    stop("Analysis data must contain at most one row per 2001 Census district.", call. = FALSE)
  }
  out <- merge(map_geometry, attributes, by = key, all.x = TRUE, sort = FALSE)

  # Derive state codes from the 2001 Census district identifiers so state
  # outlines include districts with missing analysis values.
  boundary_state <- public_map_state_code_2001(out[[key]])
}
```

```

if ("state_code_2001" %in% names(out)) {
  panel_state <- plain_chr(out$state_code_2001)
  mismatch <- !is.na(panel_state) & nzchar(panel_state) & panel_state != boundary_state
  if (any(mismatch)) {
    stop("Analysis state codes disagree with the state codes in the 2001 Census district identifiers.",
      call. = FALSE)
  }
}
out$state_code_2001 <- boundary_state
out
}

save_map_figure <- function(spec, path_base, district_panel, map_geometry, boundary_reference, formats) {
  if (is.null(spec$variable) || !spec$variable %in% names(district_panel)) {
    stop("Map figure '", spec$name, "' is missing variable '", spec$variable, "'.", call. = FALSE)
  }

  plot_data <- complete_public_map_geometry(district_panel, map_geometry)
  plot_data <- prepare_public_map_data(plot_data, spec$variable)
  p <- build_public_ggplot_map(plot_data, spec, boundary_reference = boundary_reference)
  save_plot_formats(p, path_base, formats, width = 7.2, height = 5.2, dpi = 300)
}

```

Outputs: outputs/figures/main/map_emi_share_enrolled.pdf, outputs/figures/main/map_emi_exposure.pdf, outputs/figures/main/map_public_emi_exposure.pdf, and outputs/figures/main/map_private_emi_exposure.pdf

Rendering code: R/output/make_figures.R




Figure 3: District distribution of English-medium schooling measures on 2001 Census boundaries. All panels use a common percentage scale. Source: NSS 64th Round, 2007–08 (National Sample Survey Office 2011). Tengshe (2014) and Natural Earth (2022) provide geospatial data. District boundaries reflect 2001 geography. Politically disputed regions are identified using Natural Earth, and those without data—Aksai Chin, Azad Kashmir, Gilgit-Baltistan, the Trans-Karakoram Tract, and Siachen Glacier—are colored light gray. Their inclusion does not imply a position on territorial disputes.