

Code Sample

UNDER-10MB COPY

Anonymous

This document contains excerpts from the replication code of a paper titled “Inequality in a Potential Equalizer of Opportunity: Variation in the Accessibility and Net Benefits of English-Medium Instruction in India,” with figures selectively rasterized to reduce file size. Excerpts are followed by their corresponding outputs when appropriate. The project may have been updated since this file was generated; the most up-to-date versions of this document, the full paper, and its replication package are available online, and I would be happy to provide links for those interested. This PDF can be generated by running `make samples` or `scripts/run_full_build.sh` from the repo root.

Historical Geography and Measurement

Reconstructing District Boundary Transfers Between 1991 and 2001

A historical Census table reports how each 1991 district’s population was apportioned among 2001 districts after boundary changes. The headerless source also splits long district names across physical continuation rows. The reader repairs only continuation patterns supported by the table’s structure, retains continuation rows that contain genuine transfer records, converts transfer shares to numbers, and requires the shares allocated from every 1991 source district to total 100% within 0.05 percentage points.

File: `R/io/read_inputs.R`

```
#' Rejoin district names split across rows in the historical boundary table
#'
#' Rows with a blank 1991 population field are continuations. A trailing hyphen
#' identifies a split 1991 district name; a 2001 name can continue on a row with
#' no transfer percentages.
repair_district_carveout_wrapped_rows <- function(x) {
```

```

out <- safe_df(x)
if (!nrow(out)) return(out)
blank <- function(value) is.na(value) || !nzchar(trimws(as.character(value)))
drop <- rep(FALSE, nrow(out))
for (i in seq_len(nrow(out))) {
  if (i == 1L || !blank(out$pop_1991[[i]])) next
  previous <- i - 1L
  source_continuation <- !blank(out$district_1991[[i]])
  target_continuation <- !blank(out$district_2001[[i]])
  has_transfer <- !blank(out$pct_01in91[[i]]) || !blank(out$pct_91in01[[i]])

  if (source_continuation) {
    previous_name <- plain_chr(out$district_1991[[previous]])
    if (!length(previous_name) || !grepl("-", trimws(previous_name[[1L]]))) {
      stop("Unexpected wrapped row in the district boundary table.", call. = FALSE)
    }
    continuation <- plain_chr(out$district_1991[[i]])[[1L]]
    combined <- join_district_carveout_wrapped_label(previous_name[[1L]], continuation)
    out$district_1991[[previous]] <- combined
    if (has_transfer) {
      out$district_1991[[i]] <- combined
      out$pop_1991[[i]] <- out$pop_1991[[previous]]
    }
  }

  if (!has_transfer && (source_continuation || target_continuation)) {
    if (target_continuation) {
      if (blank(out$district_2001[[previous]])) {
        stop("Unexpected continuation of a 2001 district name in the district boundary table.", call. =
          FALSE)
      }
      previous_target <- trimws(as.character(out$district_2001[[previous]]))
      continuation <- trimws(as.character(out$district_2001[[i]]))
      out$district_2001[[previous]] <- join_district_carveout_wrapped_label(
        previous_target, continuation
      )
    }
    drop[[i]] <- TRUE
  }
}
out[!drop, , drop = FALSE]

```

```

}

district_carveout_rounding_tolerance_pp <- function() 0.05

#' Validate population shares in the 1991-to-2001 district transfer table
#'
validate_district_carveout_shares <- function(
  x,
  rounding_tolerance_pp = district_carveout_rounding_tolerance_pp()) {
  out <- safe_df(x)
  if (!nrow(out)) return(invisible(TRUE))
  share_cols <- c("pct_01in91", "pct_91in01")
  for (field in share_cols) {
    values <- num(out[[field]])
    bad <- is.finite(values) & (values < 0 | values > 100)
    if (any(bad)) {
      stop("District transfer shares must lie in [0, 100].", call. = FALSE)
    }
  }
  source_share <- num(out$pct_01in91)
  if (any(!is.finite(source_share))) {
    stop("Shares of each 1991 district assigned to 2001 districts must be complete.", call. = FALSE)
  }
  source_key <- paste(out$district_1991, out$pop_1991, sep = "__")
  source_sum <- vapply(split(source_share, source_key), sum, numeric(1))
  if (any(abs(source_sum - 100) > rounding_tolerance_pp)) {
    stop("Shares of each 1991 district assigned to 2001 districts must sum to 100 within rounding
    ↪ tolerance.", call. = FALSE)
  }
  invisible(TRUE)
}

read_district_carveouts <- function(path) {
  out <- utils::read.csv(
    path,
    header = FALSE,
    col.names = c("district_1991", "pop_1991", "district_2001", "pct_01in91", "pct_91in01"),
    stringsAsFactors = FALSE,
    na.strings = c("", "NA"),
    check.names = FALSE
  )
}

```

```

out <- repair_district_carveout_wrapped_rows(out)
out$district_1991 <- fill_down_missing(out$district_1991)
out$pop_1991 <- num(gsub(",", "", fill_down_missing(out$pop_1991), fixed = TRUE))
out$pct_01in91 <- num(out$pct_01in91)
out$pct_91in01 <- num(out$pct_91in01)
validate_district_carveout_shares(out)
out
}

```

Generating High-Precision District-Match Candidates

District names vary across survey waves and administrative vintages, so standardized names generate candidates rather than establish identity by themselves. Within each state, the code scores candidate pairs using Jaro–Winkler similarity, normalized Damerau–Levenshtein similarity, trigram cosine similarity, and token Jaccard similarity. A high-precision flag additionally requires reciprocal nearest-neighbor agreement, compatible directional tokens, preferred reference vintage, metric-specific thresholds, and separation from the runner-up; downstream lineage review still decides whether a candidate becomes an accepted geographic link.

File: R/districts/lineage_matching.R

```

score_match_candidates <- function(source_roster, reference_units, excluded_source_ids = character()) {
  need_pkg("stringdist", "district-lineage candidate scores")
  source_roster <- safe_df(source_roster)
  reference_units <- safe_df(reference_units)
  if (!"raw_district" %in% names(source_roster)) source_roster$raw_district <- source_roster$district_std
  source_open <- source_roster[!source_roster$source_row_id %in% excluded_source_ids, , drop = FALSE]
  if (!nrow(source_open)) return(empty_source_candidates())

  pairs <- merge(
    source_open, reference_units,
    by = "state_std", all = FALSE, sort = FALSE,
    suffixes = c("_source", "_candidate")
  )
  if (!nrow(pairs)) return(empty_source_candidates())
  source_name <- pairs$district_std_source
  candidate_name <- pairs$district_std_candidate
  pairs$jwt <- 1 - stringdist::stringdist(source_name, candidate_name, method = "jw", p = 0.1)
  dl_distance <- stringdist::stringdist(source_name, candidate_name, method = "dl")
  denom <- pmax(nchar(source_name), nchar(candidate_name), 1L)
}

```

```

pairs$dl <- pmax(0, 1 - dl_distance / denom)
pairs$trigram <- stringdist::stringsim(source_name, candidate_name, method = "cosine", q = 3)
pairs$token <- token_jaccard_similarity(source_name, candidate_name)
pairs$score <- rowMeans(pairs[c("jw", "dl", "trigram", "token")], na.rm = TRUE)

ranked <- safe_bind_rows(lapply(split(seq_len(nrow(pairs)), pairs$source_row_id), function(i) {
  x <- pairs[i, , drop = FALSE]
  preference <- vintage_preference(x$wave[[1]])
  vintage_rank <- match(x$reference_vintage, preference, nomatch = length(preference) + 1L)
  x <- x[order(-x$score, -x$jw, vintage_rank, x$unit_id), , drop = FALSE]
  x$rank <- match(x$district_std_candidate, unique(x$district_std_candidate))
  distinct_scores <- x$score[!duplicated(x$district_std_candidate)]
  second <- if (length(distinct_scores) >= 2L) distinct_scores[[2]] else -Inf
  x$margin <- ifelse(x$rank == 1L, x$score - second, NA_real_)
  x$preferred_vintage <- !duplicated(x$district_std_candidate)
  x
}))

candidate_groups <- split(seq_len(nrow(ranked)), ranked$unit_id)
best_by_candidate <- stats::setNames(
  vapply(candidate_groups, function(i) {
    score <- ranked$score[i]
    ranked$source_row_id[i[which.max(score)]]
  }, character(1)),
  names(candidate_groups)
)

ranked$reciprocal_nearest <- ranked$rank == 1L &
  unname(best_by_candidate[ranked$unit_id]) == ranked$source_row_id
threshold <- district_match_candidate_thresholds()
ranked$high_precision_candidate <- ranked$rank == 1L & ranked$preferred_vintage &
  ranked$reciprocal_nearest &
  directional_tokens_compatible(ranked$district_std_source, ranked$district_std_candidate) &
  ranked$jw >= threshold[["jw"]] & ranked$dl >= threshold[["dl"]] &
  ranked$trigram >= threshold[["trigram"]] & ranked$margin >= threshold[["margin"]]
ranked <- ranked[ranked$rank <= 5L, , drop = FALSE]

data.frame(
  source_row_id = ranked$source_row_id,
  wave = ranked$wave,
  source_code = ranked$source_code,
  state_std = ranked$state_std,

```

```

source_name_raw = ranked$raw_district,
source_name = ranked$district_std_source,
candidate_unit = ranked$unit_id,
candidate_name = ranked$district_std_candidate,
candidate_source_id = ranked$source_id,
reference_vintage = ranked$reference_vintage,
candidate_method = "fuzzy_name_candidate",
rank = ranked$rank,
jw = ranked$jw,
dl = ranked$dl,
trigram = ranked$trigram,
token = ranked$token,
score = ranked$score,
margin = ranked$margin,
reciprocal_nearest = ranked$reciprocal_nearest,
high_precision_candidate = ranked$high_precision_candidate,
stringsAsFactors = FALSE
)
}

```

Reconstructing District Data from the 1991 Language Atlas

The 1991 Census Language Atlas repeats each district across eight language pages. Districts are identified by a Census state code plus a within-state serial, but some repeated rows omit the serial or contain errors introduced during text extraction. When a row itself has no serial, the code recovers one only if the same raw row label within the same Atlas block has exactly one distinct serial across repeated pages. The reconstructed code is then checked against the official 1991 Primary Census Abstract, and a candidate is promoted only when its Atlas population is within 1% of the corresponding Census population.

File: `scripts/build_language_atlas_1991.py`

```

def cross_page_serial_evidence(cells: list[dict[str, object]]) -> dict[tuple[int, str], tuple[int, ...]]:
    """Collect district serials printed with repeated Language Atlas row labels.

The Atlas prints each district on eight language pages. The district serial is
the within-state number paired with the Census state code. If one copy omits
that number, repeated copies can supply it only when their printed serials agree.
    """

```

```

observed: dict[tuple[int, str], set[int]] = {}
unique_rows = {
    (int(cell["block"]), int(cell["page"]), int(cell["row_sequence"]),
     str(cell["row_label_raw"]).strip(), str(cell["serial_candidate"]).strip())
    for cell in cells
}

for block, _page, _sequence, label, serial_raw in unique_rows:
    if not label or label == "INDIA":
        continue
    serial = int(serial_raw) if serial_raw else parse_leading_serial(label)
    if serial is not None:
        observed.setdefault((block, label), set()).add(serial)
return {key: tuple(sorted(values)) for key, values in observed.items()}

def reconcile_district_serial(
    cell: dict[str, object],
    evidence: dict[tuple[int, str], tuple[int, ...]],
) -> tuple[int | None, str, str]:
    """Choose a district serial from the row itself or agreeing repeated copies."""
    label = str(cell["row_label_raw"]).strip()
    serial_column = int(cell["serial_candidate"]) if str(cell["serial_candidate"]).strip() else None
    serial_label = parse_leading_serial(label)
    repeated = evidence.get((int(cell["block"]), label), ())
    repeated_text = ";".join(map(str, repeated))

    if serial_column is not None and serial_label is not None and serial_column != serial_label:
        return None, "direct_conflict", repeated_text
    direct = serial_column if serial_column is not None else serial_label
    direct_source = "serial_column" if serial_column is not None else "label_prefix"

    if direct is not None:
        # Prefer a serial printed on the current row. Use repeated copies of the
        # same label only when the current row has no serial because continuation
        # pages contain occasional OCR errors.
        return direct, direct_source, repeated_text
    if len(repeated) == 1:
        return repeated[0], "cross_page_exact_label", repeated_text
    if len(repeated) > 1:
        return None, "cross_page_conflict", repeated_text
    return None, "missing", repeated_text

```

```

def build_district_population_validation(
    cells: list[dict[str, object]],
    state_crosswalk: dict[str, dict[str, str]],
    pca91: dict[tuple[str, str], int],
) -> list[dict[str, object]]:
    rows = merge_population_continuations(cells, state_crosswalk)
    serial_evidence = cross_page_serial_evidence(cells)
    seen_state_codes: set[str] = set()
    current_state: dict[str, str] | None = None
    candidates: list[dict[str, object]] = []
    for cell in rows:
        label = str(cell["row_label_raw"]).strip()
        if label == "INDIA":
            continue
        mapped_state = state_crosswalk.get(label)
        if mapped_state is not None and mapped_state["state_code_1991"] not in seen_state_codes:
            current_state = mapped_state
            seen_state_codes.add(mapped_state["state_code_1991"])
            continue

        serial, serial_source, cross_page_serials = reconcile_district_serial(cell, serial_evidence)

        state_code = current_state["state_code_1991"] if current_state else ""
        state_name = current_state["state_name_1991"] if current_state else ""
        district_code = f"{serial:02d}" if serial is not None else ""
        pca_population = pca91.get((state_code, district_code)) if state_code and district_code else None
        atlas_population = (
            int(cell["speaker_count_candidate"])
            if str(cell["speaker_count_candidate"]).strip()
            else None
        )
        if atlas_population is not None and pca_population is not None and pca_population > 0:
            abs_diff = abs(atlas_population - pca_population)
            rel_diff = abs_diff / pca_population
        else:
            abs_diff = None
            rel_diff = None
        candidates.append({
            "block": cell["block"],

```

```

        "page": cell["page"],
        "row_sequence": cell["row_sequence"],
        "row_sequence_end": cell["row_sequence_end"],
        "row_merge_status": cell["row_merge_status"],
        "row_label_raw": label,
        "state_code_1991": state_code,
        "state_name_1991": state_name,
        "district_code_1991": district_code,
        "serial_source": serial_source,
        "cross_page_serials": cross_page_serials,
        "atlas_population_raw": cell["raw_value"],
        "atlas_population_candidate": "" if atlas_population is None else atlas_population,
        "pca91_population": "" if pca_population is None else pca_population,
        "population_abs_diff": "" if abs_diff is None else abs_diff,
        "population_relative_diff": "" if rel_diff is None else round(rel_diff, 8),
        "identity_status": "",
        "population_status": "",
        "promotion_status": "review_required",
    })

expected_state_codes = {row["state_code_1991"] for row in state_crosswalk.values()}
if seen_state_codes != expected_state_codes:
    missing = sorted(expected_state_codes - seen_state_codes)
    extra = sorted(seen_state_codes - expected_state_codes)
    raise ValueError(
        "Language Atlas state headings do not match the expected 1991 states; "
        f"missing={missing}; extra={extra}"
    )

pca_state_codes = {state_code for state_code, _ in pca91}
if pca_state_codes != expected_state_codes:
    raise ValueError("Language Atlas state names and 1991 Primary Census Abstract state codes disagree")

key_counts: dict[tuple[str, str], int] = {}
for row in candidates:
    key = (str(row["state_code_1991"]), str(row["district_code_1991"]))
    if all(key):
        key_counts[key] = key_counts.get(key, 0) + 1

for row in candidates:
    key = (str(row["state_code_1991"]), str(row["district_code_1991"]))
    if not row["state_code_1991"]:

```

```

        identity_status = "missing_state_context"
    elif row["serial_source"] in {"direct_conflict", "cross_page_conflict"}:
        identity_status = "serial_conflict"
    elif not row["district_code_1991"]:
        identity_status = "missing_district_serial"
    elif key_counts.get(key, 0) > 1:
        identity_status = "duplicate_district_code"
    elif key not in pca91:
        identity_status = "missing_pca_district"
    else:
        identity_status = "candidate_official_code"
    row["identity_status"] = identity_status

    if row["atlas_population_candidate"] == "":
        population_status = "atlas_population_unparsed"
    elif row["pca91_population"] == "":
        population_status = "pca_population_missing"
    elif float(row["population_relative_diff"]) <= POPULATION_RELATIVE_TOLERANCE:
        population_status = "within_1pct"
    else:
        population_status = "over_1pct"
    row["population_status"] = population_status
    if identity_status == "candidate_official_code" and population_status == "within_1pct":
        row["promotion_status"] = "population_validated_candidate"
return candidates

```

Bounding Historical Linguistic Distance with Unresolved Source Data

Some 1991 Language Atlas population remains unresolved after language-name reconciliation. Rather than as-signing that mass implicitly or dropping the affected district, the code bounds the speaker-weighted distance measure. It separates accepted positive-distance speakers from known zero-distance and native-English speakers, then lets unresolved population either leave the positive-distance mean unchanged or lie anywhere on the admis-sible positive-distance support. A source-only quality rule subsequently requires at least 99% accepted speaker coverage and no more than half a Shastry degree of bound width before the preferred historical measure enters the analysis.

File: R/measures/build_linguistic_distance_iv.R

```

historical_linguistic_population_coherent <- function(speakers, population) {
  speakers <- num(speakers)
  if (!is.numeric(population) || length(population) != 1L ||
      !is.finite(population) || population <= 0) {
    stop("Historical linguistic bounds require a positive district population.", call. = FALSE)
  }
  accepted <- is.finite(speakers) & speakers >= 0
  accepted_total <- sum(speakers[accepted], na.rm = TRUE)
  isTRUE(accepted_total <= population)
}

historical_linguistic_distance_bounds <- function(
  speakers, degree, language, population, nonzero_degree_range = c(1, 5)) {
  speakers <- num(speakers)
  degree <- num(degree)
  language <- normalize_language_label(language)
  population_coherent <- historical_linguistic_population_coherent(speakers, population)
  accepted <- is.finite(speakers) & speakers >= 0
  english <- accepted & language == "English"
  known_nonzero <- accepted & is.finite(degree) & degree > 0 & !english
  known_zero <- accepted & is.finite(degree) & degree == 0 & !english
  known_irrelevant <- known_zero | english

  nonzero_speakers <- sum(speakers[known_nonzero], na.rm = TRUE)
  numerator <- sum(speakers[known_nonzero] * degree[known_nonzero], na.rm = TRUE)
  known_irrelevant_speakers <- sum(speakers[known_irrelevant], na.rm = TRUE)
  resolved_speakers <- nonzero_speakers + known_irrelevant_speakers
  unresolved_mass <- max(0, population - resolved_speakers)
  point <- if (nonzero_speakers > 0) numerator / nonzero_speakers else NA_real_

  nonzero_degree_range <- sort(unique(num(nonzero_degree_range)))
  if (length(nonzero_degree_range) < 2L || any(!is.finite(nonzero_degree_range)) ||
      any(nonzero_degree_range <= 0)) {
    stop("Historical linguistic-distance bounds require a finite positive degree range.", call. = FALSE)
  }
  degree_min <- min(nonzero_degree_range)
  degree_max <- max(nonzero_degree_range)
  if (!population_coherent) {
    point <- lower <- upper <- width <- unresolved_mass <- NA_real_
  } else if (!is.finite(point) || !is.finite(degree_min) || !is.finite(degree_max)) {
    lower <- upper <- width <- NA_real_
  }
}

```

```

} else if (unresolved_mass <= 0) {
  lower <- upper <- point
  width <- 0
} else {
  lower_if_nonzero <- (numerator + unresolved_mass * degree_min) /
    (nonzero_speakers + unresolved_mass)
  upper_if_nonzero <- (numerator + unresolved_mass * degree_max) /
    (nonzero_speakers + unresolved_mass)
  # Unresolved population may instead be Hindi/English/unmapped and therefore
  # leave the accepted-speaker mean unchanged. Include both possibilities.
  lower <- min(point, lower_if_nonzero)
  upper <- max(point, upper_if_nonzero)
  width <- upper - lower
}

list(
  point = point,
  lower = lower,
  upper = upper,
  width = width,
  nonzero_speakers = nonzero_speakers,
  known_irrelevant_speakers = known_irrelevant_speakers,
  resolved_speakers = resolved_speakers,
  unresolved_mass_upper_bound = unresolved_mass,
  population_coherent = population_coherent,
  degree_min = degree_min,
  degree_max = degree_max
)
}

```

Output: outputs/figures/appendix/appendix_historical_language_persistence.pdf

Rendering code: R/output/appendix_data_construction_exhibits.R and R/output/save_figures.R




Figure 2: Persistence of speaker-weighted linguistic distance from Hindi from 1991 to 2001. Each point is a 2001 Census district with a reconstructed measure in both years; the dashed line is the 45-degree line. Sources: Census of India 1991 Language Atlas (Office of the Registrar General and Census Commissioner, India 2004) and Census of India 2001 language tables (Office of the Registrar General and Census Commissioner of India 2001).

Independent Schooling Measures and Survey Estimation

Reconstructing English-Medium Enrollment from Archived DISE Reports

District Information System for Education (DISE) report cards encode medium-of-instruction totals in more than one layout. The parser accepts either language rows under an identifiable total column or language columns paired with an explicit total row, and refuses to infer counts when those structural cues are missing. Extraction is restricted to PDF pages whose district-year provenance has already been reviewed; when multiple archived reports cover the same district-year, an explicit priority must select one unique source.

File: `scripts/build_dise_report_language_enrollment.py`

```
LANGUAGES = ("english", "hindi")
NUMBER = re.compile(r"(?<![\w.]) (?:\d{1,3} (?:,\d{3})+|\d+)(?![\w.])")
MEDIUM_HEADING = re.compile(r"medium\s+of\s+instruction", re.IGNORECASE)

def medium_block(text: str, radius: int = 28) -> list[str]:
    lines = text.splitlines()
    hit = next((i for i, line in enumerate(lines) if MEDIUM_HEADING.search(line)), None)
    if hit is None:
```

```

        raise ValueError("page has no 'medium of instruction' heading")
    return lines[hit : min(len(lines), hit + radius)]

def _has_total_header(lines: list[str], row_index: int) -> bool:
    start = max(0, row_index - 6)
    return any(re.search(r"\btotal\b", line, re.IGNORECASE) for line in lines[start:row_index])

def parse_languages_as_rows(lines: list[str]) -> dict[str, int | None] | None:
    values: dict[str, int | None] = {language: None for language in LANGUAGES}
    matched = False
    for i, line in enumerate(lines):
        stripped = line.strip()
        language = next(
            (name for name in LANGUAGES if re.match(rf"^{name}\b", stripped, re.IGNORECASE)),
            None,
        )
        if language is None:
            continue
        numbers = [parse_int(m.group(0)) for m in NUMBER.finditer(stripped)]
        if not numbers:
            continue
        if not _has_total_header(lines, i):
            raise ValueError(
                f"{language} row found without a nearby Total column; refusing to guess"
            )
        values[language] = numbers[-1]
        matched = True
    return values if matched else None

def _column_positions(header: str) -> dict[str, int]:
    positions = {}
    lower = header.lower()
    for language in LANGUAGES:
        pos = lower.find(language)
        if pos >= 0:
            positions[language] = pos
    return positions

```

```

def _nearest_number_at_column(line: str, position: int) -> int | None:
    matches = list(NUMBER.finditer(line))
    if not matches:
        return None
    nearest = min(matches, key=lambda m: abs(m.start() - position))
    return parse_int(nearest.group(0))

def parse_languages_as_columns(lines: list[str]) -> dict[str, int | None] | None:
    header_index = None
    positions: dict[str, int] = {}
    for i, line in enumerate(lines):
        candidate = _column_positions(line)
        if candidate:
            header_index = i
            positions = candidate
            break
    if header_index is None:
        return None

    total_line = next(
        (
            line
            for line in lines[header_index + 1 :]
            if re.match(r"^\\s*(?:grand\\s+)?total\\b", line, re.IGNORECASE)
        ),
        None,
    )
    if total_line is None:
        raise ValueError("language columns found without an explicit Total row; refusing to guess")

    return {
        language: _nearest_number_at_column(total_line, position)
        for language, position in positions.items()
    } | {language: None for language in LANGUAGES if language not in positions}

def parse_language_counts(text: str) -> dict[str, int | None]:
    lines = medium_block(text)
    parsed = parse_languages_as_rows(lines)

```

```

if parsed is None:
    parsed = parse_languages_as_columns(lines)
if parsed is None or parsed.get("english") is None:
    raise ValueError("could not resolve English enrollment from medium-of-instruction table")
return parsed

def preferred_rows(rows: Iterable[dict[str, str]]) -> list[dict[str, str]]:
    grouped: dict[tuple[str, str, str], list[dict[str, str]]] = defaultdict(list)
    for row in rows:
        key = (
            row["academic_year"].strip(),
            row["state_report"].strip(),
            row["district_report"].strip(),
        )
        grouped[key].append(row)

    selected = []
    for key, group in grouped.items():
        priorities = [int(row["report_priority"]) for row in group]
        best = min(priorities)
        winners = [row for row in group if int(row["report_priority"]) == best]
        if len(winners) != 1:
            raise ValueError(f"non-unique preferred report provenance for {key}")
        selected.append(winners[0])
    return sorted(
        selected,
        key=lambda row: (
            row["academic_year"],
            row["state_report"],
            row["district_report"],
        ),
    )

def rebuild(
    metadata: Path,
    registry: Path,
    archive_root: Path,
    pdftotext: str,
) -> list[dict[str, str]]:

```

```

report_map = registered_reports(registry)
rebuilt = []
for row in preferred_rows(read_csv(metadata)):
    source = row["source_pdf"].strip()
    if source not in report_map:
        raise ValueError(f"unregistered source_pdf in report metadata: {source}")
    pdf = archive_root / report_map[source]
    if not pdf.is_file():
        raise FileNotFoundError(f"missing registered DISE report PDF: {pdf}")
    page = int(row["source_page"])
    counts = parse_language_counts(extract_pdf_page(pdf, page, pdftotext))
    out = dict(row)
    out["english_enrollment"] = str(counts["english"])
    out["hindi_enrollment"] = "" if counts["hindi"] is None else str(counts["hindi"])
    rebuilt.append(out)
return rebuilt

```

Output: outputs/figures/appendix/appendix_nss_dise_agreement.pdf

Rendering code: R/output/appendix_data_construction_exhibits.R and R/output/save_figures.R




Figure 7: Agreement between NSS and DISE English-medium measures. The left panel compares district levels; the right panel subtracts each source's state mean before comparison. Sources: NSS 64th Round, 2007–08 (National Sample Survey Office 2011) and DISE District Report Cards, 2007–08 (District Information System for Education 2017).

Constructing the NSS Survey Design

The 2007–08 National Sample Survey (NSS) education survey samples households within first-stage units such as villages or urban blocks. This function constructs the joint sampling stratum from state, rural/urban sector, stratum, and substratum, supplies the household multiplier as the sampling weight, and creates the `survey::svydesign()` used for the enrollment probit. `nest = TRUE` allows first-stage-unit identifiers to be interpreted within strata, and final estimation stops if a required design field is missing.

File: R/selection/estimate_selection_probit.R

```
## List NSS fields needed to reproduce the survey sample design
##
selection_survey_design_variables <- function() {
  c("FSU_SL_NO", "weight", "STATE", "SECTOR", "STRATUM", "SUB_STRATUM_NO")
}

## Construct the NSS design used for weighted enrollment estimation
##
build_survey_design_selection <- function(selection_df, require_all = FALSE) {
  required <- selection_survey_design_variables()
  missing <- setdiff(required, names(selection_df))
  if (isTRUE(require_all) && length(missing)) {
    stop(
      "NSS survey design is missing required fields: ",
      paste(missing, collapse = ", "),
      call. = FALSE
    )
  }

  psu <- first_col(selection_df, c("FSU_SL_NO", "fsu", "PSU", "psu"))
  weight <- first_col(selection_df, c("weight", "WEIGHT", "Multiplier", "multiplier"))
  strata_cols <- intersect(c("STATE", "SECTOR", "STRATUM", "SUB_STRATUM_NO"), names(selection_df))
  if (is.null(psu) || is.null(weight) || !length(strata_cols)) {
    if (isTRUE(require_all)) {
      stop("The required NSS survey design could not be constructed.", call. = FALSE)
    }
    return(NULL)
  }

  ## NSS strata are defined by the joint state, rural/urban sector, stratum, and
  ## substratum identifiers used in the survey files.
```

```

selection_df$survey_strata <- interaction(selection_df[strata_cols], drop = TRUE)
survey::svydesign(
  ids = stats::as.formula(paste0("~", psu)),
  strata = ~.survey_strata,
  weights = stats::as.formula(paste0("~", weight)),
  data = selection_df,
  nest = TRUE
)
}

```

Average Marginal Effects on the Estimation Sample

The code recovers the exact observations and NSS sampling weights retained by the enrollment probit, checks that the weights align with those observations, and passes both to `marginaleffects::avg_slopes()`. The fitted survey design determines model estimation and its covariance matrix; the supplied weights separately determine how observation-level response-scale effects are averaged over the fitted sample.

File: R/selection/compute_average_marginal_effects.R

```

#' Recover the observations and sampling weights used to fit the enrollment model
#'
ame_model_data_and_weights <- function(model) {
  model_data <- as.data.frame(stats::model.frame(model))

  model_weights <- if (inherits(model, "svyglm") && !is.null(model$survey.design)) {
    stats::weights(model$survey.design, type = "sampling")
  } else {
    tryCatch(stats::weights(model, type = "prior"), error = function(e) NULL)
  }

  if (is.null(model_weights)) model_weights <- rep(1, nrow(model_data))
  model_weights <- as.numeric(model_weights)
  if (length(model_weights) != nrow(model_data)) {
    stop("AME averaging weights do not align with the estimation sample.", call. = FALSE)
  }
  if (any(!is.finite(model_weights)) || any(model_weights < 0) || sum(model_weights) <= 0) {
    stop("AME averaging weights must be finite, nonnegative, and have positive total weight.", call. = FALSE)
  }

  list(data = model_data, wts = model_weights)
}

```

```

}

#' Compute average marginal effects over the fitted NSS sample
#'
compute_ames_marginaleffects <- function(model) {
  # The survey design already determines model estimation and its covariance
  # matrix. Pass the fitted rows and sampling weights explicitly here so
  # avg_slopes() separately averages observation-level response-scale effects over
  # the same weighted estimation sample after the model has been saved and reloaded.
  amed <- ame_model_data_and_weights(model)
  marginaleffects::avg_slopes(
    model,
    newdata = amed$data,
    wts = amed$wts,
    vcov = TRUE,
    type = "response"
  )
}

```

Output: outputs/tables/appendix/appendix_selection_ame.tex

Rendering code: R/output/appendix_selection_exhibits.R and R/output/save_tables.R

Table 7: Average Marginal Effects for School Enrollment

Term	Enrolled (1 = yes)
Age (years)	-0.028*** (0.000)
Female (ref: Male)	0.044*** (0.003)
Household size	-0.004*** (0.001)
Religion: Muslim (ref: Hindu)	-0.126*** (0.007)
Religion: Christian	-0.043*** (0.013)
Religion: Sikh	0.040 (0.026)
Religion: Jain	-0.001 (0.012)
Religion: Buddhist	0.047* (0.024)
Religion: Zoroastrian	0.101*** (0.034)
Religion: Other	0.252*** (0.003)
Social group: Scheduled Tribe (ref: Other)	-0.102*** (0.008)
Social group: Scheduled Caste	-0.075***

Table 7: Average Marginal Effects for School Enrollment (*continued*)

Term	Enrolled (1 = yes)
	(0.006)
Social group: Other Backward Class	-0.029***
	(0.005)
Urban (ref: Rural)	-0.003
	(0.005)
Distance 1–2km (ref: <1km)	-0.077***
	(0.026)
Distance 2–3km	-0.001
	(0.010)
Distance 3–5km	0.004
	(0.066)
Distance > 5km	0.036
	(0.035)
Male household education: Literate, no school (ref: Illiterate)	0.097***
	(0.017)
Male household education: Literate, school < primary	0.127***
	(0.006)
Male household education: Primary	0.142***
	(0.005)
Male household education: Upper primary	0.195***
	(0.006)
Male household education: Secondary	0.212***
	(0.007)
Male household education: Higher secondary	0.247***
	(0.008)
Male household education: Postsecondary+	0.245***
	(0.009)

* $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Entries are average changes in predicted enrollment probability. Continuous covariates are average slopes; categorical covariates are average discrete comparisons against the stated reference category, evaluated over the observed covariate distribution. Standard errors use the NSS sampling design, including household weights, primary sampling units, and strata formed from state, stratum, and substratum identifiers. The model also controls for district averages of schooling costs and benefits among enrolled children; the table displays marginal effects for child and household covariates. Source: NSS 64th Round, 2007–08, Participation and Expenditure in Education.

Identification and Econometric Robustness

Assessing Weak-Instrument Strength with the Montiel Olea–Pflueger Effective F

The paper uses linguistic distance from Hindi as an instrument for district exposure to English-medium schooling. For specifications with one endogenous regressor, this function evaluates first-stage strength with the Montiel Olea–Pflueger effective F, using a 10% relative-bias benchmark and 5% test size. It recovers the regressors, instruments, estimation rows, and state clusters from the fitted `ivreg` model before calling `momentfit::MOPtest()`. The table pairs this strength diagnostic with two-stage least squares and Anderson–Rubin inference robust to weak identification.

```

mop_effective_f <- function(
  model, analysis_data, tau = 0.10, size = 0.05) {
  unavailable <- function(reason) {
    list(
      statistic = NA_real_, critical_value = NA_real_, p.value = NA_real_,
      effective_df = NA_real_, tau = tau, size = size,
      status = "not_estimated", reason = reason
    )
  }

  if (!inherits(model, "ivreg")) {
    return(unavailable("Montiel Olea-Pflueger effective F requires an ivreg model."))
  }

  if (!requireNamespace("momentfit", quietly = TRUE)) {
    return(unavailable("Package 'momentfit' is not installed."))
  }

  if (!is.finite(tau) || tau <= 0 || !is.finite(size) || size <= 0 || size >= 1) {
    return(unavailable("Montiel Olea-Pflueger tau and size must be valid probabilities."))
  }

  endogenous <- suppressWarnings(as.integer(model$endogenous))
  if (length(endogenous) != 1L || !is.finite(endogenous)) {
    return(unavailable("Montiel Olea-Pflueger effective F requires exactly one endogenous regressor."))
  }

  regression_formula <- tryCatch(
    stats::formula(model, component = "regressors"),
    error = function(e) NULL
  )

  instrument_formula <- tryCatch(
    stats::formula(stats::delete.response(stats::terms(model, component = "instruments")),
    error = function(e) NULL
  )

  if (is.null(regression_formula) || is.null(instrument_formula)) {
    return(unavailable("Could not recover ivreg component formulas for Montiel Olea-Pflueger effective F."))
  }

  needed <- unique(c(all.vars(regression_formula), all.vars(instrument_formula)))
  data <- iv_analysis_frame(analysis_data, needed)
  missing <- setdiff(needed, names(data))

```

```

if (length(missing)) {
  return(unavailable(paste(
    "Missing variables for Montiel Olea-Pflueger effective F:",
    paste(missing, collapse = ", ")
  )))
}
moment_data <- data[needed]
if (any(!stats::complete.cases(moment_data))) {
  return(unavailable("Regressors or instruments used by the fitted IV model contain missing values.))
}

stored_cluster <- attr(model, "cluster_state", exact = TRUE)
cluster <- if (
  length(stored_cluster) == stats::nobs(model) &&
  !anyNA(stored_cluster) && length(unique(stored_cluster)) >= 2L
) {
  as.vector(stored_cluster)
} else {
  iv_model_cluster(model, data)
}
covariance <- if (is.null(cluster)) "MDS" else "CL"
covariance_options <- if (is.null(cluster)) {
  list(type = "HCO")
} else {
  # MOPtest() uses clustered HCO moment covariance with its finite-cluster
  # adjustment here. The separately reported first-stage Wald F uses HC1, so the
  # two statistics follow different covariance conventions.
  list(
    cluster = data.frame(cluster = cluster),
    type = "HCO", cadjust = TRUE, multi0 = FALSE
  )
}

result <- tryCatch({
  moment_model <- momentfit::momentModel(
    regression_formula,
    instrument_formula,
    data = moment_data,
    vcov = covariance,
    vcovOptions = covariance_options
  )

```

```

momentfit::MOPtest(
  moment_model,
  tau = tau,
  size = size,
  estMethod = "TSLS",
  simplified = TRUE,
  print = FALSE
)
}, error = function(e) e)
if (inherits(result, "error")) return(unavailable(conditionMessage(result)))

values <- suppressWarnings(as.numeric(result[c("Feff", "critValue", "pvalue", "Keff"))])
names(values) <- c("statistic", "critical_value", "p.value", "effective_df")
if (length(values) != 4L || any(!is.finite(values))) {
  return(unavailable("momentfit::MOPtest() did not return finite effective-F statistics."))
}

c(
  as.list(values),
  list(tau = tau, size = size, status = "estimated", reason = NA_character_)
)
}

```

Output: outputs/tables/appendix/appendix_iv_weak_inference.tex

Rendering code: R/output/appendix_identification_exhibits.R and R/output/save_tables.R

Table 10: Inference with Weak Instruments for Long-Run Consumption Change

	Real log consumption change, 2004-05 to 2022-23	Real log consumption change, 2004-05 to 2023-24
EMI exposure	0.117 (0.182)	0.137 (0.281)
MOP effective F	0.69	0.41
AR p -value at zero	0.029	0.081
AR 95% accepted set	$[-0.193, -0.030]^{\dagger} \cup [0.009, 0.193]^{\dagger}$	$[-0.186, 0.186]^{\dagger}$
Direct effect threshold for $\beta = 0$ (share of reduced form)	16.6%	0.0%
Observations	524	522

The table reports conventional 2SLS coefficients with state-clustered standard errors. Instrument strength is summarized by the Montiel Olea–Pflueger effective F statistic, and Anderson–Rubin tests provide inference robust to weak instruments. The 95% AR accepted sets are displayed as unions of closed intervals. A dagger marks a component that reaches an evaluated search boundary. Its outer endpoint is the numerical search limit and should not be interpreted as a finite confidence bound. The direct effect threshold is the smallest permitted direct effect, expressed as a share of the absolute reduced form, for which the AR set includes $\beta = 0$. Sources: Census of India 2001; NSS 61st Round, 2004–05; NSS 64th Round, 2007–08; HCES 2022–23; HCES 2023–24.

Constructing Falsification-Adaptive Sets for Overidentified IV Designs

With one endogenous regressor and several excluded instruments, the exclusion-based falsification-adaptive set is formed from just-identified IV estimands that use each excluded instrument in turn while conditioning on the others. The code constructs every constituent on one common complete-case sample, estimates state-clustered inference and its conditional first stage, and forms the set from the minimum and maximum constituent estimates only when every registered constituent is estimable. Weak conditional first stages are reported as cautions but do not cause instruments to be dropped, which would make the identified set depend on a sample-based selection rule.

File: R/iv/falsification_adaptive_set.R

```
iv_fas_component_formula <- function(specification, instrument) {
  spec <- as_single_iv_specification(specification)
  excluded <- plain_chr(unlist(spec$excluded_instruments[[1L]], use.names = FALSE))
  instrument <- plain_chr(instrument)[[1L]]
  if (!instrument %in% excluded) {
    stop("FAS component instrument is not registered as excluded by the IV specification.", call. = FALSE)
  }
  included <- unique(c(
    unlist(spec$included_language_controls[[1L]], use.names = FALSE),
    unlist(spec$controls[[1L]], use.names = FALSE),
```

```

    setdiff(excluded, instrument)
  ))
  make_iv_formula(
    spec$outcome[[1L]], spec$treatment[[1L]], instrument,
    controls = included,
    fixed_effects = iv_fixed_effect_terms(spec$fixed_effect[[1L]])
  )
}

estimate_iv_fas_component <- function(data, specification, instrument) {
  spec <- as_single_iv_specification(specification)
  excluded <- plain_chr(unlist(spec$excluded_instruments[[1L]], use.names = FALSE))
  instrument <- plain_chr(instrument)[[1L]]
  other_instruments <- setdiff(excluded, instrument)

  # The iv_specification_variables() function resolves transformed formula terms such as
  # factor(region) to their underlying data columns. Keep sample construction on
  # that canonical contract instead of treating formula expressions as columns.
  needed <- iv_specification_variables(spec)
  missing <- setdiff(needed, names(data))
  if (length(missing)) {
    stop(
      "FAS specification ", spec$specification_id[[1L]],
      " is missing columns: ", paste(missing, collapse = ", "),
      call. = FALSE
    )
  }
  x <- as.data.frame(data)
  x <- x[stats::complete.cases(x[needed]), , drop = FALSE]
  if (nrow(x) < 3L) {
    return(data.frame(
      specification_id = spec$specification_id[[1L]],
      instrument = instrument, estimate = NA_real_, std.error = NA_real_,
      p.value = NA_real_, first_stage_estimate = NA_real_,
      first_stage_std.error = NA_real_, first_stage_f = NA_real_,
      first_stage_p.value = NA_real_, n = nrow(x), status = "not_estimated",
      reason = "Insufficient complete observations.", stringsAsFactors = FALSE
    ))
  }

  formula <- iv_fas_component_formula(spec, instrument)
  fit <- tryCatch(

```

```

ivreg::ivreg(formula, data = x, model = TRUE, x = TRUE, y = TRUE),
error = function(e) e
)
if (inherits(fit, "error")) {
  return(data.frame(
    specification_id = spec$specification_id[[1L]],
    instrument = instrument, estimate = NA_real_, std.error = NA_real_,
    p.value = NA_real_, first_stage_estimate = NA_real_,
    first_stage_std.error = NA_real_, first_stage_f = NA_real_,
    first_stage_p.value = NA_real_, n = nrow(x), status = "not_estimated",
    reason = conditionMessage(fit), stringsAsFactors = FALSE
  ))
}

cluster <- iv_specification_cluster(x, spec)
inference <- iv_clustered_inference(fit, cluster)
treatment_term <- model_term_inference(
  fit, spec$treatment[[1L]], inference$vcov
)

included <- unique(c(
  unlist(spec$included_language_controls[[1L]], use.names = FALSE),
  unlist(spec$controls[[1L]], use.names = FALSE),
  other_instruments,
  iv_fixed_effect_terms(spec$fixed_effect[[1L]])
))

first_stage <- stats::lm(
  stats::reformulate(unique(c(instrument, included)), response = spec$treatment[[1L]]),
  data = x
)

first_stage_inference <- iv_clustered_inference(first_stage, cluster)
first_stage_term <- model_term_inference(
  first_stage, instrument, first_stage_inference$vcov
)

first_stage_joint <- clustered_joint_wald_test(
  first_stage, instrument, cluster, inference = first_stage_inference
)

first_stage_f <- unname(first_stage_joint[["statistic"]])

status <- if (
  identical(inference$status, "estimated") &&

```

```

    identical(first_stage_inference$status, "estimated") &&
    all(is.finite(treatment_term[c("estimate", "std.error", "p.value")])) &&
    all(is.finite(first_stage_term[c("estimate", "std.error", "p.value")]))
  ) "estimated" else "inference_unavailable"
reasons <- unique(na.omit(c(
  if (identical(inference$status, "unavailable")) inference$reason else NA_character_,
  if (identical(first_stage_inference$status, "unavailable")) first_stage_inference$reason else
    ↪ NA_character_
)))

data.frame(
  specification_id = spec$specification_id[[1L]],
  instrument = instrument,
  estimate = treatment_term[["estimate"]],
  std.error = treatment_term[["std.error"]],
  p.value = treatment_term[["p.value"]],
  first_stage_estimate = first_stage_term[["estimate"]],
  first_stage_std.error = first_stage_term[["std.error"]],
  first_stage_f = first_stage_f,
  first_stage_p.value = unname(first_stage_joint[["p.value"]]),
  n = stats::nobs(fit),
  status = status,
  reason = if (length(reasons)) paste(reasons, collapse = "; ") else NA_character_,
  stringsAsFactors = FALSE
)
}

estimate_iv_falsification_adaptive_set_spec <- function(data, specification) {
  spec <- as_single_iv_specification(specification)
  excluded <- plain_chr(unlist(spec$excluded_instruments[[1L]], use.names = FALSE))
  if (spec$n_endogenous[[1L]] != 1L || length(excluded) <= 1L) {
    stop("FAS requires one endogenous regressor and multiple excluded instruments.", call. = FALSE)
  }
  components <- safe_bind_rows(lapply(excluded, function(instrument) {
    estimate_iv_fas_component(data, spec, instrument)
  }))
  estimated <- components$status == "estimated" & is.finite(components$estimate)
  complete <- length(estimated) == length(excluded) && all(estimated)

  lower <- if (complete) min(components$estimate) else NA_real_
  upper <- if (complete) max(components$estimate) else NA_real_

```

```

lower_instrument <- if (complete) components$instrument[which.min(components$estimate)] else NA_character_
upper_instrument <- if (complete) components$instrument[which.max(components$estimate)] else NA_character_
first_stage_f <- num(components$first_stage_f)

summary <- data.frame(
  specification_id = spec$specification_id[[1L]],
  adjustment_id = spec$adjustment_id[[1L]],
  construction_id = spec$construction_id[[1L]],
  n_instruments = length(excluded),
  n_components_estimated = sum(estimated),
  fas_lower = lower,
  fas_upper = upper,
  fas_width = if (complete) upper - lower else NA_real_,
  fas_contains_zero = if (complete) lower <= 0 && upper >= 0 else NA,
  lower_endpoint_instrument = lower_instrument,
  upper_endpoint_instrument = upper_instrument,
  min_conditional_first_stage_f = if (all(is.finite(first_stage_f))) min(first_stage_f) else NA_real_,
  max_conditional_first_stage_f = if (all(is.finite(first_stage_f))) max(first_stage_f) else NA_real_,
  n_conditional_first_stage_f_below_10 = sum(is.finite(first_stage_f) & first_stage_f < 10),
  constituent_relevance_caution = if (all(is.finite(first_stage_f))) any(first_stage_f < 10) else NA,
  n = if (nrow(components) && length(unique(components$n)) == 1L) components$n[[1L]] else NA_integer_,
  status = if (complete) "estimated" else "incomplete",
  reason = if (complete) NA_character_ else "At least one just-identified constituent could not be
  ↪ estimated with clustered inference.",
  stringsAsFactors = FALSE
)
list(summary = summary, components = components)
}

```

Price Harmonization and Real Consumption

Linking Rural and Urban State CPI Series Across Index Bases

The historical Consumer Price Index for Rural Labourers (CPI-RL) and Consumer Price Index for Industrial Workers (CPI-IW) must be put onto the scale of the later state rural/urban CPI (CPI-R/U). For each state and rural/urban sector, the code uses the median monthly ratio of CPI-R/U to the historical index over January 2013–December 2014 as a multiplicative splice factor, which limits the influence of a single anomalous overlap month. It requires at least six common months and retains the overlap count and dispersion diagnostics alongside each link factor.

```

build_temporal_price_series <- function(
  price_sources,
  switch_date = as.Date("2013-01-01"),
  overlap_start = as.Date("2013-01-01"),
  overlap_end = as.Date("2014-12-01"),
  minimum_link_months = 6L,
  pre_switch_start = NULL,
  pre_switch_end = NULL) {
  switch_date <- price_boundary(switch_date)
  old <- select_pre_2013_price_series(price_sources)
  new <- select_post_2013_price_series(price_sources)
  links <- summarise_price_links(old, new, overlap_start, overlap_end)

  minimum_link_months <- as.integer(minimum_link_months)
  if (length(minimum_link_months) != 1L || is.na(minimum_link_months) || minimum_link_months < 1L) {
    stop("minimum_link_months must be one positive integer.", call. = FALSE)
  }

  old_groups <- unique(old[old$period < switch_date, c("state_code", "sector"), drop = FALSE])
  linked_groups <- links[links$link_months >= minimum_link_months, c("state_code", "sector"), drop = FALSE]
  linked_key <- paste(linked_groups$state_code, linked_groups$sector, sep = "\r")
  missing_key <- !paste(old_groups$state_code, old_groups$sector, sep = "\r") %in% linked_key
  if (any(missing_key)) {
    bad <- old_groups[missing_key, , drop = FALSE]
    stop(
      "Pre-2013 price series have fewer than the required overlap months with CPI-R/U: ",
      paste(paste(bad$state_code, bad$sector, sep = "/"), collapse = ", "),
      call. = FALSE
    )
  }

  old <- merge(old, links, by = c("state_code", "sector"), all.x = TRUE, sort = FALSE)
  old <- old[old$period < switch_date, , drop = FALSE]
  if (!is.null(pre_switch_start) || !is.null(pre_switch_end)) {
    if (is.null(pre_switch_start) || is.null(pre_switch_end)) {
      stop("Both pre_switch_start and pre_switch_end are required when trimming the old price series.",
        call. = FALSE)
    }
    pre_switch_start <- price_boundary(pre_switch_start)
  }

```

```

pre_switch_end <- price_boundary(pre_switch_end)
if (pre_switch_end < pre_switch_start || pre_switch_end >= switch_date) {
  stop("The requested pre-2013 period must end before the CPI switch date.", call. = FALSE)
}
old <- old[old$period >= pre_switch_start & old$period <= pre_switch_end, , drop = FALSE]
}
old$index_unlinked <- old$index
# Put each historical CPI series on the CPI-R/U scale before appending the
# post-2013 observations.
old$index <- num(old$index) * num(old$link_factor)

new <- new[new$period >= switch_date, , drop = FALSE]
new$link_factor <- 1
new$link_months <- NA_integer_
new$link_ratio_mad <- NA_real_
new$link_ratio_min <- NA_real_
new$link_ratio_max <- NA_real_
new$index_unlinked <- new$index

columns <- c(
  "state_code", "sector", "period", "index", "index_unlinked", "price_source",
  "link_factor", "link_months", "link_ratio_mad", "link_ratio_min", "link_ratio_max"
)
index <- rbind(old[columns], new[columns])
index$year <- as.integer(format(index$period, "%Y"))
index$month <- as.integer(format(index$period, "%m"))
index <- index[order(index$state_code, index$sector, index$period), , drop = FALSE]
rownames(index) <- NULL
validate_temporal_price_chain(index, switch_date)

structure(
  list(index = index, links = links, switch_date = switch_date),
  class = "emi_temporal_price_series"
)
}

```

Constructing Comparable State-Sector Price Deflators

Linking the CPI series handles price change over time, while Tendulkar state-sector poverty-line ratios provide a cross-sectional price-level adjustment relative to the all-India rural 2011–12 reference. The displayed core validates those spatial and temporal inputs, constructs or accepts a reference-period CPI for each state-sector

series, preserves documented donor-series lineage when needed, and multiplies the spatial price relative by the temporal CPI relative. The resulting deflator places household consumption on a common spatial and temporal purchasing-power basis; the figure compares district estimates from the 2022–23 and 2023–24 Household Consumption Expenditure Survey rounds.

File: R/prices/price_deflators.R

```
build_state_sector_deflator <- function(
  temporal_index, spatial_relatives, reference_period = NULL, reference_index = NULL) {
  idx <- safe_df(temporal_index)
  spatial <- safe_df(spatial_relatives)
  validate_price_index(idx, keys = c("state_code", "sector", "period"))
  required_spatial <- c("state_code", "sector", "spatial_price_relative")
  missing <- setdiff(required_spatial, names(spatial))
  if (length(missing)) {
    stop("Spatial price relatives are missing columns: ", paste(missing, collapse = ", "), call. = FALSE)
  }
  if (anyDuplicated(spatial[c("state_code", "sector")])) {
    stop("Spatial price relatives must have one row per state and sector.", call. = FALSE)
  }
  if (any(!positive_finite(num(spatial$spatial_price_relative)))) {
    stop("Spatial price relatives must be positive and finite.", call. = FALSE)
  }

  if (is.null(reference_index)) {
    if (is.null(reference_period)) stop("A price reference period or reference-index table is required.",
      ↪ call. = FALSE)
    reference_period <- as.Date(reference_period)
    ref <- idx[idx$period %in% reference_period, , drop = FALSE]
    if (!nrow(ref)) stop("No observations fall in the requested price reference period.", call. = FALSE)
    ref_mean <- stats::aggregate(index ~ state_code + sector, ref, function(x) mean(num(x), na.rm = TRUE))
    names(ref_mean)[names(ref_mean) == "index"] <- "reference_index"
  } else {
    ref_mean <- safe_df(reference_index)
    required_reference <- c("state_code", "sector", "reference_index")
    missing_reference <- setdiff(required_reference, names(ref_mean))
    if (length(missing_reference)) {
      stop("Reference-price table is missing columns: ", paste(missing_reference, collapse = ", "), call. =
        ↪ FALSE)
    }
  }
}
```

```

if (anyDuplicated(ref_mean[c("state_code", "sector")])) ||
  ↪ any(!positive_finite(ref_mean$reference_index))) {
  stop("Reference-price table must have one positive row per state and rural/urban sector.", call. =
    ↪ FALSE)
}

ref_mean <- ref_mean[required_reference]
}

out <- merge(
  idx,
  spatial,
  by = c("state_code", "sector"), all.x = TRUE, sort = FALSE
)

out <- merge(out, ref_mean, by = c("state_code", "sector"), all.x = TRUE, sort = FALSE)
if ("source_state_code" %in% names(out) && any(!positive_finite(out$reference_index))) {
  donor_reference <- ref_mean
  names(donor_reference)[names(donor_reference) == "state_code"] <- "source_state_code"
  names(donor_reference)[names(donor_reference) == "reference_index"] <- "donor_reference_index"
  out <- merge(out, donor_reference, by = c("source_state_code", "sector"), all.x = TRUE, sort = FALSE)
  missing_reference <- !positive_finite(out$reference_index)
  out$reference_index[missing_reference] <- out$donor_reference_index[missing_reference]
  out$donor_reference_index <- NULL
}

# The product adjusts for the rural/urban price level in each state and for
# price change from the reference period to the survey month.
out$temporal_price_relative <- num(out$index) / num(out$reference_index)
out$price_deflator <- num(out$spatial_price_relative) * num(out$temporal_price_relative)
if (any(!positive_finite(out$price_deflator))) {
  stop("Price deflators must be positive and finite for every retained state, sector, and month.", call. =
    ↪ FALSE)
}

out <- out[order(out$state_code, out$sector, out$period), , drop = FALSE]
rownames(out) <- NULL
out
}

```

Output: outputs/figures/appendix/appendix_consumption_hces_consistency.pdf

Rendering code: R/output/appendix_data_construction_exhibits.R and R/output/save_figures.R




Figure 8: Agreement between district real-consumption estimates in the 2022-23 and 2023-24 HCES rounds. Each panel uses districts retained in both rounds under the consumption sample criteria; the dashed line is the 45-degree line. Sources: National Sample Survey Office, HCES 2022-23 and HCES 2023-24 (National Sample Survey Office 2024, 2025).

Testing CPI Links Against Known Answers

A small rural example has monthly new-to-old CPI ratios of 2, 2, and 3, so its registered median link factor must equal 2 even though the final overlap month implies 3; the urban example has a constant ratio of 0.5. The behavioral test verifies those median links, the first and last monthly ratios retained for diagnostics, and the number of overlap months.

File: `tests/testthat/test-prices-and-controls.R`

```
test_that("CPI link factors equal median overlap ratios", {
  months <- as.Date(c("2013-01-01", "2013-02-01", "2013-03-01"))
  old <- rbind(
    data.frame(state_code = "A", sector = "rural", period = months, index = c(100, 110, 120)),
    data.frame(state_code = "A", sector = "urban", period = months, index = c(200, 220, 240))
  )
  new <- rbind(
    data.frame(state_code = "A", sector = "rural", period = months, index = c(200, 220, 360)),
    data.frame(state_code = "A", sector = "urban", period = months, index = c(100, 110, 120))
  )
  links <- summarise_price_links(old, new, months[1], months[3])
  expect_equal(links$link_factor, c(2, 0.5))
  expect_equal(links$link_factor_first, c(2, 0.5))
})
```

```

expect_equal(links$link_factor_last, c(3, 0.5))
expect_equal(links$link_months, c(3L, 3L))
})

```

Spatial Presentation on a Common Geography

Preserving the Complete District Universe in Published Maps

Published district maps start from the complete 2001 Census geometry and left-join analysis attributes onto it rather than plotting only observations retained by a particular estimation sample. Districts without estimates therefore remain in the national outline as no-data areas. The code also reconstructs state identity from the canonical 2001 district key, checks it against any analysis-side state code, and passes the completed geometry to the shared public-map renderer.

File: R/output/save_figures.R

```

complete_public_map_geometry <- function(district_panel, map_geometry) {
  if (!has_sf_geometry(map_geometry)) {
    stop("District maps require the complete 2001 Census district geometry.", call. = FALSE)
  }
  key <- "target_unit_2001"
  if (!key %in% names(map_geometry) && "unit_id" %in% names(map_geometry)) {
    map_geometry[[key]] <- map_geometry$unit_id
  }
  if (!key %in% names(map_geometry) || !key %in% names(district_panel)) {
    stop("Map geometry and analysis data must contain a 2001 Census district key.", call. = FALSE)
  }
  if (anyDuplicated(map_geometry[[key]])) {
    stop("Complete 2001 Census map geometry must contain one row per district.", call. = FALSE)
  }
  attributes <- if (inherits(district_panel, "sf")) sf::st_drop_geometry(district_panel) else
  ↪ safe_df(district_panel)
  attributes <- attributes[!duplicated(attributes[[key]]), , drop = FALSE]
  out <- merge(map_geometry, attributes, by = key, all.x = TRUE, sort = FALSE)

  # Read the state code from each 2001 Census district identifier. This preserves
  # complete state and national borders when a district has no value to plot.
  boundary_state <- public_map_state_code_2001(out[[key]])
  if ("state_code_2001" %in% names(out)) {

```

```

panel_state <- plain_chr(out$state_code_2001)
mismatch <- !is.na(panel_state) & nzchar(panel_state) & panel_state != boundary_state
if (any(mismatch)) {
  stop("Analysis state codes disagree with the state codes in the 2001 Census district identifiers.",
    ↪ call. = FALSE)
}
}
out$state_code_2001 <- boundary_state
out
}

save_map_plot_formats <- function(map_plot, path_base, formats, width = 8, height = 6, dpi = 300) {
  save_plot_formats(map_plot, path_base, formats, width = width, height = height, dpi = dpi)
}

save_map_figure <- function(spec, path_base, district_panel, map_geometry, boundary_reference, formats) {
  if (!has_sf_geometry(district_panel)) {
    stop("Map figure '", spec$name, "' requires an sf district_panel with validated geometry.", call. =
      ↪ FALSE)
  }
  if (is.null(spec$variable) || !spec$variable %in% names(district_panel)) {
    stop("Map figure '", spec$name, "' is missing variable '", spec$variable, "'.", call. = FALSE)
  }

  plot_data <- complete_public_map_geometry(district_panel, map_geometry)
  plot_data <- prepare_public_map_data(plot_data, spec$variable)
  p <- build_public_ggplot_map(plot_data, spec, boundary_reference = boundary_reference)
  save_map_plot_formats(p, path_base, formats, width = 7.2, height = 5.2, dpi = 300)
}

```

Outputs: outputs/figures/main/map_emi_share_enrolled.png, outputs/figures/main/map_emi_exposure.png, outputs/figures/main/map_public_emi_exposure.png, and outputs/figures/main/map_private_emi_exposure.png

Rendering code: R/output/make_figures.R




Figure 3: District distribution of English-medium schooling measures on 2001 Census boundaries. All panels use a common percentage scale. Source: NSS 64th Round, 2007–08 (National Sample Survey Office 2011). Tengshe (2014) and Natural Earth (2022) provide geospatial data. District boundaries reflect 2001 geography. Politically disputed regions are identified using Natural Earth, and those without data—Aksai Chin, Azad Kashmir, Gilgit-Baltistan, the Trans-Karakoram Tract, and Siachen Glacier—are colored light gray. Their inclusion does not imply a position on territorial disputes.