

Inequality in a Potential Equalizer of Opportunity

Variation in the Accessibility and Net Benefits of English-Medium Instruction in India

Anonymous

WRITING SAMPLE: 15-PAGE COPY

This excerpt contains Introduction, Proxies and Persistence of Linguistic Conditions, Linguistic Conditions and Reported Language Choice, Measuring Access to English-Medium Schooling, and EMI Access as a Function of Social Factors from the full paper. Links to the full paper and repository are omitted here. This writing sample can be generated using `make samples` or `bash scripts/run_full_build.sh`.

Abstract

English-medium instruction (EMI), or the teaching of school subjects in English, is often viewed as a potential palliative for the intergenerational rigidity of socioeconomic status in India, where English skills command substantial wage premia. I study whether the conditions governing households' access to and potential net gains from EMI are themselves unequally distributed. I organize my analysis around three such conditions: inherited linguistic costs to acquiring English, access to English-medium schooling, and the local economic setting in which English skills may be rewarded. By harmonizing district data on languages, schooling, consumption, migration, human capital, and sectoral employment, I find inequality across all three. Inherited linguistic conditions are highly persistent and predict future human capital gains, financial asset access, and structural change, but their relationship with language acquisition behavior varies across institutional settings. EMI exposure and provision are strongly organized by schooling system and sharply stratified along socioeconomic lines, even within the same local education market. Consumption growth is more associated with EMI enrollment than enrollment alone, and the strength of their association increases in pre-existing human capital and urbanization. In sum, a high-return educational opportunity need not equalize opportunity when access to it and the resources which help benefit from it are unequally distributed.

1 Introduction

In an Indian economy characterized by rapid GDP growth alongside wage stagnation, profuse informality, and high rates of both youth unemployment and underemployment (International Labour Organization 2024), what can a

household do to improve the future prospects of their children? The history of economics highlights *education* as one answer (Mincer 1974; Card 1999; Björklund and Kjellström 2002; Bhandari and Bordoloi 2006). The history of India highlights another: *English*. In a context of extraordinary linguistic diversity (Gurevich et al. 2025) and thus extraordinary communication and coordination costs (Nakagawa and Sugawara 2021), English has come to serve as a primary domestic lingua franca (Clingsmith 2014). In its capacity as a global lingua franca (Melitz 2016), meanwhile, English opens pathways both internationally—e.g., to education at prestigious universities or employment at multinational firms (Ramanathan 2014; Bedi 2019, 32–33)—and at home. The domestic expansion of internationally traded services and information technology, in particular, has led to a flurry of jobs whose higher socioeconomic returns (Chamarbagwala and Sharma 2011; Smokotin et al. 2014) are largely accessible only to those who can use English (Mankiw and Swagel 2006; Shastri 2012).

Empirical study corroborates the notion that English expands economic mobility (Itani et al. 2015). Using nationally representative household data, Azam et al. (2013) find that English fluency was, at least up until 2005, associated with a 34% hourly wage premium for men and a 22% premium for women; even limited English proficiency was associated with premia of 13% and 10%, respectively. Refeque and Azad (2022) find similar differentials extending through 2011–12. Evidence for causality in this association comes from Chakraborty and Bakshi (2016), who interpret West Bengal’s 1983–2004 ban on English teaching in public primary schools as a natural experiment where cohorts more *exposed* to the policy earned lower subsequent wages. “Language skills are human capital,” after all (Chiswick and Miller 1995, 4), and in a context like India especially, “English-language skills [specifically] are a form of human capital” (Azam et al. 2013, 344).

It is thus by combining our two answers—education and English—that we can begin to understand the demand for *English-medium instruction (EMI)*, or for the teaching of all school subjects in English, currently seen in India (Gooptu and Mukherjee 2023; Lahoti and Mukhopadhyay 2019; Chattopadhyay and Roy 2017). For its potential to expand equality of opportunity, EMI has been touted as a tool of social development (National Knowledge Commission 2009) and described as a route to escape caste (Mathur 2013; Mathew and Srivastava 2011; Vij 2020) alongside the rigid wage, employment, and occupation distributions it enforces across generations (Madheswaran and Singhari 2016; Munshi and Rosenzweig 2006).

But it is ultimately households who decide whether or not a child attends EMI, and the determinants of their choice—the costs they face, the benefits they expect, and the forces constraining the choice set in front of them—are affected by factors which are not just national and historical, but also specific and local. My focus is on whether such factors are associated with the same inequalities that English skills are often expected to help households escape. I study inequalities which exist along dimensions such as social group, religion, parental education, household resources, and urban or rural residence. If the ability to access or benefit from EMI is systematically

related to these characteristics, even within the same local education market, then we have evidence that a high economic payoff to English skills alone is not enough for EMI to become an equalizer of opportunity.

Though I never estimate a structural model of household choice, their investment problem still provides a helpful organizing framework. I will first study a type of cost whose variation is predetermined relative to the institutional and economic factors that otherwise affect the household’s problem: **inherited linguistic costs to learning English**. These will reflect historically persistent features of the local linguistic environment (see Rowe and Weisleder 2020) which affect the opportunity cost of learning English over other domestic bridge languages, namely Hindi. This work, detailed in Section 2, owes much credit to Shastry (2012); in line with her methodology, I create several different measures of local linguistic distance from *Hindi* to capture this opportunity cost. In Section 3 I analyze the **accessibility of English-medium instruction** through enrollment, several measures of EMI exposure, and school management. Private schooling is a distinct household choice with its own costs and selection patterns, so I do not treat private management as a measure of instructional quality (Endow 2019; Muralidharan and Sundararaman 2015). Finally, in Section 5 I study potential determinants of **expected returns to English skills**. My goal here is to find evidence on the value of English skills in the local labor market, with the assumption that households incorporate the local history of such returns into their beliefs regarding English, akin to households having adaptive expectations. I ask whether EMI exposure predicts later district consumption and whether that association is stronger where pre-existing human capital and urbanization are greater. I then examine whether inherited linguistic conditions accompany human-capital growth and skilled in-migration, as well as changes in financial access and economic structure.

The local linguistic environment, the local education market, and the local labor market—all three of the factors I focus on are studied locally. These choices are justifiable: the linguistic conditions which households pass down to their children are highly persistent at the local level, local schooling conditions constrain the menu of educational opportunities households choose from, and local economic conditions plausibly cause households to update their beliefs on the economic returns to learning English. I will use a household’s district to proxy for their ‘local’ environment.¹ Though the district is the most granular unit at which I can harmonize the longitudinal data sources I rely on, it must be noted that interdistrict migration would confound and attenuate any association between the characteristics of one district’s education market and its future economic outcomes. Older studies found interdistrict migration rates (and particularly rates of economically-motivated interdistrict migration) to be so low even after the expansion of IT and services in the 1990s (Topalova 2005; National Statistical Office 2021) that they “assume[d] that district of current residence (or of employment) of an individual is approximately the same as the schooling district” (Chakraborty and Bakshi 2016, 6). More recent work, however, complicates

¹A district in India is a political unit akin to a county in the United States, albeit with an average population of around 2 million. Table 1 contains summary statistics derived from the distribution of district populations.

(Mahapatro 2020; Rajan et al. 2020) and even outright repudiates (Lucas 2021) this notion. These studies emphasize the role of the IT expansion in driving interdistrict migration (Economic Divison 2015; Ghose 2026), and my own results associate potential determinants of household investment into EMI with the skill composition of recent work migrants. I return to those results in Section B. Regardless, it is evident that district outcomes are not an aggregation of individual returns to EMI, but a function of places and their changing populations.

The results show that the conditions surrounding English-skill investment are themselves unequally distributed. The relative cost of acquiring a bridge language is partly inherited, and language behavior varies with these inherited linguistic conditions. Formal English-medium schooling is strongly organized by state boundaries, particularly in the public sector. Within the same district, social-group differences remain large in English-medium and private-school participation, with larger differences in urban schooling markets. EMI exposure is more strongly associated with later local consumption than enrollment alone, and this association is stronger in districts with higher initial human capital and urbanization. The inherited linguistic gradient is also associated with later changes in human capital, the skill composition of recent work migrants, financial access, and sectoral composition. These broader associations are selective rather than evidence of general employment expansion.

Taken together, the evidence limits the equalizing promise of a high-return educational opportunity. A child may face high future payoffs to English skills while the household inherits a less favorable linguistic starting point, has less access to EMI, or lives in a district with fewer complementary resources through which English skills can be converted into economic opportunity. English skills can be valuable while access to English-medium schooling remains unequal. A high-return educational opportunity does not automatically equalize opportunity when entry into it, and the resources needed to benefit from it, are themselves unequally distributed.

2 Inherited Linguistic Conditions

2.1 Proxies and Persistence of Linguistic Conditions

Why measure distance from Hindi? India’s linguistic diversity gives languages of wider communication an important role, and both Hindi and English have served that role across groups (Clingsmith 2014; Shastry 2012). Shastry’s argument is comparative: linguistic diversity creates incentives to learn English or Hindi, and greater distance between a mother tongue and Hindi raises the cost of learning Hindi relative to English (Shastry 2012). The speaker-weighted measure defined in Section 2 summarizes a persistent linguistic component of this relative cost.

Most district variation in the measure occurs between states. In the 573-district schooling panel used later in the paper, state indicators alone explain 84.4% of its cross-district variation. The preferred district specifications

therefore compare districts within states and adjust for the baseline Census-2001 covariates. Figure 1 shows the observed measure and the residual after removing state indicators and the linear contribution of those covariates.




Figure 1: District linguistic distance before and after adjustment for state fixed effects and baseline Census-2001 covariates. Sources: Census of India 2001 language and district characteristics (Office of the Registrar General and Census Commissioner of India 2001; Asher et al. 2021). Tengshe (2014) and Natural Earth (2022) provide geospatial data. District boundaries reflect 2001 geography. Politically disputed regions are identified using Natural Earth, and those for which we lack data—Aksai Chin, Azad Kashmir, Gilgit-Baltistan, the Trans-Karakoram Tract, and Siachen Glacier—are colored light gray. Their inclusion does not imply a position on territorial disputes.

I call these conditions *inherited* because the district pattern is highly persistent across the independently reconstructed 1991 and 2001 measures. Among the 91 Census-2001 districts meeting the historical reconstruction requirements in both years, the correlation between the 1991 and 2001 speaker-weighted distance measures is 0.898; weighting districts by their 1991 populations raises the correlation to 0.928. The population-weighted correlation for the share of speakers whose mother tongues are at least three Shastry degrees from Hindi is 0.919. Figure 2 shows the district comparison. These estimates agree with Shastry’s finding that district linguistic distance is highly correlated across earlier Census years (Shastry 2012).




Figure 2: Persistence of speaker-weighted linguistic distance from Hindi from 1991 to 2001. Each point is a Census-2001 district with a reconstructed measure in both years; the dashed line is the 45-degree line. Sources: Census of India 1991 Language Atlas (Office of the Registrar General and Census Commissioner, India 2004) and Census of India 2001 language tables (Office of the Registrar General and Census Commissioner of India 2001).

An independent reconstruction provides a second validation of the 1991 series. Helms and Lim’s replication data include a Shastry-derived 1991 district linguistic-distance measure (Helms and Lim 2025). Across the 148 districts shared by their extract and my reconstruction, the two measures correlate at 0.989. The agreement shows that the independently assembled series contain nearly identical cross-district variation.

The measure captures inherited linguistic distance from Hindi. Family members, peers, schools, media, and other social settings also shape the language input children receive (Rowe and Weisleder 2020).

2.2 Linguistic Conditions and Reported Language Choice

I next test whether the Shastry relative-cost measure is associated with the subsidiary languages people report. The Census-2001 bilingualism and trilingualism tables record native-language populations and reported subsidiary languages (Office of the Registrar General and Census Commissioner of India 2001). For English and Hindi, the outcome is the share of multilingual native speakers who report the language as a first subsidiary language or, where applicable, a second subsidiary language. The multilingual outcome is the share of native speakers who report any subsidiary language. Each regression observation is a state-by-native-language cell, weighted by its native-speaker population. The specifications compare language groups within states and include the within-state language-prevalence controls defined in Section 2; the registered continuous and threshold specifications also include the Hindi/Urdu reference indicator.

Table 1

Linguistic Distance and Reported Language Choice

	All states				Hindi-belt states		
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
Linguistic distance	2.234 (1.448)		-9.694*** (2.573)	-3.581*** (0.916)	-0.037 (2.942)		-9.024* (5.321)
Distant-language indicator		16.727*** (4.394)				10.165 (10.807)	
Outcome	Reported English	Reported English	Reported Hindi	Multilingualism	Reported English	Reported English	Reported Hindi
Population	Multilingual speakers	Multilingual speakers	Multilingual speakers	Native speakers	Multilingual speakers	Multilingual speakers	Multilingual speakers
Observations	1,566	1,566	1,566	1,587	560	560	560
Partial R^2	0.025	0.128	0.094	0.139	0.000	0.046	0.121
State fixed effects	Yes	Yes	Yes	Yes	Yes	Yes	Yes
Language-prevalence controls	Yes	Yes	Yes	Yes	Yes	Yes	Yes

Each column is a weighted state-by-native-language regression. Standard errors in parentheses use HC1 heteroskedasticity-robust inference. All specifications include state fixed effects, native-language state share, an indicator for the state's most common native language, and a Hindi/Urdu reference indicator. Continuous-distance coefficients are percentage-point changes per one Shastri distance degree; the distant-language indicator equals one for languages at least three degrees from Hindi. Partial R^2 reports the incremental explanatory power of the displayed linguistic-distance term. Columns (5)-(7) use the prespecified Hindi-belt sample. The coefficients describe associations with reported subsidiary-language outcomes. * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$. Source: Census of India 2001 bilingualism and trilingualism tables.

The continuous English estimate is 2.23 percentage points per Shastry distance degree, with a standard error of 1.45 ($p = 0.123$). The threshold specification is much larger: speakers of languages at least three degrees from Hindi report English at a rate 16.73 percentage points higher than speakers of the remaining classified mother tongues (SE 4.39; $p < 0.001$). This specification parallels Shastry’s district measure based on the share of speakers whose languages are at least three degrees from Hindi (Shastry 2012).

The largest continuous coefficient in the table is for reported Hindi. Each additional Shastry degree is associated with 9.69 percentage points less reported Hindi among multilingual speakers (SE 2.57; $p < 0.001$). Greater distance is also associated with a 3.58 percentage-point decline in the share of native speakers who report any subsidiary language (SE 0.92; $p < 0.001$). The pattern therefore includes both the reported subsidiary language and the prevalence of multilingualism.

The alternative linguistic measures produce different English estimates. Glottolog genealogical distance yields a coefficient of 0.21 percentage points (SE 0.43; $p = 0.627$), and Dyen lexical noncognacy yields 0.11 percentage points (SE 0.08; $p = 0.167$) in the corresponding continuous specifications (Hammarström et al. 2026; Dyen et al. 1992; “Comparative Indo-European Database” 1997). The Shastry threshold specification yields the large English association.

3 Access to English-Medium Instruction

3.1 Measuring Access to English-Medium Schooling

Inherited linguistic conditions can change the relative cost of acquiring English, but households can invest only through schooling options they can reach. I therefore use the measures defined in Section 2 to describe realized access to English-medium schooling. The NSS 64th Round records household choices in 2007–08: whether children ages 5–19 are enrolled, the reported medium of instruction among enrolled children, and school management among enrolled children with reported management (National Sample Survey Office 2011). These outcomes reflect available schools, household preferences, and household resources. They are not direct measures of school supply.

The schooling measures separate distinct margins of realized access. Enrollment records whether an age-eligible child attends school. EMI among enrolled children records medium of instruction among children already in school, and all-child EMI exposure incorporates both enrollment and reported medium. I also separate public- and private-school EMI exposure and private-school enrollment. School management and medium of instruction identify different schooling choices: households choosing private schools often cite perceived learning quality or school environment more often than English medium itself (Joshi and Kumar 2017). I therefore keep enrollment,


EMI among enrolled children, all-child EMI exposure, public EMI, private EMI, and private enrollment separate. Reported medium of instruction is also distinct from English proficiency and instructional quality. In the rural sample studied by Lahoti and Mukhopadhyay (2019), 57% of children described as attending English-medium schools were actually taught in the dominant regional language, while another 18% used English books with explanation in the regional language. I therefore interpret the NSS and DISE measures as classifications of reported school medium rather than direct measures of English acquisition or the amount of English used in class. The NSS also does not observe private tutoring or other supplemental English instruction.

DISE provides a school-reported perspective. Its district aggregates describe the composition of the local school system, including enrollment by medium and management, rather than household survey choices. Enrollment shares still combine available provision with family demand, so I use DISE as a measure of the reported school system rather than a pure supply measure. The agreement between NSS and DISE later in this section is informative because the two datasets measure the same schooling margin from different respondents (Mehta 2022; District Information System for Education 2017).


Figure 3 compares the four EMI measures directly. EMI among enrolled children isolates the medium-of-instruction margin among children already in school, while all-child EMI also incorporates enrollment. Splitting all-child exposure by school management shows that public-school EMI is concentrated especially in northern and northeastern districts, whereas private-school EMI is more geographically dispersed. The four panels use the same percentage bins so that differences in both levels and geography can be compared directly.

3.2 EMI Access as a Function of Social Factors


The existence of options within a district does not imply equal access to them. Figure 4 compares Scheduled Tribe (ST), Scheduled Caste (SC), and Other Backward Class (OBC) children with children in the residual Other category within common districts. The social-group differences are present for enrollment but substantially larger for English-medium and private schooling. Relative to Other children, mean district enrollment is lower by about 13.6 percentage points for ST children, 9.8 points for SC children, and 6.4 points for OBC children. EMI among enrolled children is lower by about 13.3, 17.2, and 13.2 points respectively. Private enrollment is lower by about 18.1 points for ST children, 23.5 for SC children, and 14.5 for OBC children. Household resources and social institutions can shape these schooling choices (Endow 2019; Raju 2023). In Mumbai, Munshi and Rosenzweig (2006) show that caste networks continued to direct lower-caste boys toward local-language schools associated with traditional occupations as returns to nontraditional white-collar work rose, while lower-caste girls shifted more rapidly toward English schools.




(a) EMI share among enrolled children.



(b) EMI exposure among all children.



(c) Public-school EMI exposure among all children.



(d) Private-school EMI exposure among all children.

Figure 3: District distribution of English-medium schooling measures on Census-2001 boundaries. All panels use a common percentage scale. Source: NSS 64th Round, 2007-08 (National Sample Survey Office 2011). Tengshe (2014) and Natural Earth (2022) provide geospatial data. District boundaries reflect 2001 geography. Politically disputed regions are identified using Natural Earth, and those for which we lack data—Aksai Chin, Azad Kashmir, Gilgit-Baltistan, the Trans-Karakoram Tract, and Siachen Glacier—are colored light gray. Their inclusion does not imply a position on territorial disputes.




Figure 4: Social-group differences in access to English-medium and private schooling. Panel A reports mean within-district percentage-point differences for OBC, SC, and ST children relative to the Other category in enrollment, EMI among enrolled children, private enrollment, and all-child private EMI exposure. Panel B separates rural and urban differences for EMI among enrolled children and private EMI exposure. Source: NSS 64th Round, 2007–08 (National Sample Survey Office 2011).

The rural/urban comparison makes the social stratification more pronounced. For SC children, for example, the difference in EMI among enrolled children is about 8.4 percentage points in rural districts and 24.6 points in urban districts; the corresponding difference in private EMI exposure is about 7.5 and 21.4 points. OBC and ST children show the same qualitative pattern. A plausible interpretation is that more differentiated urban schooling markets offer a wider range of formal options while household resources and social position still condition which options are reachable. I treat that as a descriptive sorting pattern rather than evidence that urbanization itself causes the larger differences.

The corresponding decompositions for boys and girls are directionally similar. Government programs operating

before the NSS round explicitly targeted girls' education (Press Information Bureau 2008), so I also examine the sex split directly. In this national cross-section, social group and rural/urban setting distinguish access more strongly than sex.

The result is a constraint on how the schooling menu should be interpreted. A district can contain English-medium and private schools without placing those options within equal reach of every household. The next section asks whether these observed EMI choices are associated with later economic outcomes that might help explain household demand for them.