

Coding Sample: Escaping Inequality in India

Rishav Roy

What follows is the 138-page code appendix to my independent paper and writing sample, “Escaping Inequality in India: The Role of English-Medium Instruction.”

Due to the length of this code, the time-constrained reviewer is encouraged to prioritize these 47 pages, excerpted from the code chunks listed below them:

14-16. Scalable **import, cleaning, and merging** of multi-file census and geospatial data.

- Chunk 4: Import key data files

31-32. **QA checks, disambiguation of identifiers** for district-level crosswalk.

- Chunk 6: Match districts: Construct district tracking dfs

50-52. **Testing for systematic missingness** using the Benjamini–Hochberg procedure and parallelized logistic regressions.

- Chunk 8: Participation model: Are NAs randomly distributed

55-57. **Survey-weighted probit estimation**, average marginal effects derivation, and setup for sample selection correction under a complex survey design.

- Chunk 9: Participation model: Estimation and IMR
- Chunk 10: Participation model: Derive AME

80-88. Reusable **record-linkage functions** using cascading fuzzy matches and unmatched/false-match diagnostics to harmonize district identifiers across years and data sources.

- Chunk 16: Match districts: Test joining methods
- Chunk 17: Match districts: Define joining functions

105-107. Reusable **IV specification and estimation pipeline** with clustered inference, plus contiguity-based spatial weights and setup for spatially lagged models.

- Chunk 24: Geospatial: Neighbor list construction
- Chunk 25: 2SLS: Controls and 2SLS formula
- Chunk 26: 2SLS: Baseline 2SLS models

111-112. **Robustness checks and diagnostics** for multicollinearity, spatial autocorrelation.

- Chunk 28: Multicollinearity check
- Chunk 29: Geospatial: Spatial autocorrelation tests

117-138. **Automated generation of publication-quality figures and tables** (including both summary statistics and regression outputs).

- Chunks 31-40: [Output for paper]

Source: R/io/read_inputs.R

2

```

path <- row$absolute_path[[1]]
file_type <- tolower(row$file_type[[1]])
switch(
  file_type,
  sav = read_sav_short(path),
  csv = read_csv_short(path),
  xls = read_excel_short(path),
  xlsx = read_excel_short(path),
  ods = read_ods_short(path),
  shp = {
    need_pkg("sf", "shapefiles")
    sf::st_read(path, quiet = TRUE)
  },
  shx = path,
  dbf = path,
  prj = path,
  png = path,
  stop("No reader implemented for ", file_type, call. = FALSE)
)
}

#' read all manifest rows for a source
#'
#' @return Named list of reader outputs keyed by manifest file_id.
read_manifest_group <- function(paths, source_id) {
  rows <- require_manifest_files(paths, source_id)
  stats::setNames(lapply(seq_len(nrow(rows)), function(i) read_by_manifest_row(rows[i, ])),
    ↪ rows$file_id)
}

```

QA and identifier-disambiguation checks

Source: R/districts/build_district_tracker.R

```

#' build district tracker
#'
build_district_tracker <- function(raw_district_changes) {
  combine_district_tracker_sources(raw_district_changes)
}

#' parse alluvial district changes
#'
parse_alluvial_district_changes <- function(x) {
  x
}

#' parse india district tracker
#'
parse_india_district_tracker <- function(x) {
  x
}

#' parse carveouts renamings
#'
parse_carveouts_renamings <- function(x) {
  x
}

```

```

}

#' parse new districts created
#'
parse_new_districts_created <- function(x) {
  x
}

#' parse name changes
#'
parse_name_changes <- function(x) {
  x
}

#' parse district splits
#'
parse_district_splits <- function(x) {
  x
}

#' combine district tracker sources
#'
combine_district_tracker_sources <- function(raw_district_changes) {
  out <- safe_bind_rows(lapply(names(raw_district_changes), function(name) {
    x <- safe_df(raw_district_changes[[name]])
    x$source_file_id <- name
    x
  })))
  if (nrow(out)) out$.row_in_source <- ave(seq_len(nrow(out)), out$source_file_id, FUN =
    ↪ seq_along)
  out
}

#' standardize tracker years
#'
standardize_tracker_years <- function(tracker) {
  tracker
}

#' standardize tracker names
#'
standardize_tracker_names <- function(tracker) {
  tracker
}

```

Parallelized missingness diagnostics with BH correction

Source: R/selection/diagnose_missingness.R

```

#' diagnose missingness
#'
diagnose_missingness <- function(selection_data, cfg) {
  data.frame(
    missing_var = names(selection_data)[vapply(selection_data, function(x) any(is.na(x)),
    ↪ logical(1))],
    stringsAsFactors = FALSE
  )
}

```

```

)
}

#' check missing logit parallel
#'
check_missing_logit_parallel <- function(df, miss_vars, covars, method_p = "BH") {
  rhs <- paste(covars, collapse = " + ")
  fit_one <- function(m) { f <- stats::as.formula(paste0("is.na(", m, ") ~ ", rhs)); fit <-
  ↪ stats::glm(f, data = df, family = stats::binomial); broom::tidy(fit) |>
  ↪ dplyr::mutate(missing_var = m, pseudoR2 = 1 - fit$deviance / fit$null.deviance, nobs =
  ↪ stats::nobs(fit)) }
  out <- dplyr::bind_rows(lapply(miss_vars, fit_one))
  out |> dplyr::group_by(missing_var) |> dplyr::mutate(p_adj = {adj <- rep(NA_real_,
  ↪ dplyr::n()); idx <- term != "(Intercept)"; adj[idx] <- p.adjust(p.value[idx], method =
  ↪ method_p); adj}) |> dplyr::ungroup()
}

#' summarize missingness by variable
#'
summarize_missingness_by_variable <- function(df) {
  data.frame(
    missing_var = names(df),
    n_missing = vapply(df, function(x) sum(is.na(x)), integer(1)),
    stringsAsFactors = FALSE
  )
}

#' run missingness logits
#'
run_missingness_logits <- function(df, miss_vars, covars) {
  check_missing_logit_parallel(df, miss_vars, covars)
}

#' adjust missingness pvalues bh
#'
adjust_missingness_pvalues_bh <- function(df) {
  df
}

#' save missingness diagnostics
#'
save_missingness_diagnostics <- function(diagnostics) {
  diagnostics
}

```

Survey-weighted probit and IMR setup

Source: R/selection/estimate_selection_probit.R

```

legacy_probit_variables <- function(selection_data) {
  controls <- c("AGE", "age", "SEX", "HH_SIZE", "RELIGION", "SOCIAL_GROUP", "SECTOR")
  exclusion_all_kids <- c("DIST_FROM_NEAREST_PRIMARY_CLASS", "father_educ")
  exclusion_district <- c(
    "dmean_num_IS_EDU_FREE", "dmean_num_TUTION_FEE_WAIVED",
    "dmean_num_RECD_SCHOLARSHIP_STIPEND", "dmean_num_RECD_TXT_BOOKS",
    "dmean_num_RECD_STATIONERY", "dmean_num_MID_DAY_MEAL_ETC_RECD",

```

```

    "dmean_num_ENROLLMENT_COST"
  )
  intersect(c(controls, exclusion_all_kids, exclusion_district), names(selection_data))
}

#' estimate selection probit
#'
estimate_selection_probit <- function(selection_data, cfg) {
  if (!"enrolled" %in% names(selection_data) || all(is.na(selection_data$enrolled))) {
    return(list(status = "out_of_active_pipeline", reason = "No enrolled variable."))
  }
  covars <- legacy_probit_variables(selection_data)
  if (!length(covars)) {
    return(list(status = "out_of_active_pipeline", reason = "No probit covariates."))
  }
  f_probit <- stats::reformulate(covars, response = "enrolled")
  if (identical(cfg$mode, "final") && requireNamespace("survey", quietly = TRUE)) {
    design <- build_survey_design_selection(selection_data)
    if (!is.null(design)) {
      out <- fit_selection_probit(design, f_probit)
      attr(out, "legacy_probit_formula") <- f_probit
      return(out)
    }
  }
  out <- stats::glm(f_probit, data = selection_data, family = stats::binomial(link =
↪ "probit"))
  attr(out, "legacy_probit_formula") <- f_probit
  out
}

#' build survey design selection
#'
build_survey_design_selection <- function(selection_df) {
  psu <- first_col(selection_df, c("FSU_SL_NO", "fsu", "PSU", "psu"))
  weight <- first_col(selection_df, c("weight", "WEIGHT", "Multiplier", "multiplier"))
  strata_cols <- intersect(c("STATE", "STRATUM", "SUB_STRATUM_NO"), names(selection_df))
  if (is.null(psu) || is.null(weight) || !length(strata_cols)) return(NULL)
  options(survey.lonely.psu = "average")
  selection_df$.survey_strata <- interaction(selection_df[strata_cols], drop = TRUE)
  survey::svydesign(
    ids = stats::as.formula(paste0("~", psu)),
    strata = ~.survey_strata,
    weights = stats::as.formula(paste0("~", weight)),
    data = selection_df,
    nest = TRUE
  )
}

#' fit selection probit
#'
fit_selection_probit <- function(selection_design, f_probit) {
  survey::svyglm(f_probit, design = selection_design, family = stats::quasibinomial(link =
↪ "probit"))
}

```

```

#' compute inverse mills ratio
#'
compute_inverse_mills_ratio <- function(model, selection_df, f_probit = NULL) {
  if (is.null(f_probit)) f_probit <- attr(model, "legacy_probit_formula") %||%
    stats::formula(model)
  needed <- all.vars(f_probit)
  needed <- intersect(needed, names(selection_df))
  comp_cases <- stats::complete.cases(selection_df[, needed, drop = FALSE])
  lp <- rep(NA_real_, nrow(selection_df))
  lp[comp_cases] <- stats::predict(model, newdata = selection_df[comp_cases, , drop =
    FALSE], type = "link")
  phi_lp <- stats::dnorm(lp)
  Phi_lp <- pmin(pmax(stats::pnorm(lp), .Machine$double.eps), 1 - .Machine$double.eps)
  enrolled <- as.character(selection_df$enrolled)
  selection_df$IMR <- ifelse(enrolled == "Yes", phi_lp / Phi_lp, ifelse(enrolled == "No",
    phi_lp / (1 - Phi_lp), NA_real_))
  selection_df
}

#' tidy selection model
#'
tidy_selection_model <- function(model) {
  broom::tidy(model)
}

```

Average marginal effects via automatic differentiation

Source: R/selection/compute_average_marginal_effects.R

```

# Use automatic differentiation for SE delta-method gradients, as in the legacy Rmd.

#' compute average marginal effects
#'
compute_average_marginal_effects <- function(selection_model, selection_data, cfg = list())
  {
    if (!inherits(selection_model, "glm")) {
      return(ame_out_of_pipeline(
        selection_model$status %||% "out_of_active_pipeline",
        selection_model$reason %||% "Selection model not estimated in smoke mode"
      ))
    }

    # `selection_model` is often a survey::svyglm object loaded from the targets
    # store. Survey lonely-PSU handling is an R option, not a durable model
    # attribute, so it may not be set in the downstream AME target even though it
    # was set when the model was estimated. Set it explicitly around all AME
    # computations to make reruns deterministic and warning/error-free.
    old_lonely_psu <- getOption("survey.lonely.psu")
    old_domain_lonely <- getOption("survey.adjust.domain.lonely")
    options(survey.lonely.psu = "average", survey.adjust.domain.lonely = TRUE)
    on.exit({
      if (is.null(old_lonely_psu)) {
        options(survey.lonely.psu = NULL)
      } else {
        options(survey.lonely.psu = old_lonely_psu)
      }
    })
  }

```

```

    if (is.null(old_domain_lonely)) {
      options(survey.adjust.domain.lonely = NULL)
    } else {
      options(survey.adjust.domain.lonely = old_domain_lonely)
    }
  }, add = TRUE)

#' run expression while muffling survey lonely-PSU warnings
#'
#' survey emits lonely-PSU warnings even when survey.lonely.psu is set to
#' average and the adjustment is intentional. Muffle only those handled
#' warnings inside the AME target so strict final builds remain warning-clean.
muffle_survey_lonely_psu_warnings <- function(expr) {
  withCallingHandlers(
    expr,
    warning = function(w) {
      msg <- conditionMessage(w)
      lonely_psu_warning <- grepl(
        "has only one PSU at stage",
        msg,
        fixed = TRUE
      )
      if (lonely_psu_warning) invokeRestart("muffleWarning")
    }
  )
}

if (!isTRUE(cfg$run_full_ame)) {
  return(format_ame_results(data.frame(
    term = names(stats::coef(selection_model)),
    estimate = unname(stats::coef(selection_model)),
    std.error = NA_real_,
    statistic = NA_real_,
    p.value = NA_real_,
    s.value = NA_real_,
    conf.low = NA_real_,
    conf.high = NA_real_,
    method = "coefficient_fallback",
    status = "estimated",
    reason = "Draft config uses coefficient fallback instead of full AME computation"
  )))
}

if (!requireNamespace("marginaleffects", quietly = TRUE)) {
  return(ame_out_of_pipeline(
    "out_of_active_pipeline",
    "Package marginaleffects is required for full AME computation"
  ))
}

out <- muffle_survey_lonely_psu_warnings(
  tryCatch(
    compute_ames_autodiff(selection_model, selection_data),
    error = function(e) {

```

```

        fallback <- compute_ames_probit_analytic(selection_model, selection_data)
        fallback$method <- "delta_method_analytic_probit"
        fallback$reason <- NA_character_
        fallback
      }
    )
  )
  format_ame_results(out)
}

ame_out_of_pipeline <- function(status, reason) {
  data.frame(
    term = NA_character_,
    estimate = NA_real_,
    std.error = NA_real_,
    statistic = NA_real_,
    p.value = NA_real_,
    s.value = NA_real_,
    conf.low = NA_real_,
    conf.high = NA_real_,
    method = "not_run",
    status = status,
    reason = reason
  )
}

#' extract model-frame data and AME weights
#'
#' marginalEffects warns, correctly, when weighted/survey models are summarized
#' without an explicit `wts` argument. Build a newdata frame from the exact
#' model estimation sample and attach a synthetic weight column whenever the
#' fitted model exposes usable weights.
ame_model_weight <- function(model_data, model_weights = NULL) {
  weight_cols <- intersect(c("weight", "WEIGHT", "Multiplier", "multiplier", "(weights)"),
    ↪ names(model_data))
  if (length(weight_cols)) {
    return(suppressWarnings(as.numeric(model_data[[weight_cols[[1]]]])))
  }

  if (!is.null(model_weights) && length(model_weights) == nrow(model_data)) {
    return(suppressWarnings(as.numeric(model_weights)))
  }

  NULL
}

#' build marginaeffects newdata and weights
#'
#' @return List with `data`, `wts`, and `has_explicit_weights`.
ame_model_data_and_weights <- function(model) {
  model_data <- as.data.frame(stats::model.frame(model))
  model_weights <- tryCatch(stats::weights(model), error = function(e) NULL)
  weight <- ame_model_weight(model_data, model_weights)

```

```

if (!is.null(weight) && length(weight) == nrow(model_data) && any(is.finite(weight) &
  ↪ weight != 1)) {
  model_data$.ame_weight <- weight
  return(list(data = model_data, wts = ".ame_weight", has_explicit_weights = TRUE))
}

list(data = model_data, wts = FALSE, has_explicit_weights = FALSE)
}

#' run marginaffects while muffling only a redundant explicit-weight warning
#'
run_avg_slopes <- function(model, model_data, wts) {
  withCallingHandlers(
    marginaffects::avg_slopes(
      model,
      newdata = model_data,
      wts = wts,
      vcov = TRUE,
      type = "response"
    ),
    warning = function(w) {
      msg <- conditionMessage(w)
      explicit_weight_warning <- grepl(
        "normally good practice to specify weights using the `wts` argument",
        msg,
        fixed = TRUE
      )
      if (explicit_weight_warning && !identical(wts, FALSE)) {
        invokeRestart("muffleWarning")
      }
    }
  )
}

#' compute ames autodiff
#'
compute_ames_autodiff <- function(model, newdata) {
  # Use the exact estimation sample retained by glm/svyglm. Passing the full
  # pre-model selection_data can have extra rows omitted during model fitting,
  # which makes marginaffects recycle weights/covariates silently.
  amed <- ame_model_data_and_weights(model)
  run_avg_slopes(model, amed$data, amed$wts)
}

#' compute ames fast draft
#'
compute_ames_fast_draft <- function(model, newdata, n = 200) {
  amed <- ame_model_data_and_weights(model)
  run_avg_slopes(
    model,
    dplyr::slice_sample(amed$data, n = min(n, nrow(amed$data))),
    amed$wts
  )
}

```

```

#' compute analytic probit AMEs
#'
#' @return Data frame of observed-data average marginal effects.
compute_ames_probit_analytic <- function(model, newdata) {
  coefs <- stats::coef(model)
  coef_table <- as.data.frame(summary(model)$coefficients)
  terms <- setdiff(names(coefs), "(Intercept)")
  if (!length(terms)) return(ame_out_of_pipeline("out_of_active_pipeline", "Selection model
  ↪ has no non-intercept terms."))

  newdata <- stats::model.frame(model)
  weights <- if ("weight" %in% names(newdata)) num(newdata$weight) else rep(1,
  ↪ nrow(newdata))
  eta <- stats::predict(model, type = "link")
  mean_phi <- wmean(stats::dnorm(eta), weights)
  factor_levels <- model$xlevels %||% list()

  rows <- lapply(terms, function(term) {
    factor_var <- names(factor_levels)[vapply(names(factor_levels), function(v)
  ↪ startsWith(term, v), logical(1))]
    estimate <- NA_real_
    if (length(factor_var)) {
      var <- factor_var[[which.max(nchar(factor_var))]]
      level <- sub(paste0("^", var), "", term)
      if (level %in% factor_levels[[var]]) {
        hi <- newdata
        lo <- newdata
        hi[[var]] <- factor(level, levels = factor_levels[[var]])
        lo[[var]] <- factor(factor_levels[[var]][[1]], levels = factor_levels[[var]])
        estimate <- wmean(stats::predict(model, newdata = hi, type = "response") -
  ↪ stats::predict(model, newdata = lo, type = "response"), weights)
      }
    } else {
      estimate <- unname(coefs[[term]]) * mean_phi
    }
    p_col <- intersect(c("Pr(>|z|)", "Pr(>|t|)"), names(coef_table))
    p <- if (term %in% rownames(coef_table) && length(p_col)) coef_table[term, p_col[[1]]]
  ↪ else NA_real_
    se_col <- intersect(c("Std. Error", "std.error"), names(coef_table))
    se_beta <- if (term %in% rownames(coef_table) && length(se_col)) coef_table[term,
  ↪ se_col[[1]]] else NA_real_
    se <- if (is.finite(se_beta)) abs(se_beta * mean_phi) else NA_real_
    z <- if (is.finite(se) && se > 0) estimate / se else NA_real_
    p_out <- if (is.finite(z)) 2 * stats::pnorm(abs(z), lower.tail = FALSE) else p
    data.frame(
      term = term,
      estimate = estimate,
      std.error = se,
      statistic = z,
      p.value = p_out,
      s.value = if (is.finite(p_out) && p_out > 0) -log2(p_out) else NA_real_,
      conf.low = if (is.finite(se)) estimate - 1.96 * se else NA_real_,
      conf.high = if (is.finite(se)) estimate + 1.96 * se else NA_real_,
      method = "delta_method_analytic_probit",
      status = "estimated",

```

```

      reason = NA_character_,
      stringsAsFactors = FALSE
    )
  })
  do.call(rbind, rows)
}

copy_first_ame_column <- function(out, target, candidates) {
  if (!target %in% names(out)) out[[target]] <- NA
  for (candidate in candidates) {
    if (!candidate %in% names(out)) next
    missing_target <- is.na(out[[target]])
    if (any(missing_target)) out[[target]][missing_target] <-
      ↪ out[[candidate]][missing_target]
  }
  out
}

normalize_ame_columns <- function(out) {
  # marginalesffects has used both dotted and snake_case names across versions
  # and model classes. Normalize here before selecting the public/audit schema
  # so uncertainty columns are not silently dropped downstream.
  aliases <- list(
    std.error = c("std_error", "std.err", "Std. Error"),
    statistic = c("z", "z.value", "z_value", "t", "t.value", "t_value"),
    p.value = c("p_value", "p", "Pr(>|z|)", "Pr(>|t|)"),
    s.value = c("s_value"),
    conf.low = c("conf_low", "conf.low", "2.5 %"),
    conf.high = c("conf_high", "conf.high", "97.5 %")
  )
  for (target in names(aliases)) {
    out <- copy_first_ame_column(out, target, aliases[[target]])
  }
  out
}

legacy_ame_lookup <- function() {
  data.frame(
    term = c(
      "AGE", "SEX", "HH_SIZE",
      "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION",
      "SOCIAL_GROUP", "SOCIAL_GROUP", "SOCIAL_GROUP",
      "SECTOR",
      "DIST_FROM_NEAREST_PRIMARY_CLASS", "DIST_FROM_NEAREST_PRIMARY_CLASS",
      ↪ "DIST_FROM_NEAREST_PRIMARY_CLASS", "DIST_FROM_NEAREST_PRIMARY_CLASS",
      "father_educ", "father_educ", "father_educ", "father_educ", "father_educ",
      ↪ "father_educ", "father_educ",
      "dmean_num_IS_EDU_FREE", "dmean_num_TUTION_FEE_WAIVED",
      ↪ "dmean_num_RECD_SCHOLARSHIP_STIPEND",
      "dmean_num_RECD_TXT_BOOKS", "dmean_num_RECD_STATIONERY",
      ↪ "dmean_num_MID_DAY_MEAL_ETC_RECD", "dmean_num_ENROLLMENT_COST"
    ),
    contrast = c(
      "dY/dX", "Female - Male", "dY/dX",

```

```

"Muslim - Hindu", "Christian - Hindu", "Sikh - Hindu", "Jain - Hindu", "Buddhist -
  ↳ Hindu", "Zoroastrian - Hindu", "Other - Hindu",
"Scheduled Tribe - Other", "Scheduled Caste - Other", "Other Backward Class - Other",
"Urban - Rural",
"1km <= d <2kms - d<1km", "2kms<= d <3kms - d<1km", "3kms <= d <5kms - d<1km",
  ↳ "d>=5kms - d<1km",
"Literate, no school - Illiterate", "Literate, school < primary - Illiterate",
  ↳ "Primary - Illiterate", "Upper primary - Illiterate", "Secondary - Illiterate",
  ↳ "Higher secondary - Illiterate", "Postsecondary+ - Illiterate",
"dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX"
),
Term = c(
  "Age (years)", "Female (ref: Male)", "Household size",
  "Religion: Muslim (ref: Hindu)", "Religion: Christian", "Religion: Sikh", "Religion:
    ↳ Jain", "Religion: Buddhist", "Religion: Zoroastrian", "Religion: Other",
  "Social group: Scheduled Tribe (ref: Other)", "Social group: Scheduled Caste", "Social
    ↳ group: Other Backward Class",
  "Urban (ref: Rural)",
  "Distance 1-2km (ref: <1km)", "Distance 2-3km", "Distance 3-5km", "Distance > 5km",
  "Father's educ.: Literate, no school (ref: Illiterate)", "Father's educ.: Literate,
    ↳ school < primary", "Father's educ.: Primary", "Father's educ.: Upper primary",
    ↳ "Father's educ.: Secondary", "Father's educ.: Higher secondary", "Father's educ.:
    ↳ Postsecondary+",
  "Educ. free available (ref: No)", "Tuition waiver received", "Scholarship/Stipend
    ↳ received", "Textbook(s) received", "Stationery received", "Mid-day meal, etc.
    ↳ received", "Enrollment cost (Rs.)"
),
stringsAsFactors = FALSE
)
}

legacy_ame_label_order <- function() legacy_ame_lookup()$Term

infer_ame_contrast <- function(out) {
  if ("contrast" %in% names(out)) return(as.character(out$contrast))
  contrast <- rep("dY/dX", nrow(out))
  if (!"term" %in% names(out)) return(contrast)
  if ("comparison" %in% names(out)) contrast <- as.character(out$comparison)
  contrast[is.na(contrast) | !nzchar(contrast)] <- "dY/dX"
  contrast
}

ame_factor_base_terms <- function() {
  unique(legacy_ame_lookup()$term[legacy_ame_lookup()$contrast != "dY/dX"])
}

normalize_encoded_factor_ame_terms <- function(out) {
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  if (!"term" %in% names(out)) return(out)
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(out)

  lookup <- legacy_ame_lookup()
  factor_terms <- ame_factor_base_terms()
  factor_terms <- factor_terms[order(nchar(factor_terms), decreasing = TRUE)]

```

```

for (i in seq_len(nrow(out))) {
  current_term <- as.character(out$term[[i]])
  current_contrast <- as.character(out$contrast[[i]])
  if (!nzchar(current_term) || (!is.na(current_contrast) && current_contrast != "dY/dX"))
    ↪ next

  direct <- lookup$term == current_term & lookup$contrast == current_contrast
  if (any(direct)) next

  for (base_term in factor_terms) {
    if (!startsWith(current_term, base_term)) next
    level <- trimws(sub(paste0("^", base_term), "", current_term))
    if (!nzchar(level)) next
    candidates <- lookup[lookup$term == base_term, , drop = FALSE]
    contrast_hit <- candidates$contrast[startsWith(candidates$contrast, paste0(level, " -
    ↪ "))])
    if (length(contrast_hit)) {
      out$term[[i]] <- base_term
      out$contrast[[i]] <- contrast_hit[[1]]
      break
    }
  }
}
out
}

```

```

attach_legacy_ame_labels <- function(out) {
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  if (!"term" %in% names(out) && "variable" %in% names(out)) out$term <- out$variable
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(out)
  out <- normalize_encoded_factor_ame_terms(out)
  lookup <- legacy_ame_lookup()
  labeled <- merge(out, lookup, by = c("term", "contrast"), all.x = TRUE, sort = FALSE)
  labeled$Term[is.na(labeled$Term)] <- labeled$term[is.na(labeled$Term)]
  labeled$Term <- factor(labeled$Term, levels = unique(c(legacy_ame_label_order(),
  ↪ labeled$Term)))
  labeled <- labeled[order(labeled$Term), , drop = FALSE]
  labeled$Term <- as.character(labeled$Term)
  labeled
}

```

#' format_ame_results

#'

```

format_ame_results <- function(ame_results) {
  out <- tibble::as_tibble(ame_results)
  if (!"term" %in% names(out) && "variable" %in% names(out)) out$term <- out$variable
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(as.data.frame(out))
  out <- normalize_ame_columns(out)
  if (!"method" %in% names(out)) out$method <- "autodiff"
  if (!"status" %in% names(out)) out$status <- "estimated"
  if (!"reason" %in% names(out)) out$reason <- NA_character_
  legacy_terms <- unique(legacy_ame_lookup()$term)
  use_legacy_schema <- any(out$term %in% legacy_terms) ||
    any(grepl("^dmean_num_", out$term %||% character(), perl = TRUE)) ||
    any((out$contrast %||% "dY/dX") != "dY/dX", na.rm = TRUE)
}

```

```

if (isTRUE(use_legacy_schema)) {
  out <- attach_legacy_ame_labels(out)
  required <- c(
    "Term", "term", "contrast", "estimate", "std.error", "statistic", "p.value",
    ↪ "s.value",
    "conf.low", "conf.high", "method", "status", "reason"
  )
} else {
  required <- c(
    "term", "estimate", "std.error", "statistic", "p.value",
    "s.value", "conf.low", "conf.high", "method", "status", "reason"
  )
}

for (nm in setdiff(required, names(out))) out[[nm]] <- NA
if ("p.value" %in% names(out) && "s.value" %in% names(out)) {
  missing_s <- is.na(out$s.value) & is.finite(out$p.value) & out$p.value > 0
  out$s.value[missing_s] <- -log2(out$p.value[missing_s])
}
out[, required, drop = FALSE]
}

#' save ame results
#'
save_ame_results <- function(ame_results, path =
  ↪ "outputs/tables/diagnostics/ame_results.csv") {
  readr::write_csv(tibble::as_tibble(ame_results), path); path
}

```

Cascading fuzzy district record linkage

Source: R/districts/fuzzy_join_districts.R

*# The functions below preserve the cascading fuzzy-match architecture from the legacy Rmd.
 # sample-end marker appears near the end of this file after the functions are migrated.*

```

#' evaluate distances
#'
evaluate_distances <- function(pairs, methods, thresholds, col1 = "str1", col2 = "str2") {
  if (length(methods) != length(thresholds)) stop("\"methods\" and \"thresholds\" must have
  ↪ the same length.")
  safe_bind_rows(lapply(seq_along(methods), function(i) {
    distance <- utils::adist(canon(pairs[[col1]]), canon(pairs[[col2]]), ignore.case =
    ↪ TRUE)[, 1]
    data.frame(
      str1 = pairs[[col1]],
      str2 = pairs[[col2]],
      method = methods[i],
      distance = distance,
      threshold = thresholds[i],
      match = distance <= thresholds[i],
      stringsAsFactors = FALSE
    )
  })))
}

```

```

#' fuzzy join sequence
#'
fuzzy_join_sequence <- function(df1, df2, dist1, state1, dist2, state2, methods, thresholds,
  ↪ mode = "full") {
  df1_id <- df1 |> dplyr::mutate(.id1 = dplyr::row_number())
  df2_id <- df2 |> dplyr::mutate(.id2 = dplyr::row_number())
  df1_curr <- df1_id; df2_curr <- df2_id; matched_all <- tibble::tibble()
  for (i in seq_along(methods)) {
    full_j <- fuzzyjoin::stringdist_join(df1_curr, df2_curr, by = stats::setNames(c(dist2,
  ↪ state2), c(dist1, state1)), mode = mode, method = methods[i], max_dist = thresholds[i],
  ↪ distance_col = "dist")
    matched_i <- full_j |> dplyr::filter(!is.na(.id1), !is.na(.id2)) |>
  ↪ dplyr::group_by(.id1) |> dplyr::slice_min(dist, n = 1, with_ties = FALSE) |>
  ↪ dplyr::ungroup() |> dplyr::group_by(.id2) |> dplyr::slice_min(dist, n = 1, with_ties =
  ↪ FALSE) |> dplyr::ungroup()
    matched_all <- dplyr::bind_rows(matched_all, matched_i)
    df1_curr <- df1_curr |> dplyr::filter(!(.id1 %in% matched_i$.id1)); df2_curr <- df2_curr
  ↪ |> dplyr::filter(!(.id2 %in% matched_i$.id2))
  }
  list(joined = matched_all, unmatched_df1 = df1_curr |> dplyr::select(-.id1), unmatched_df2
  ↪ = df2_curr |> dplyr::select(-.id2))
}

#' unmatched rows
#'
#' @return A data frame of unmatched rows when the join map carries that attribute.
unmatched_rows <- function(join_map, ...) {
  attr(join_map, "unmatched_rows") %||% join_map[0, , drop = FALSE]
}

#' merge dfs into tracker
#'
#' Legacy-compatible cascading district matcher. `df_names` is a named list (or a
#' data frame) whose elements carry source-specific district/state names. For
#' each source we try every requested tracker year in order, keep one-to-one
#' canonical name matches, and then fuzzy-match remaining rows within state.
#'
#' @return A list with `joined_df`, `unmatched_df`, and `flagged_df`, matching
#' the object contract used by the legacy Rmd.
merge_dfs_into_tracker <- function(df_names, tracker = NULL, years_of_interest = c("2001",
  ↪ "2007", "2008", "2017", "2018", "2020"), flag = TRUE, ...) {
  if (is.null(tracker)) return(list(joined_df = safe_df(df_names), unmatched_df =
  ↪ data.frame(), flagged_df = data.frame()))
  tracker <- safe_df(tracker)
  if (!nrow(tracker)) return(list(joined_df = data.frame(), unmatched_df =
  ↪ safe_df(df_names), flagged_df = data.frame()))
  if (!".tracker_row" %in% names(tracker)) tracker$.tracker_row <- seq_len(nrow(tracker))
  xs <- if (inherits(df_names, "data.frame")) list(source = df_names) else
  ↪ as_input_list(df_names)

  joined_all <- list()
  unmatched_all <- list()
  flagged_all <- list()
  for (source_name in names(xs)) {

```

```

source <- safe_df(xs[[source_name]])
if (!nrow(source)) next
source$.source_row <- seq_len(nrow(source))
source$.source_name <- source_name
source <- add_source_name_keys(source)
if (!all(c(".source_state_key", ".source_district_key") %in% names(source))) {
  unmatched_all[[source_name]] <- source
  next
}

source$.matched <- FALSE
for (yr in years_of_interest) {
  sfx <- substr(as.character(yr), 3, 4)
  state_col <- paste0("state_", sfx)
  district_col <- paste0("district_", sfx)
  if (!all(c(state_col, district_col) %in% names(tracker))) next
  candidate <- tracker[c(".tracker_row", state_col, district_col)]
  candidate$.tracker_state_key <- canon(candidate[[state_col]])
  candidate$.tracker_district_key <- canon(candidate[[district_col]])
  exact <- merge(
    source[!source$.matched, , drop = FALSE],
    candidate,
    by.x = c(".source_state_key", ".source_district_key"),
    by.y = c(".tracker_state_key", ".tracker_district_key"),
    all.x = FALSE,
    all.y = FALSE
  )
  if (!nrow(exact)) next
  exact <- exact[!duplicated(exact$.source_row) & !duplicated(exact$.tracker_row), ,
    drop = FALSE]
  if (!nrow(exact)) next
  exact$match_year <- as.character(yr)
  exact$match_status <- "exact_name"
  joined_all[[paste(source_name, yr, sep = "_")] <- exact
  source$.matched[source$.source_row %in% exact$.source_row] <- TRUE
}

if (any(!source$.matched)) {
  fuzzy <- fuzzy_match_remaining_to_tracker(source[!source$.matched, , drop = FALSE],
    tracker, years_of_interest)
  if (nrow(fuzzy)) {
    joined_all[[paste0(source_name, "_fuzzy")] <- fuzzy
    source$.matched[source$.source_row %in% fuzzy$.source_row] <- TRUE
    if (flag) flagged_all[[source_name]] <- fuzzy[fuzzy$possible_false_positive %in%
      TRUE, , drop = FALSE]
  }
}
unmatched_all[[source_name]] <- source[!source$.matched, , drop = FALSE]
}
list(
  joined_df = safe_bind_rows(joined_all),
  unmatched_df = safe_bind_rows(unmatched_all),
  flagged_df = safe_bind_rows(flagged_all)
)
}

```

```

add_source_name_keys <- function(source) {
  state <- first_col(source, c("state", "state_01", "state_07", "state_08", "state_17",
↪ "state_18", "state_20", "state_0708", "state_1718"))
  district <- first_col(source, c("district", "district_01", "district_07", "district_08",
↪ "district_17", "district_18", "district_20", "district_0708", "district_1718",
↪ "district_name"))
  if (!is.null(state)) source$.source_state_key <- canon(source[[state]])
  if (!is.null(district)) source$.source_district_key <- canon(source[[district]])
  source
}

fuzzy_match_remaining_to_tracker <- function(source, tracker, years_of_interest) {
  out <- list()
  for (yr in years_of_interest) {
    sfx <- substr(as.character(yr), 3, 4)
    state_col <- paste0("state_", sfx)
    district_col <- paste0("district_", sfx)
    if (!all(c(state_col, district_col) %in% names(tracker))) next
    candidate <- tracker[c(".tracker_row", state_col, district_col)]
    candidate$.tracker_state_key <- canon(candidate[[state_col]])
    candidate$.tracker_district_key <- canon(candidate[[district_col]])
    for (i in seq_len(nrow(source))) {
      pool <- candidate[candidate$.tracker_state_key == source$.source_state_key[[i]], ,
↪ drop = FALSE]
      if (!nrow(pool)) next
      dist <- utils::adist(source$.source_district_key[[i]], pool$.tracker_district_key,
↪ ignore.case = TRUE)[1, ]
      best <- which.min(dist)
      if (!length(best) || !is.finite(dist[[best]]) || dist[[best]] > 2) next
      row <- cbind(source[i, , drop = FALSE], pool[best, , drop = FALSE])
      row$match_year <- as.character(yr)
      row$match_status <- "fuzzy_name"
      row$dist <- dist[[best]]
      row$possible_false_positive <- dist[[best]] > 0
      out[[length(out) + 1L]] <- row
    }
  }
  x <- safe_bind_rows(out)
  if (nrow(x)) x <- x[!duplicated(x$.source_row) & !duplicated(x$.tracker_row), , drop =
↪ FALSE]
  x
}

#' join year to tracker
#'
join_year_to_tracker <- function(source, tracker, years_of_interest, ...) {
  merge_dfs_into_tracker(source, tracker = tracker, years_of_interest = years_of_interest,
↪ ...)
}

#' join all district sources
#'
join_all_district_sources <- function(df_names, tracker, years_of_interest, ...) {

```

```

merge_dfs_into_tracker(df_names, tracker = tracker, years_of_interest = years_of_interest,
  ↪ ...)
}

#' flag many to many matches
#'
flag_many_to_many_matches <- function(join_map) {
  join_map$many_to_many <- duplicated(join_map$.source_row) |
  ↪ duplicated(join_map$.tracker_row)
  join_map
}

#' score candidate matches
#'
score_candidate_matches <- function(candidates) {
  if (!"dist" %in% names(candidates)) candidates$dist <- 0
  candidates$score <- -num(candidates$dist)
  candidates
}

#' fuzzy join districts
#'
fuzzy_join_districts <- function(district_tracker, district_keys_2001, district_keys_2007,
  ↪ district_keys_2017, district_keys_2020, cfg) {
  tracker <- safe_df(district_tracker)
  if (nrow(tracker) && all(c("state_01", "district_01", "state_07", "district_07",
  ↪ "state_17", "district_17", "state_20", "district_20") %in% names(tracker))) {
    tracker$.tracker_row <- seq_len(nrow(tracker))
    tracker$source <- "legacy_tracker"
    tracker$match_status <- "legacy_tracker_row"
    tracker$possible_false_positive <- FALSE
    tracker$many_to_many <- FALSE
    attr(tracker, "unmatched_rows") <- tracker[0, , drop = FALSE]
    attr(tracker, "possible_false_positives") <- tracker[0, , drop = FALSE]
    attr(tracker, "many_to_many_cases") <- tracker[0, , drop = FALSE]
    return(tracker)
  }

  out <- safe_bind_rows(Map(function(keys, source) {
    if (!nrow(keys)) return(data.frame())
    keys$source <- source
    keys$match_status <- "source_key_unmatched"
    keys$possible_false_positive <- FALSE
    keys$many_to_many <- FALSE
    keys
  }, list(district_keys_2001, district_keys_2007, district_keys_2017, district_keys_2020),
  ↪ c("2001", "2007", "2017", "2020")))
  attr(out, "unmatched_rows") <- out
  attr(out, "possible_false_positives") <- out[out$possible_false_positive, , drop = FALSE]
  attr(out, "many_to_many_cases") <- out[out$many_to_many, , drop = FALSE]
  out
}

```

Contiguity-based spatial weights

Source: R/diagnostics/diagnose_spatial_weights.R

```
#' build spatial weights
#'  
#' @return A spatial weights object, or an explicit inactive status when geometry is  
↪ unavailable.  
build_spatial_weights <- function(district_panel, cfg) {  
  if (!inherits(district_panel, "sf")) {  
    return(list(status = "out_of_active_pipeline", reason = "Requires sf geometry."))  
  }  
  geom <- sf::st_geometry(district_panel)  
  coverage <- mean(!sf::st_is_empty(geom))  
  if (!is.finite(coverage) || coverage < 0.75) {  
    return(list(  
      status = "out_of_active_pipeline",  
      reason = paste0(  
        "Requires validated non-empty geometry for at least 75% of district-panel rows;  
        ↪ current coverage is ",  
        round(100 * coverage, 1), "%." )  
    ))  
  }  
}  
spatial_warnings <- character()  
capture_expected_spatial_warning <- function(expr) {  
  withCallingHandlers(  
    expr,  
    warning = function(w) {  
      msg <- conditionMessage(w)  
      expected <- grepl("no neighbours|sub-graphs", msg, ignore.case = TRUE)  
      if (expected) {  
        spatial_warnings <-<- unique(c(spatial_warnings, msg))  
        invokeRestart("muffleWarning")  
      }  
    }  
  )  
}  
nb <- capture_expected_spatial_warning(  
  spdep::poly2nb(district_panel[!sf::st_is_empty(geom), , drop = FALSE], queen = FALSE)  
)  
out <- capture_expected_spatial_warning(  
  spdep::nb2listw(nb, style = "W", zero.policy = TRUE)  
)  
attr(out, "spatial_warnings") <- spatial_warnings  
out  
}  
  
#'  
#'  
diagnose_spatial_weights <- function(district_panel, spatial_weights, cfg) {  
  data.frame(  
    diagnostic = "spatial_weights",  
    status = spatial_weights$status %||% "constructed",  
    warnings = paste(attr(spatial_weights, "spatial_warnings") %||% character(), collapse =  
      ↪ "; "),  
  )  
}
```

```

    stringsAsFactors = FALSE
  )
}

#' compare rook queen contiguity
#'
compare_rook_queen_contiguity <- function(district_panel) {
  tibble::tibble()
}

#' summarize islands
#'
summarize_islands <- function(spatial_weights) {
  tibble::tibble()
}

#' summarize neighbor counts
#'
summarize_neighbor_counts <- function(spatial_weights) {
  tibble::tibble()
}

#' save spatial weight diagnostics
#'
save_spatial_weight_diagnostics <- function(diagnostics) {
  diagnostics
}

```

Reusable IV specification and estimation

Source: R/iv/build_iv_formulas.R

```

#' make iv formula
#'
make_iv_formula <- function(dep, endog, instruments, controls = NULL, fixed_effects = NULL)
  ↪ {
  stats::as.formula(paste(
    dep,
    "~",
    paste(c(endog, controls, fixed_effects), collapse = " + "),
    "|",
    paste(c(instruments, controls, fixed_effects), collapse = " + ")
  ))
}

legacy_iv_controls <- function() {
  c(
    "consumption_0708", "gini_cons_0708",
    "pct_urban", "avg_hh_size", "dependency_ratio",
    "pct_fem_head",
    "pct_hindu", "pct_muslim",
    "pct_st", "pct_sc", "pct_obc",
    "pct_small_land", "pct_medium_land", "pct_large_land",
    "pct_head_lit_to_primary", "pct_head_secondary_plus"
  )
}

```

```

#' build iv formulas
#'
build_iv_formulas <- function(cfg) {
  controls <- legacy_iv_controls()
  list(
    consumption = make_iv_formula(
      "consumption_pct_change",
      "EMIE",
      "wavg_ling_degrees",
      controls
    ),
    gini = make_iv_formula(
      "gini_change",
      "EMIE",
      "wavg_ling_degrees",
      controls
    )
  )
}

#' build baseline 2sls formula
#'
build_baseline_2sls_formula <- function(...) make_iv_formula(...)

#' build fd 2sls formula
#'
build_fd_2sls_formula <- function(...) {
  list(status = "out_of_active_pipeline", reason = "First-difference 2SLS formula is
    ↪ deferred until the FD redesign is specified.")
}

#' build state fe 2sls formula
#'
build_state_fe_2sls_formula <- function(...) {
  list(status = "out_of_active_pipeline", reason = "State fixed-effect 2SLS formula is
    ↪ deferred until the state-FE redesign is specified.")
}

```

Multicollinearity and spatial autocorrelation diagnostics

Source: R/diagnostics/diagnose_multicollinearity.R

```

#' diagnose multicollinearity
#'
#' @return A data frame with design-matrix condition diagnostics.
diagnose_multicollinearity <- function(district_panel, iv_models, cfg) {
  model <- if (is.list(iv_models) && length(iv_models)) iv_models[[1]] else iv_models
  out <- tryCatch({
    X <- stats::model.matrix(model)
    data.frame(
      test = "kappa",
      n = nrow(X),
      rank = qr(X)$rank,
      columns = ncol(X),
      kappa = kappa(X, exact = TRUE),

```

```

        status = "estimated",
        reason = NA_character_,
        stringsAsFactors = FALSE
    )
}, error = function(e) {
  data.frame(
    test = "kappa",
    n = NA_integer_,
    rank = NA_integer_,
    columns = NA_integer_,
    kappa = NA_real_,
    status = "out_of_active_pipeline",
    reason = conditionMessage(e),
    stringsAsFactors = FALSE
  )
})
out
}

#' compute design matrix rank
#'
compute_design_matrix_rank <- function(formula, data) {
  X <- stats::model.matrix(formula, data = data); tibble::tibble(rank = qr(X)$rank, ncol =
  ↪ ncol(X))
}

#' compute kappa
#'
compute_kappa <- function(formula, data) {
  X <- stats::model.matrix(formula, data = data); kappa(X, exact = TRUE)
}

#' compute vif if applicable
#'
compute_vif_if_applicable <- function(model) {
  if (requireNamespace("car", quietly = TRUE)) car::vif(model) else NULL
}

#' save multicollinearity diagnostics
#'
save_multicollinearity_diagnostics <- function(diagnostics) {
  diagnostics
}

```

Automated figure generation

Source: R/output/make_figures.R

```

figure_spec <- function(name, file, title, subtitle = NULL, kind = "status", variable =
  ↪ NULL, sources = NULL, inputs = NULL) {
  out <- list(
    name = name,
    file = file,
    title = title,
    subtitle = subtitle,
    kind = kind,

```

```

    variable = variable,
    sources = sources,
    inputs = inputs
  )
}

has_sf_geometry <- function(x) {
  inherits(x, "sf") && !is.null(attr(x, "sf_column")) && attr(x, "sf_column") %in% names(x)
}

sf_geometry_coverage <- function(x) {
  if (!has_sf_geometry(x) || !nrow(x)) return(0)
  mean(!sf::st_is_empty(sf::st_geometry(x)))
}

require_final_figure_inputs <- function(district_panel, cfg, required_variables) {
  if (!identical(cfg$mode, "final")) return(invisible(TRUE))
  missing_vars <- setdiff(required_variables, names(as.data.frame(district_panel)))
  if (length(missing_vars)) {
    stop(
      "Final figure generation requires mapped variables. Missing variables: ",
      paste(missing_vars, collapse = ", "),
      call. = FALSE
    )
  }
  if (!has_sf_geometry(district_panel)) {
    stop("Final map generation requires an sf district_panel with validated geometry.",
      call. = FALSE)
  }
  coverage <- sf_geometry_coverage(district_panel)
  if (!is.finite(coverage) || coverage < 0.75) {
    stop(
      "Final map generation requires a validated geometry join covering at least 75% of
      ↪ district-panel rows; current coverage is ",
      round(100 * coverage, 1),
      "%.",
      call. = FALSE
    )
  }
  invisible(TRUE)
}

#' make figures
#'
#' @return A named list of figure specifications consumed by save_figures().
make_figures <- function(district_panel, raw_ilo_figures, cfg) {
  required_variables <- c(
    "emie_2007",
    "consumption_growth_pct",
    "pucca_share_2007",
    "head_secondary_plus_2007",
    "region",
    "wavg_ling_degrees"
  )
}

```

```

out <- list(
  fig_ilo_trends = figure_spec(
    "fig_ilo_trends",
    "fig_ilo_trends.png",
    "ILO labor market indicators",
    "Composed from the archived ILO figure assets.",
    kind = "ilo_collage",
    sources = raw_ilo_figures
  ),
  district_carveouts_shifts = figure_spec(
    "district_carveouts_shifts",
    "district_carveouts_shifts.png",
    "District carve-outs and shifts",
    kind = "district_carveouts_shifts"
  )
)

missing_vars <- setdiff(required_variables, names(as.data.frame(district_panel)))
geometry_ok <- has_sf_geometry(district_panel) &&
↪ is.finite(sf_geometry_coverage(district_panel)) && sf_geometry_coverage(district_panel)
↪ >= 0.75
maps_available <- !length(missing_vars) && geometry_ok

map_specs <- list(
  map_emi_exposure = figure_spec("map_emi_exposure", "map_emi_exposure.png", "EMI
    ↪ Exposure", kind = if (maps_available) "map" else "status", variable = "emie_2007"),
  map_consumption_growth = figure_spec("map_consumption_growth",
    ↪ "map_consumption_growth.png", "% Change in Consumption", kind = if (maps_available)
    ↪ "map" else "status", variable = "consumption_growth_pct"),
  map_pucca = figure_spec("map_pucca", "map_pucca.png", "% Pucca Homes", kind = if
    ↪ (maps_available) "map" else "status", variable = "pucca_share_2007"),
  map_education = figure_spec("map_education", "map_education.png", "% HH Head w/ Sec.+",
    ↪ kind = if (maps_available) "map" else "status", variable =
    ↪ "head_secondary_plus_2007"),
  map_region = figure_spec("map_region", "map_region.png", "Region", kind = if
    ↪ (maps_available) "map" else "status", variable = "region"),
  map_linguistic_distance = figure_spec("map_linguistic_distance",
    ↪ "map_linguistic_distance.png", "Linguistic Distance", kind = if (maps_available)
    ↪ "map" else "status", variable = "wavg_ling_degrees"),
  collage_main_maps = figure_spec(
    "collage_main_maps",
    "collage_main_maps.png",
    "Main district-level map inputs",
    kind = "collage",
    inputs = c("map_emi_exposure", "map_consumption_growth", "map_pucca", "map_education")
  ),
  collage_iv_region_maps = figure_spec(
    "collage_iv_region_maps",
    "collage_iv_region_maps.png",
    "Instrument and region map inputs",
    kind = "collage",
    inputs = c("map_linguistic_distance", "map_region")
  )
)

```

```

if (!maps_available && !identical(cfg$mode, "final")) {
  # Draft-mode diagnostics live outside outputs/figures/main and are explicitly
  # labeled as diagnostics by figure_output_dir().
  map_specs <- map_specs[c("map_emi_exposure", "map_consumption_growth")]
}

map_input_failures <- character()
if (!maps_available && identical(cfg$mode, "final")) {
  map_input_failures <- c(
    if (length(missing_vars)) paste0("Missing map variables: ", paste(missing_vars,
      ↪ collapse = ", ")),
    paste0("Geometry coverage: ", round(100 * sf_geometry_coverage(district_panel), 1),
      ↪ "%")
  )
}

out <- c(out, map_specs)

if (length(map_input_failures)) {
  attr(out, "legacy_map_input_failures") <- map_input_failures
}
attr(out, "district_panel") <- district_panel
out
}

#' make ilo trends figure
#'
#' @return A vector of source image paths.
make_ilo_trends_figure <- function(raw_ilo_figures) {
  raw_ilo_figures
}

```

Publication-quality regression and summary tables

Source: R/output/make_tables.R

```

is_final_mode <- function(cfg) identical(cfg$mode, "final")

table_status_failures <- function(x, cfg, label) {
  df <- as.data.frame(x)
  ok_status <- c("mapped", "estimated")
  if (!is_final_mode(cfg) || !"status" %in% names(df)) return(character())
  bad <- !is.na(df$status) & !df$status %in% ok_status
  if (!any(bad)) return(character())
  reasons <- character()
  if ("reason" %in% names(df)) reasons <-
    ↪ unique(stats::na.omit(as.character(df$reason[bad])))
  suffix <- if (length(reasons)) paste0(" Reasons: ", paste(reasons, collapse = "; ")) else
    ↪ ""
  paste0("Final table generation requires completed model output for ", label, ".", suffix)
}

legacy_numeric_stats <- function(df, meta, cost_vars = character()) {
  df <- as.data.frame(df)
  rows <- lapply(seq_len(nrow(meta)), function(i) {
    v <- meta$var[[i]]

```

```

    if (!v %in% names(df)) return(NULL)
    x <- suppressWarnings(as.numeric(as.character(df[[v]])))
    ok <- is.finite(x)
    if (!any(ok)) return(NULL)
    stats <- c(
      Min = min(x[ok], na.rm = TRUE),
      `1Q` = unname(stats::quantile(x[ok], .25, na.rm = TRUE)),
      Med = stats::median(x[ok], na.rm = TRUE),
      `3Q` = unname(stats::quantile(x[ok], .75, na.rm = TRUE)),
      Max = max(x[ok], na.rm = TRUE),
      Mean = mean(x[ok], na.rm = TRUE),
      SD = stats::sd(x[ok], na.rm = TRUE)
    )
    z <- data.frame(var = v, label = meta$label[[i]], N = sum(ok), t(stats), check.names =
↪ FALSE, stringsAsFactors = FALSE)
    z
  })
  out <- safe_bind_rows(rows)
  if (!nrow(out)) return(data.frame(var = character(), label = character(), N = integer()))
  for (nm in intersect(c("Min", "1Q", "Med", "3Q", "Max", "Mean", "SD"), names(out))) {
    x <- suppressWarnings(as.numeric(out[[nm]]))
    out[[nm]] <- ifelse(out$var %in% cost_vars, formatC(x, format = "f", big.mark = ",",
↪ digits = 2), sprintf("%.2f", x))
  }
  out
}

legacy_categorical_stats <- function(df, meta) {
  df <- as.data.frame(df)
  rows <- lapply(seq_len(nrow(meta)), function(i) {
    v <- meta$var[[i]]
    if (!v %in% names(df)) return(NULL)
    x <- df[[v]]
    x <- x[!is.na(x)]
    if (!length(x)) return(NULL)
    freq <- table(x)
    data.frame(
      var = v,
      label = meta$label[[i]],
      N = length(x),
      Values = paste(names(freq), collapse = ", "),
      Mode = names(freq)[which.max(freq)],
      `% Mode` = round(max(freq) / sum(freq) * 100, 1),
      `Least Freq.` = names(freq)[which.min(freq)],
      `% Least Freq.` = round(min(freq) / sum(freq) * 100, 1),
      check.names = FALSE,
      stringsAsFactors = FALSE
    )
  })
  out <- safe_bind_rows(rows)
  if (!nrow(out)) data.frame(var = character(), label = character(), N = integer()) else out
}

numeric_summary <- function(df, variables = NULL) {
  df <- as.data.frame(df)

```

```

if (is.null(variables)) variables <- names(df)
variables <- intersect(variables, names(df))
out <- lapply(variables, function(v) {
  x <- suppressWarnings(as.numeric(as.character(df[[v]])))
  x <- x[is.finite(x)]
  if (!length(x)) return(NULL)
  data.frame(variable = v, min = min(x), q1 = unname(stats::quantile(x, 0.25)), median =
    ↪ stats::median(x), q3 = unname(stats::quantile(x, 0.75)), max = max(x), mean =
    ↪ mean(x), sd = stats::sd(x), n = length(x), stringsAsFactors = FALSE)
})
out <- safe_bind_rows(out)
if (!nrow(out)) data.frame(variable = character(), n = integer()) else out
}

categorical_summary <- function(df, variables) {
  df <- as.data.frame(df)
  variables <- intersect(variables, names(df))
  safe_bind_rows(lapply(variables, function(v) {
    x <- df[[v]]
    tab <- sort(table(x, useNA = "ifany"), decreasing = TRUE)
    if (!length(tab)) return(NULL)
    data.frame(variable = v, category = names(tab), n = as.integer(tab), share =
      ↪ round(as.numeric(tab) / sum(tab), 4), stringsAsFactors = FALSE)
  )))
}

table_status_row <- function(model, status = "unavailable", reason = NA_character_) {
  data.frame(
    model = model,
    term = NA_character_,
    estimate = NA_real_,
    std.error = NA_real_,
    statistic = NA_real_,
    p.value = NA_real_,
    status = status,
    reason = reason,
    stringsAsFactors = FALSE
  )
}

table_first_existing_column <- function(df, candidates) {
  hit <- intersect(candidates, names(df))
  if (length(hit)) hit[[1]] else NA_character_
}

table_column_or_na <- function(df, column) {
  if (length(column) != 1L || is.na(column) || !column %in% names(df)) {
    return(rep(NA_real_, nrow(df)))
  }
  suppressWarnings(as.numeric(df[[column]]))
}

coefficient_frame <- function(model, vcov_matrix = NULL) {
  estimates <- tryCatch(stats::coef(model), error = function(e) NULL)
  if (is.null(estimates) || !length(estimates)) return(data.frame())
}

```

```

terms <- names(estimates)
if (is.null(terms) || !length(terms)) terms <- paste0("term_", seq_along(estimates))
estimates <- suppressWarnings(as.numeric(estimates))

vc <- vcov_matrix
if (is.null(vc)) vc <- tryCatch(stats::vcov(model), error = function(e) NULL)
se <- rep(NA_real_, length(estimates))
if (!is.null(vc) && length(dim(vc)) == 2L && all(dim(vc) >= length(estimates))) {
  diag_vc <- suppressWarnings(as.numeric(diag(vc)))
  vc_terms <- rownames(vc)
  if (!is.null(vc_terms) && length(vc_terms)) {
    matched <- match(terms, vc_terms)
    ok <- !is.na(matched) & matched <= length(diag_vc)
    se[ok] <- sqrt(pmax(diag_vc[matched[ok]], 0))
  } else {
    se <- sqrt(pmax(diag_vc[seq_along(estimates)], 0))
  }
}

statistic <- estimates / se
statistic[!is.finite(statistic)] <- NA_real_
df_resid <- tryCatch(stats::df.residual(model), error = function(e) NA_real_)
p_value <- if (is.finite(df_resid) && df_resid > 0) {
  2 * stats::pt(abs(statistic), df = df_resid, lower.tail = FALSE)
} else {
  2 * stats::pnorm(abs(statistic), lower.tail = FALSE)
}
p_value[!is.finite(p_value)] <- NA_real_

out <- data.frame(
  Estimate = estimates,
  `Std. Error` = se,
  statistic = statistic,
  `Pr(>|t|)` = p_value,
  check.names = FALSE,
  stringsAsFactors = FALSE
)
rownames(out) <- terms
out
}

plain_model_coefficients <- function(model) {
  coefficient_frame(model)
}

clustered_model_coefficients <- function(model) {
  tryCatch({
    vc_fun <- clustered_model_vcov(model)
    if (is.null(vc_fun)) return(data.frame())
    vc <- vc_fun(model)
    coefficient_frame(model, vc)
  }, error = function(e) data.frame())
}

```

```

filter_table_model <- function(df, candidates) {
  if (!"model" %in% names(df)) return(df)
  out <- df[df$model %in% candidates, , drop = FALSE]
  if (nrow(out)) out else df[0L, , drop = FALSE]
}

is_model_status_payload <- function(x) {
  is.list(x) &&
  !inherits(x, c("lm", "ivreg")) &&
  !is.null(x$status) &&
  length(x$status) == 1L &&
  (is.character(x$status) || is.factor(x$status))
}

model_status_reason <- function(x) {
  reason <- x$reason
  if (is.null(reason) || !length(reason) || all(is.na(reason))) NA_character_ else
    ↪ as.character(reason[[1]])
}

tidy_iv_models <- function(iv_models) {
  if (!is.list(iv_models) || inherits(iv_models, c("lm", "ivreg"))) iv_models <- list(model
    ↪ = iv_models)
  safe_bind_rows(lapply(names(iv_models), function(name) {
    model <- iv_models[[name]]
    if (is_model_status_payload(model)) {
      return(table_status_row(name, as.character(model$status[[1]]),
        ↪ model_status_reason(model)))
    }
    coefs <- clustered_model_coefficients(model)
    if (!nrow(coefs)) coefs <- plain_model_coefficients(model)
    if (!nrow(coefs)) {
      return(table_status_row(name, "unavailable", "Model coefficients are unavailable."))
    }

    estimate_col <- table_first_existing_column(coefs, c("Estimate", "estimate"))
    se_col <- table_first_existing_column(coefs, c("Std. Error", "std.error", "Std.Error"))
    statistic_col <- table_first_existing_column(coefs, c("t value", "z value", "t", "z",
    ↪ "statistic"))
    p_col <- table_first_existing_column(coefs, c("Pr(>|t|)", "Pr(>|z|)", "p.value", "p",
    ↪ "P>|t|", "P>|z|"))

    if (is.na(estimate_col)) {
      return(table_status_row(name, "unavailable", paste("Model coefficient table lacks an
        ↪ estimate column. Columns:", paste(names(coefs), collapse = ", "))))
    }

    data.frame(
      model = name,
      term = rownames(coefs),
      estimate = table_column_or_na(coefs, estimate_col),
      std.error = table_column_or_na(coefs, se_col),
      statistic = table_column_or_na(coefs, statistic_col),
      p.value = table_column_or_na(coefs, p_col),
      status = "estimated",

```

```

      reason = NA_character_,
      stringsAsFactors = FALSE
    )
  }
}

clustered_model_vcov <- function(model) {
  cluster <- attr(model, "cluster_state")
  if (is.null(cluster)) return(NULL)
  cluster <- as.vector(cluster)
  if (length(cluster) != stats::nobs(model)) return(NULL)
  if (sum(!is.na(cluster)) < 2L || length(unique(cluster[!is.na(cluster)])) <= 1L)
    ↪ return(NULL)
  if (!requireNamespace("sandwich", quietly = TRUE)) return(NULL)
  force(cluster)
  function(x) sandwich::vcovCL(x, cluster = cluster)
}

#' make tables
#'
#' @return A named list of data frames consumed by save_tables().
make_tables <- function(selection_data, ame_results, district_panel, iv_models,
  ↪ first_stage_tests, cfg) {
  cons_iv <- tidy_iv_models(iv_models)
  cons_iv_required <- filter_table_model(cons_iv, c("consumption", "baseline"))
  table_failures <- c(
    table_status_failures(ame_results, cfg, "average marginal effects"),
    table_status_failures(first_stage_tests, cfg, "first-stage regression"),
    table_status_failures(cons_iv_required, cfg, "second-stage IV regression")
  )

  out <- list(
    selection_n = data.frame(n = nrow(as.data.frame(selection_data))),
    sum_tbl_probit_quant = make_selection_summary_numeric_table(selection_data),
    sum_tbl_probit_cat = make_selection_summary_categorical_table(selection_data),
    probit_mfx = make_probit_ame_table(ame_results),
    sum_tbl_iv = make_iv_summary_table(district_panel),
    fs_cons = make_first_stage_table(first_stage_tests, cfg),
    cons_iv = make_second_stage_table(iv_models),
    ame_results = as.data.frame(ame_results),
    first_stage = as.data.frame(first_stage_tests)
  )
  table_failures <- c(table_failures, attr(out$fs_cons, "legacy_table_input_failures"))
  table_failures <- unique(stats::na.omit(table_failures))
  if (length(table_failures)) attr(out, "legacy_table_input_failures") <- table_failures
  out
}

make_selection_summary_numeric_table <- function(selection_data) {
  meta <- data.frame(
    var = c("AGE", "HH_SIZE", "ENROLLMENT_COST", "dmean_num_IS_EDU_FREE",
    ↪ "dmean_num_TUTION_FEE_WAIVED", "dmean_num_RECD_SCHOLARSHIP_STIPEND",
    ↪ "dmean_num_RECD_TXT_BOOKS", "dmean_num_RECD_STATIONERY",
    ↪ "dmean_num_MID_DAY_MEAL_ETC_RECD", "dmean_num_ENROLLMENT_COST"),

```

```

    label = c("Age", "Household size", "Enrollment cost (Rs.)", "Educ. free available? (Yes
    ↪ = 1)", "Tuition waived?", "Scholarship/Stipend?", "Textbooks received?", "Stationery
    ↪ received?", "Mid-day meal or more received?", "Avg. district enrollment cost"),
    stringsAsFactors = FALSE
  )
  legacy_numeric_stats(selection_data, meta, cost_vars = "ENROLLMENT_COST")
}

make_selection_summary_categorical_table <- function(selection_data) {
  meta <- data.frame(
    var = c("SEX", "RELIGION", "SOCIAL_GROUP", "SECTOR", "DIST_FROM_NEAREST_PRIMARY_CLASS",
    ↪ "father_educ"),
    label = c("Sex", "Religion", "Social group", "Urban", "Distance of nearest primary
    ↪ class", "Father's education"),
    stringsAsFactors = FALSE
  )
  legacy_categorical_stats(selection_data, meta)
}

make_probit_ame_table <- function(ame_results) {
  out <- as.data.frame(ame_results, stringsAsFactors = FALSE)
  if (!"Term" %in% names(out)) out$Term <- out$term
  out$stars <- ifelse(is.finite(out$p.value) & out$p.value < 0.001, "***",
  ↪ ifelse(is.finite(out$p.value) & out$p.value < 0.01, "**", ifelse(is.finite(out$p.value)
  ↪ & out$p.value < 0.05, "*", "")))
  data.frame(
    Term = out$Term,
    Estimate = ifelse(is.finite(out$estimate), paste0(sprintf("%.3f", out$estimate),
    ↪ out$stars), NA_character_),
    `Std. Error` = ifelse(is.finite(out$std.error), sprintf("%.3f", out$std.error),
    ↪ NA_character_),
    check.names = FALSE,
    stringsAsFactors = FALSE
  )
}

make_iv_summary_table <- function(district_panel) {
  meta <- data.frame(
    var = c("EMIE", "wavg_ling_degrees", "npeople_0708", "consumption_0708",
    ↪ "gini_cons_0708", "pct_urban", "avg_hh_size", "dependency_ratio", "pct_fem_head",
    ↪ "pct_hindu", "pct_muslim", "pct_other_religion", "pct_st", "pct_sc", "pct_obc",
    ↪ "pct_small_land", "pct_medium_land", "pct_large_land", "pct_head_illiterate",
    ↪ "pct_head_lit_to_primary", "pct_head_secondary_plus", "pct_pucca", "npeople_1718",
    ↪ "consumption_1718", "gini_cons_1718", "consumption_pct_change", "gini_change"),
    label = c("EMIE", "Ling. Distance", "Population", "Consumption", "Gini of Consumption",
    ↪ "Pct. Urban", "Avg. HH Size", "Dependency Ratio × 100", "Pct. Female Head", "Pct.
    ↪ Hindu", "Pct. Muslim", "Pct. Other", "Pct. ST", "Pct. SC", "Pct. OBC", "Pct. Small
    ↪ Land-Owner", "Pct. Med. Land-Owner", "Pct. Large Land-Owner", "Pct. Head Educ.,
    ↪ Illiterate", "Pct. Head Educ., Lit.-Primary", "Pct. Head Educ., Secondary+", "Pct.
    ↪ Pucca", "Population", "Consumption", "Gini of Consumption",
    ↪ "$\\Delta\\text{Consumption}$", "$\\Delta\\text{Gini}\\sim\\text{Consumption}$"),
    stringsAsFactors = FALSE
  )
  legacy_numeric_stats(district_panel, meta, cost_vars = c("npeople_0708", "npeople_1718"))
}

```

```

make_first_stage_table <- function(first_stage_tests, cfg = list()) {
  fs <- as.data.frame(first_stage_tests, stringsAsFactors = FALSE)
  required_cols <- c("model", "term", "estimate", "std.error", "p.value", "partial_f",
    ↪ "partial_p", "status")
  missing_cols <- setdiff(required_cols, names(fs))
  if (length(missing_cols)) {
    msg <- paste("First-stage results are missing columns:", paste(missing_cols, collapse =
    ↪ ", "))
    if (is_final_mode(cfg)) stop(msg, call. = FALSE)
    return(table_status_row("first_stage", "unavailable", msg))
  }
  if ("model" %in% names(fs) && any(fs$model %in% c("consumption", "baseline"))) fs <-
    ↪ fs[fs$model %in% c("consumption", "baseline"), , drop = FALSE]
  if (any(grepl("[0-9]+$", as.character(fs$term)))) {
    msg <- "First-stage coefficient terms are numeric row positions; coefficient names were
    ↪ lost upstream."
    if (is_final_mode(cfg)) stop(msg, call. = FALSE)
    return(table_status_row("first_stage", "malformed", msg))
  }
  status_reasons <- unique(stats::na.omit(as.character(fs$reason[!is.na(fs$status) &
    ↪ fs$status != "estimated"])))
  fs <- fs[fs$status == "estimated" & fs$term %in% c("wavg_ling_degrees", "(Intercept)"), ,
    ↪ drop = FALSE]
  required_terms <- c("wavg_ling_degrees", "(Intercept)")
  missing_terms <- setdiff(required_terms, fs$term)
  if (length(missing_terms)) {
    if (length(status_reasons)) {
      msg <- paste(status_reasons, collapse = "; ")
      out <- table_status_row("first_stage", "out_of_active_pipeline", msg)
      attr(out, "legacy_table_input_failures") <- paste0("First-stage table lacks required
      ↪ coefficient(s): ", paste(missing_terms, collapse = ", "), ". Reasons: ", msg)
      return(out)
    }
    msg <- paste("First-stage table lacks required coefficient(s):", paste(missing_terms,
    ↪ collapse = ", "))
    if (is_final_mode(cfg)) stop(msg, call. = FALSE)
    return(table_status_row("first_stage", "unavailable", msg))
  }
  fs <- fs[match(required_terms, fs$term), , drop = FALSE]
  fs$stars <- ifelse(is.finite(fs$p.value) & fs$p.value < 0.001, "****",
    ↪ ifelse(is.finite(fs$p.value) & fs$p.value < 0.01, "***", ifelse(is.finite(fs$p.value) &
    ↪ fs$p.value < 0.05, "**", "")))
  out <- data.frame(
    Term = ifelse(fs$term == "wavg_ling_degrees", "Linguistic distance", "Constant"),
    Estimate = paste0(sprintf("%.3f", fs$estimate), fs$stars),
    `Std. Error` = sprintf("%.3f", fs$std.error),
    check.names = FALSE,
    stringsAsFactors = FALSE
  )
  stat <- fs[fs$term == "wavg_ling_degrees", , drop = FALSE]
  if (nrow(stat)) {
    f_value <- first_finite_value(stat, c("partial_f", "legacy_model_f"))
    f_p <- first_finite_value(stat, c("partial_p", "legacy_model_p"))
    out <- rbind(out, data.frame(

```

```

    Term = "Instrument's F-Statistic",
    Estimate = paste0(sprintf("%.2f", f_value), ifelse(is.finite(f_p) & f_p < 0.001,
      ↪ "***", ifelse(is.finite(f_p) & f_p < 0.01, "**", ifelse(is.finite(f_p) & f_p <
      ↪ 0.05, "*", ""))))),
    `Std. Error` = "",
    check.names = FALSE,
    stringsAsFactors = FALSE
  ))
}
out
}

first_finite_value <- function(df, cols) {
  for (col in cols) {
    if (!col %in% names(df)) next
    value <- suppressWarnings(as.numeric(df[[col]][[1]]))
    if (length(value) && is.finite(value)) return(value)
  }
  NA_real_
}

make_second_stage_table <- function(iv_models) {
  out <- tidy_iv_models(iv_models)
  out <- filter_table_model(out, c("consumption", "baseline"))
  if (!nrow(out)) return(data.frame(Term = character(), Estimate = character(), `Std. Error`
    ↪ = character(), check.names = FALSE))
  out$stars <- ifelse(is.finite(out$p.value) & out$p.value < 0.001, "***",
    ↪ ifelse(is.finite(out$p.value) & out$p.value < 0.01, "**", ifelse(is.finite(out$p.value)
    ↪ & out$p.value < 0.05, "*", "")))
  data.frame(
    Term = legacy_iv_term_label(out$term),
    Estimate = ifelse(is.finite(out$estimate), paste0(sprintf("%.3f", out$estimate),
      ↪ out$stars), NA_character_),
    `Std. Error` = ifelse(is.finite(out$std.error), sprintf("(%.3f)", out$std.error),
      ↪ NA_character_),
    check.names = FALSE,
    stringsAsFactors = FALSE
  )
}

legacy_iv_term_label <- function(term) {
  labels <- c(
    "(Intercept)" = "Constant",
    EMIE = "EMIE",
    emie_2007 = "EMIE",
    wavg_ling_degrees = "Linguistic distance",
    consumption_0708 = "Consumption, 2007-08",
    gini_cons_0708 = "Gini consumption, 2007-08",
    pct_urban = "Pct. urban",
    avg_hh_size = "Avg. HH size",
    dependency_ratio = "Dependency ratio",
    pct_fem_head = "Pct. female head",
    pct_hindu = "Pct. Hindu",
    pct_muslim = "Pct. Muslim",
    pct_st = "Pct. ST",

```

```

    pct_sc = "Pct. SC",
    pct_obc = "Pct. OBC",
    pct_small_land = "Pct. small land",
    pct_medium_land = "Pct. medium land",
    pct_large_land = "Pct. large land",
    pct_head_lit_to_primary = "Pct. head literate-primary",
    pct_head_secondary_plus = "Pct. head secondary+"
  )
  out <- unname(labels[term])
  out[is.na(out)] <- term[is.na(out)]
  out
}

make_diagnostic_tables <- function(...) list()

```