

Coding Sample: Escaping Inequality in India

Selected Excerpts

Rishav Roy

What follows are excerpts from the 138-page code appendix to my independent paper and writing sample, “Escaping Inequality in India: The Role of English-Medium Instruction.”

The next 25 pages illustrate data ingestion and harmonization, complex survey econometrics, IV estimation, and careful diagnostics. Each excerpt is a part of the code chunks listed below:

14-16. Scalable **import, cleaning, and merging** of multi-file census and geospatial data.

- Chunk 4: Import key data files

31-32. **QA and identifier-disambiguation checks** for district-level crosswalk.

- Chunk 6: Match districts: Construct district tracking dfs

50-52. **Testing for systematic missingness** using the Benjamini–Hochberg procedure and parallelized logistic regressions.

- Chunk 8: Participation model: Are NAs randomly distributed

55-57. **Survey-weighted enrollment probit estimation** and average marginal effects derivation under a complex survey design. The enrollment model is reported as a stand-alone descriptive selection model; district-level inverse-Mills-ratio aggregation is not part of the active analysis.

- Chunk 9: Participation model: Survey-weighted estimation
- Chunk 10: Participation model: Derive AME

80-88. A reviewed **district-harmonization map contract** with explicit many-to-many QA, while source attachment and candidate diagnostics remain separated from the public crosswalk authority.

- Chunk 16: Match districts: Test joining methods
- Chunk 17: Match districts: Define joining functions

105-107. Reusable **IV specification and estimation pipeline** with clustered inference, plus contiguity-based spatial weights and setup for spatially lagged models.

- Chunk 24: Geospatial: Neighbor list construction
- Chunk 25: 2SLS: Controls and 2SLS formula
- Chunk 26: 2SLS: Baseline 2SLS models

111-112. **Robustness checks and diagnostics** for multicollinearity, spatial autocorrelation.

- Chunk 28: Multicollinearity check
- Chunk 29: Geospatial: Spatial autocorrelation tests

Scalable import, cleaning, and merging of Census/geospatial data

Source: R/io/read_inputs.R

```
## read nss 2007 education
##
read_nss_2007_education <- function(paths) {
  read_manifest_group(paths, "nss_2007_education")
}

## read nss 2007 consumption
##
read_nss_2007_consumption <- function(paths) {
  read_manifest_group(paths, "nss_2007_consumption")
}

## read nss 2017 education
##
read_nss_2017_education <- function(paths) {
  read_manifest_group(paths, "nss_2017_education")
}

## read census 2001 mother tongue
##
read_census_2001_mother_tongue <- function(paths) {
  read_manifest_group(paths, "census_2001_mother_tongue")
}

## read district boundaries 2020
##
read_district_boundaries_2020 <- function(paths) {
  rows <- require_manifest_files(paths, "district_boundaries_2020")
  shp <- rows[tolower(rows$file_type) == "shp", , drop = FALSE]
  if (!nrow(shp)) stop("The district-boundary manifest includes shapefile sidecars but no
    ↪ .shp row.", call. = FALSE)
  need_pkg("sf", "district_boundaries")
  sf::st_read(shp$absolute_path[[1]], quiet = TRUE)
}

## read district change sources
##
read_district_change_sources <- function(paths) {
  read_manifest_group(paths, "district_changes")
}

## list ilo figure paths
##
list_ilo_figure_paths <- function(paths) {
  rows <- require_manifest_files(paths, "ilo_figures")
  stats::setNames(rows$absolute_path, rows$file_id)
}

## Join one label split by table line wrapping
join_district_carveout_wrapped_label <- function(previous, continuation) {
  previous <- trimws(as.character(previous))
  continuation <- trimws(as.character(continuation))
```

```

if (!nzchar(previous) || !nzchar(continuation)) {
  stop("District carve-out wrapped labels must be nonblank.", call. = FALSE)
}
if (!grepl("-", previous)) return(paste(previous, continuation))
if (grepl("^[:lower:]", continuation)) {
  return(paste0(sub("-", "", previous), continuation))
}
paste0(previous, continuation)
}

#' Read the headerless 1961-2001 district carve-out source
#'
#' The source has five data columns and no header row. Reading it with the
#' ordinary CSV default would consume the first Anantapur observation as names.
repair_district_carveout_wrapped_rows <- function(x) {
  out <- safe_df(x)
  if (!nrow(out)) return(out)
  blank <- function(value) is.na(value) || !nzchar(trimws(as.character(value)))
  drop <- rep(FALSE, nrow(out))
  for (i in seq_len(nrow(out))) {
    if (i == 1L || !blank(out$pop_1991[[i]])) next
    previous <- i - 1L
    source_continuation <- !blank(out$district_1991[[i]])
    target_continuation <- !blank(out$district_2001[[i]])
    has_transfer <- !blank(out$pct_01in91[[i]]) || !blank(out$pct_91in01[[i]])

    if (source_continuation) {
      previous_name <- plain_chr(out$district_1991[[previous]])
      if (!length(previous_name) || !grepl("-", trimws(previous_name[[1L]]))) {
        stop("Unexpected wrapped row in the district carve-out source.", call. = FALSE)
      }
      continuation <- plain_chr(out$district_1991[[i]])[[1L]]
      combined <- join_district_carveout_wrapped_label(previous_name[[1L]], continuation)
      out$district_1991[[previous]] <- combined
      if (has_transfer) {
        out$district_1991[[i]] <- combined
        out$pop_1991[[i]] <- out$pop_1991[[previous]]
      }
    }

    if (!has_transfer && (source_continuation || target_continuation)) {
      if (target_continuation) {
        if (blank(out$district_2001[[previous]])) {
          stop("Unexpected target-only continuation in the district carve-out source.",
               call. = FALSE)
        }
        previous_target <- trimws(as.character(out$district_2001[[previous]]))
        continuation <- trimws(as.character(out$district_2001[[i]]))
        out$district_2001[[previous]] <- join_district_carveout_wrapped_label(
          previous_target, continuation
        )
      }
      drop[[i]] <- TRUE
    }
  }
}

```

```

  out[!drop, , drop = FALSE]
}

district_carveout_rounding_tolerance_pp <- function() 0.05

validate_district_carveout_shares <- function(
  x,
  rounding_tolerance_pp = district_carveout_rounding_tolerance_pp()) {
  out <- safe_df(x)
  if (!nrow(out)) return(invisible(TRUE))
  share_cols <- c("pct_01in91", "pct_91in01")
  for (field in share_cols) {
    values <- num(out[[field]])
    bad <- is.finite(values) & (values < 0 | values > 100)
    if (any(bad)) {
      stop("District carve-out transfer shares must lie in [0, 100].", call. = FALSE)
    }
  }
  source_share <- num(out$pct_01in91)
  if (any(!is.finite(source_share))) {
    stop("District carve-out source shares must be complete.", call. = FALSE)
  }
  source_key <- paste(out$district_1991, out$pop_1991, sep = "__")
  source_sum <- vapply(split(source_share, source_key), sum, numeric(1))
  if (any(abs(source_sum - 100) > rounding_tolerance_pp)) {
    stop("District carve-out source shares must sum to 100 within rounding tolerance.",
      ↪ call. = FALSE)
  }
  invisible(TRUE)
}

read_district_carveouts <- function(path) {
  out <- utils::read.csv(
    path,
    header = FALSE,
    col.names = c("district_1991", "pop_1991", "district_2001", "pct_01in91",
      ↪ "pct_91in01"),
    stringsAsFactors = FALSE,
    na.strings = c("", "NA"),
    check.names = FALSE
  )
  out <- repair_district_carveout_wrapped_rows(out)
  fill_down <- function(x) {
    if (!length(x)) return(x)
    for (i in seq_along(x)) {
      missing <- is.na(x[[i]]) || (is.character(x) && !nzchar(trimws(x[[i]])))
      if (missing && i > 1L) x[[i]] <- x[[i - 1L]]
    }
    x
  }
  out$district_1991 <- fill_down(out$district_1991)
  out$pop_1991 <- num(gsub(",", "", fill_down(out$pop_1991), fixed = TRUE))
  out$pct_01in91 <- num(out$pct_01in91)
  out$pct_91in01 <- num(out$pct_91in01)
  validate_district_carveout_shares(out)
}

```

```

    out
  }

  #' read one manifest row
  #'
  #' @return Reader output for raw data files, or a validated path for sidecars/assets.
  read_by_manifest_row <- function(row) {
    path <- row$absolute_path[[1]]
    reader <- as.character(row$reader_function[[1]] %||% "")
    if (identical(reader, "read_district_carveouts")) {
      return(read_district_carveouts(path))
    }
    file_type <- tolower(row$file_type[[1]])
    switch(
      file_type,
      sav = read_sav_short(path),
      csv = read_csv_short(path),
      xls = read_excel_short(path),
      xlsx = read_excel_short(path),
      ods = read_ods_short(path),
      shp = {
        need_pkg("sf", "shapefiles")
        sf::st_read(path, quiet = TRUE)
      },
      shx = path,
      dbf = path,
      prj = path,
      png = path,
      stop("No reader implemented for ", file_type, call. = FALSE)
    )
  }

  #' read all manifest rows for a source
  #'
  #' @return Named list of reader outputs keyed by manifest file_id.
  read_manifest_group <- function(paths, source_id) {
    rows <- require_manifest_files(paths, source_id)
    stats::setNames(lapply(seq_len(nrow(rows)), function(i) read_by_manifest_row(rows[i,
    ↪ ]))), rows$file_id)
  }

```

QA and identifier-disambiguation checks

Source: R/districts/build_district_tracker.R

```

  #' build district tracker
  #'
  build_district_tracker <- function(raw_district_changes) {
    source_aliases <- list(
      alluvial = c("alluvial", "district_changes_alluvial"),
      carveouts_renamings = c("carveouts_renamings", "district_changes_carveouts"),
      new_districts_created = c("new_districts_created", "district_changes_new_districts"),
      name_changes = c("name_changes", "district_changes_name_changes"),
      district_splits = c("district_splits", "district_changes_splits"),
      india_district_tracker = c("india_district_tracker", "tracker",
      ↪ "district_changes_tracker")
    )
  }

```

```

)
source_value <- function(id) {
  hit <- source_aliases[[id]][source_aliases[[id]] %in% names(raw_district_changes)]
  if (length(hit)) raw_district_changes[[hit[[1]]]] else data.frame()
}
parsed <- list(
  alluvial = parse_alluvial_district_changes(source_value("alluvial")),
  carveouts_renamings = parse_carveouts_renamings(source_value("carveouts_renamings")),
  new_districts_created =
    ↪ parse_new_districts_created(source_value("new_districts_created")),
  name_changes = parse_name_changes(source_value("name_changes")),
  district_splits = parse_district_splits(source_value("district_splits")),
  india_district_tracker =
    ↪ parse_india_district_tracker(source_value("india_district_tracker"))
)
consumed <- unique(unlist(source_aliases, use.names = FALSE))
extras <- setdiff(names(raw_district_changes), consumed)
for (nm in extras) parsed[[nm]] <-
  ↪ parse_district_change_source(raw_district_changes[[nm]], source_type = nm)
standardize_tracker_names(standardize_tracker_years(combine_district_tracker_sources
  ↪ (parsed)))
}

#' parse alluvial district changes
#'
parse_alluvial_district_changes <- function(x) {
  x <- safe_df(x)
  year_state <- grep("[0-9]{4}.*state", canon(names(x)), value = FALSE)
  years <- suppressWarnings(as.integer(sub("^[0-9]{4}).*$", "\\1",
    ↪ canon(names(x)[year_state]))))
  years <- sort(unique(years[is.finite(years)]))
  if (length(years) < 2L) return(parse_district_change_source(x, source_type =
    ↪ "alluvial"))
  parse_wide_year_lineages(x, years, source_type = "alluvial")
}

#' parse india district tracker
#'
parse_india_district_tracker <- function(x) {
  x <- safe_df(x)
  years <- suppressWarnings(as.integer(names(x)[grepl("^[0-9]{4}$", names(x))]))
  years <- sort(unique(years[is.finite(years)]))
  if (length(years) < 2L) return(parse_district_change_source(x, source_type =
    ↪ "india_district_tracker"))
  parse_tracker_year_triplets(x, years)
}

#' parse carveouts renamings
#'
parse_carveouts_renamings <- function(x) {
  x <- safe_df(x)
  required <- c("district_1991", "pop_1991", "district_2001", "pct_01in91", "pct_91in01")
  if (!all(required %in% names(x))) {
    return(parse_district_change_source(x, source_type = "carveouts_renamings"))
  }
}

```

```

data.frame(
  source_type = "carveouts_renamings",
  source_state_raw = NA_character_,
  source_district_raw = plain_chr(x$district_1991),
  target_state_raw = NA_character_,
  target_district_raw = plain_chr(x$district_2001),
  source_year_raw = 1991L,
  target_year_raw = 2001L,
  change_type = "population_transfer_1991_2001",
  event_year = 2001L,
  source_population = num(gsub(",", "", plain_chr(x$pop_1991), fixed = TRUE)),
  source_share_to_target = num(x$pct_01in91) / 100,
  target_share_from_source = num(x$pct_91in01) / 100,
  stringsAsFactors = FALSE
)
}

#' parse new districts created
#'
parse_new_districts_created <- function(x) {
  x <- safe_df(x)
  if (!all(c("State/UT", "Old districts", "New District") %in% names(x))) {
    return(parse_district_change_source(x, source_type = "new_districts_created"))
  }
  data.frame(
    source_type = "new_districts_created",
    source_state_raw = plain_chr(x[["State/UT"]]),
    source_district_raw = plain_chr(x[["Old districts"]]),
    target_state_raw = plain_chr(x[["State/UT"]]),
    target_district_raw = plain_chr(x[["New District"]]),
    source_year_raw = suppressWarnings(as.integer(x$Year)) - 1L,
    target_year_raw = suppressWarnings(as.integer(x$Year)),
    change_type = "new_district",
    event_year = suppressWarnings(as.integer(x$Year)),
    decade = plain_chr(x$Decade),
    stringsAsFactors = FALSE
  )
}

#' parse name changes
#'
parse_name_changes <- function(x) {
  x <- safe_df(x)
  if (!all(c("State/UT", "Old Name", "New Name") %in% names(x))) {
    return(parse_district_change_source(x, source_type = "name_changes"))
  }
  decade_end <- suppressWarnings(as.integer(sub(".*-", "", plain_chr(x$Decade))))
  data.frame(
    source_type = "name_changes",
    source_state_raw = plain_chr(x[["State/UT"]]),
    source_district_raw = plain_chr(x[["Old Name"]]),
    target_state_raw = plain_chr(x[["State/UT"]]),
    target_district_raw = plain_chr(x[["New Name"]]),
    source_year_raw = NA_integer_,
    target_year_raw = decade_end,

```

```

    change_type = paste0("name_change:", canon(x$Type)),
    event_year = NA_integer_,
    decade = plain_chr(x$Decade),
    stringsAsFactors = FALSE
  )
}

#' parse district splits
#'
parse_district_splits <- function(x) {
  x <- safe_df(x)
  if (!all(c("State/UT", "District-Before", "District-After") %in% names(x))) {
    return(parse_district_change_source(x, source_type = "district_splits"))
  }
  year <- suppressWarnings(as.integer(x$Year))
  data.frame(
    source_type = "district_splits",
    source_state_raw = plain_chr(x[["State/UT"]]),
    source_district_raw = plain_chr(x[["District-Before"]]),
    target_state_raw = plain_chr(x[["State/UT"]]),
    target_district_raw = plain_chr(x[["District-After"]]),
    source_year_raw = year - 1L,
    target_year_raw = year,
    change_type = "split_or_carveout",
    event_year = year,
    decade = plain_chr(x$Decade),
    stringsAsFactors = FALSE
  )
}

find_year_field <- function(x, year, kind) {
  keys <- canon(names(x))
  pattern <- paste0("^", year, " ", kind)
  hit <- which(grepl(pattern, keys))
  if (length(hit)) names(x)[hit[[1]]] else NULL
}

parse_wide_year_lineages <- function(x, years, source_type) {
  rows <- lapply(seq_len(length(years) - 1L), function(i) {
    source_year <- years[[i]]
    target_year <- years[[i + 1L]]
    ss <- find_year_field(x, source_year, "state")
    sd <- find_year_field(x, source_year, "district")
    ts <- find_year_field(x, target_year, "state")
    td <- find_year_field(x, target_year, "district")
    if (any(vapply(list(ss, sd, ts, td), is.null, logical(1)))) return(data.frame())
    out <- data.frame(
      source_type = source_type,
      source_state_raw = plain_chr(x[[ss]]),
      source_district_raw = plain_chr(x[[sd]]),
      target_state_raw = plain_chr(x[[ts]]),
      target_district_raw = plain_chr(x[[td]]),
      source_year_raw = source_year,
      target_year_raw = target_year,

```



```

    stringsAsFactors = FALSE
  )
  source_pair <- paste(canonicalize_state_name(out$source_state_raw),
    ↪ canon(out$source_district_raw), sep = "__")
  target_pair <- paste(canonicalize_state_name(out$target_state_raw),
    ↪ canon(out$target_district_raw), sep = "__")
  out$change_type <- ifelse(source_pair == target_pair, "continuity",
    ↪ "lineage_candidate")
  out$.source_wide_row <- seq_len(nrow(x))
  keep <- !is.na(out$source_district_raw) & nzchar(trimws(out$source_district_raw)) &
    !is.na(out$target_district_raw) & nzchar(trimws(out$target_district_raw))
  out[keep, , drop = FALSE]
})
safe_bind_rows(rows)
}

parse_tracker_year_triplets <- function(x, years) {
  # The Jaacks sheet has state, district, and code columns in repeating
  # three-column year blocks, with the first data row containing field labels.
  if (nrow(x) && any(canon(unlist(x[1, ], use.names = FALSE)) == "statename")) x <- x[-1,
    ↪ , drop = FALSE]
  rows <- lapply(seq_len(length(years) - 1L), function(i) {
    source_year <- years[[i]]
    target_year <- years[[i + 1L]]
    source_col <- match(as.character(source_year), names(x))
    target_col <- match(as.character(target_year), names(x))
    if (!is.finite(source_col) || !is.finite(target_col)) return(data.frame())
    out <- data.frame(
      source_type = "india_district_tracker",
      source_state_raw = plain_chr(x[[source_col]]),
      source_district_raw = plain_chr(x[[source_col + 1L]]),
      source_code_raw = plain_chr(x[[source_col + 2L]]),
      target_state_raw = plain_chr(x[[target_col]]),
      target_district_raw = plain_chr(x[[target_col + 1L]]),
      target_code_raw = plain_chr(x[[target_col + 2L]]),
      source_year_raw = source_year,
      target_year_raw = target_year,
      stringsAsFactors = FALSE
    )
    source_pair <- paste(canonicalize_state_name(out$source_state_raw),
    ↪ canon(out$source_district_raw), sep = "__")
    target_pair <- paste(canonicalize_state_name(out$target_state_raw),
    ↪ canon(out$target_district_raw), sep = "__")
    out$change_type <- ifelse(source_pair == target_pair, "continuity",
    ↪ "annual_lineage_candidate")
    out$.source_wide_row <- seq_len(nrow(x))
    keep <- !is.na(out$source_district_raw) & nzchar(trimws(out$source_district_raw)) &
      !is.na(out$target_district_raw) & nzchar(trimws(out$target_district_raw))
    out[keep, , drop = FALSE]
  })
  safe_bind_rows(rows)
}

parse_district_change_source <- function(x, source_type) {
  x <- safe_df(x)

```

```

  if (!nrow(x)) return(data.frame(source_type = character(), stringsAsFactors = FALSE))
  x$source_type <- source_type
  x$source_state_raw <- first_present_value(x, c("state", "State", "state_raw",
↪ "state_from", "state_01", "state_07", "state_08", "state_17", "state_18",
↪ "state_20"))
  x$source_district_raw <- first_present_value(x, c("district", "District",
↪ "district_raw", "district_from", "district_01", "district_07", "district_08",
↪ "district_17", "district_18", "district_20", "district_name"))
  x$target_state_raw <- first_present_value(x, c("state_to", "new_state", "target_state",
↪ "state_20", "state_18", "state_17", "state_08", "state_07", "state"))
  x$target_district_raw <- first_present_value(x, c("district_to", "new_district",
↪ "target_district", "district_20", "district_18", "district_17", "district_08",
↪ "district_07", "district"))
  x$source_year_raw <- first_present_value(x, c("source_year", "year_from", "from_year",
↪ "start_year", "year"))
  x$target_year_raw <- first_present_value(x, c("target_year", "year_to", "to_year",
↪ "end_year", "year"))
  x$change_type <- first_present_value(x, c("change_type", "type", "event_type",
↪ "status"))
  x
}

first_present_value <- function(df, candidates) {
  hit <- first_col(df, candidates)
  if (is.null(hit)) return(rep(NA_character_, nrow(df)))
  as.character(df[[hit]])
}

#' combine district tracker sources
#'
combine_district_tracker_sources <- function(raw_district_changes) {
  out <- safe_bind_rows(lapply(names(raw_district_changes), function(name) {
    x <- safe_df(raw_district_changes[[name]])
    if (!nrow(x)) return(data.frame())
    x$source_file_id <- if ("source_file_id" %in% names(x)) x$source_file_id else name
    if ("source_type" %in% names(x)) x$source_type <- name
    x
  }))
  if (nrow(out)) out$.row_in_source <- ave(seq_len(nrow(out)), out$source_file_id, FUN =
↪ seq_along)
  out
}

#' standardize tracker years
#'
standardize_tracker_years <- function(tracker) {
  tracker <- safe_df(tracker)
  if (!nrow(tracker)) return(tracker)
  if ("source_year_raw" %in% names(tracker) && !"source_year" %in% names(tracker))
    ↪ tracker$source_year <- suppressWarnings(as.integer(tracker$source_year_raw))
  if ("target_year_raw" %in% names(tracker) && !"target_year" %in% names(tracker))
    ↪ tracker$target_year <- suppressWarnings(as.integer(tracker$target_year_raw))
  tracker
}

```

```

#' standardize tracker names
#'
standardize_tracker_names <- function(tracker) {
  tracker <- safe_df(tracker)
  if (!nrow(tracker)) return(tracker)
  if ("source_state_raw" %in% names(tracker)) tracker$source_state_key <-
    ↪ canonicalize_state_name(tracker$source_state_raw)
  if ("source_district_raw" %in% names(tracker)) tracker$source_district_key <-
    ↪ canon(tracker$source_district_raw)
  if ("target_state_raw" %in% names(tracker)) tracker$target_state_key <-
    ↪ canonicalize_state_name(tracker$target_state_raw)
  if ("target_district_raw" %in% names(tracker)) tracker$target_district_key <-
    ↪ canon(tracker$target_district_raw)
  tracker
}

```

Parallelized missingness diagnostics with BH correction

Source: R/selection/diagnose_missingness.R

```

missingness_variables <- function(selection_data) {
  df <- as.data.frame(selection_data, stringsAsFactors = FALSE)
  all_child_missing_vars <- c("DIST_FROM_NEAREST_PRIMARY_CLASS",
    ↪ "dmean_num_ENROLLMENT_COST", "father_educ")
  enrolled_only_missing_vars <- c("TUTION_FEE", "EXAMINATION_FEE", "OTHER_FEES_PAYMENTS",
    ↪ "BOOKS", "STATIONERY", "UNIFORM", "TRANSPORT")
  schooling_benefit_vars <- c(
    "IS_EDU_FREE", "TUTION_FEE_WAIVED", "RECD_SCHOLARSHIP_STIPEND",
    "RECD_TXT_BOOKS", "RECD_STATIONERY", "MID_DAY_MEAL_ETC_RECD"
  )

  # Legacy Chunk 8 distinguished variables used in the probit from
  # enrolled-only expenditure variables. Keep that distinction so the total
  # row labelled "probit-model" is not inflated by fee variables that are
  # undefined for non-enrolled children.
  probit_vars <- c(
    "enrolled", "ENROLLED", "AGE", "age", "SEX", "HH_SIZE", "RELIGION", "SOCIAL_GROUP",
    "SECTOR", "state_0708", "region_0708", all_child_missing_vars
  )
  probit_vars <- intersect(probit_vars, names(df))
  if (!length(probit_vars)) probit_vars <- setdiff(names(df),
    ↪ intersect(enrolled_only_missing_vars, names(df)))

  list(
    probit_vars = probit_vars,
    miss_vars_all = intersect(all_child_missing_vars, names(df)),
    miss_vars_enrolled = intersect(enrolled_only_missing_vars, names(df)),
    schooling_benefit_vars = intersect(schooling_benefit_vars, names(df)),
    group_vars = intersect(c("SECTOR", "SEX", "RELIGION", "SOCIAL_GROUP", "state_0708"),
      ↪ names(df)),
    cts_vars = intersect(c("AGE", "HH_SIZE"), names(df)),
    regional_vars = intersect(c("state_0708", "region_0708"), names(df))
  )
}

```

```

#' diagnose missingness
#'
#' Port the legacy Chunk 8 missingness diagnostics into an opt-in diagnostic
#' object. The returned list preserves the legacy sequence: variable counts,
#' regional rankings, missingness correlation matrices, BH-adjusted logit screens,
#' and notes for commented case-study / chi-square checks.
diagnose_missingness <- function(selection_data, cfg) {
  df <- as.data.frame(selection_data, stringsAsFactors = FALSE)
  vars <- missingness_variables(df)
  probit_vars <- vars$probit_vars

  counts <- summarize_missingness_by_variable(df[probit_vars])
  if (length(probit_vars)) {
    total_na <- sum(!stats::complete.cases(df[probit_vars]))
    total_complete <- sum(stats::complete.cases(df[probit_vars]))
  } else {
    total_na <- 0L
    total_complete <- nrow(df)
  }
  counts <- safe_bind_rows(list(
    counts,
    data.frame(missing_var = "Total probit-model with NA", n_missing = total_na,
      ↪ pct_missing = if (nrow(df)) total_na / nrow(df) else NA_real_),
    data.frame(missing_var = "Total probit-model with no NA", n_missing = total_complete,
      ↪ pct_missing = if (nrow(df)) total_complete / nrow(df) else NA_real_)
  ))

  regional <- summarize_missingness_regions(df, probit_vars, vars)
  corr_all <- missingness_correlation_matrix(df, vars$miss_vars_all, vars$group_vars,
  ↪ vars$cts_vars)
  enrolled_rows <- enrolled_schooling_rows(df)
  corr_enrolled <- missingness_correlation_matrix(df[enrolled_rows, , drop = FALSE],
  ↪ vars$miss_vars_enrolled, vars$group_vars, vars$cts_vars)
  benefit_missingness <- summarize_missingness_by_variable(
    df[enrolled_rows, vars$schooling_benefit_vars, drop = FALSE]
  )
  if (nrow(benefit_missingness)) benefit_missingness$scope <- "enrolled_children"

  covars <- intersect(c("SECTOR", "SEX", "AGE", "HH_SIZE", "RELIGION", "SOCIAL_GROUP",
  ↪ "state_0708"), names(df))
  miss_all <- intersect(vars$miss_vars_all, names(df))
  miss_enrolled <- intersect(vars$miss_vars_enrolled, names(df))
  logit_all <- if (length(miss_all) && length(covars)) run_missingness_logits(df,
  ↪ miss_all, covars) else data.frame()
  logit_enrolled <- if (length(miss_enrolled) && length(covars) && any(enrolled_rows))
  ↪ run_missingness_logits(df[enrolled_rows, , drop = FALSE], miss_enrolled, covars) else
  ↪ data.frame()
  logit_summary <- summarize_missingness_logits(safe_bind_rows(list(logit_all,
  ↪ logit_enrolled)))
  case_study <- summarize_missingness_case_study(df, vars)
  chi_square <- summarize_missingness_chisq(df, probit_vars)

  notes <- data.frame(
    diagnostic = c(
      "rajasthan_southern_case_study",

```

```

    "chi_square_any_na_by_state",
    "chi_square_any_na_by_region",
    "parallel_missingness_logit"
  ),
  legacy_status = c(
    "commented diagnostic View() code preserved as documented note",
    "commented diagnostic test not run by default",
    "commented diagnostic test not run by default",
    "ported using lapply for reproducible targets execution; legacy mclapply/parLapply
      ↪ choice documented"
  ),
  stringsAsFactors = FALSE
)

out <- list(
  missing_counts = counts,
  regional_cost = regional$cost,
  regional_distance = regional$distance,
  regional_father_education = regional$father_education,
  corr_all = corr_all,
  corr_enrolled = corr_enrolled,
  logit_all = logit_all,
  logit_enrolled = logit_enrolled,
  logit_summary = logit_summary,
  case_study = case_study,
  chi_square = chi_square,
  schooling_benefit_missingness = benefit_missingness,
  notes = notes
)
class(out) <- c("emi_missingness_diagnostics", class(out))
out
}

enrolled_schooling_rows <- function(df) {
  if (!"enrolled" %in% names(df)) return(rep(TRUE, nrow(df)))
  value <- tolower(as.character(df$enrolled))
  value %in% c("yes", "1", "true", "enrolled")
}

summarize_missingness_by_variable <- function(df) {
  df <- as.data.frame(df, stringsAsFactors = FALSE)
  if (!length(names(df))) {
    return(data.frame(missing_var = character(), n_missing = integer(), pct_missing =
      ↪ numeric(), stringsAsFactors = FALSE))
  }
  data.frame(
    missing_var = names(df),
    n_missing = vapply(df, function(x) sum(is.na(x)), integer(1)),
    pct_missing = if (nrow(df)) vapply(df, function(x) mean(is.na(x)), numeric(1)) else
      ↪ NA_real_,
    stringsAsFactors = FALSE
  )
}

```

```

summarize_missingness_case_study <- function(df, vars) {
  df <- as.data.frame(df, stringsAsFactors = FALSE)
  if (!nrow(df)) return(data.frame())
  state_col <- intersect(c("state_0708", "state"), names(df))
  region_col <- intersect(c("region_0708", "region"), names(df))
  if (!length(state_col)) return(data.frame(note = "No state column available for
  ↪ Rajasthan/Southern case study.", stringsAsFactors = FALSE))

  keep <- grepl("rajasthan", as.character(df[[state_col[[1]]]]), ignore.case = TRUE)
  case_scope <- "Rajasthan"
  if (length(region_col)) {
    region_keep <- grepl("southern", as.character(df[[region_col[[1]]]]), ignore.case =
  ↪ TRUE)
    if (any(keep & region_keep, na.rm = TRUE)) {
      keep <- keep & region_keep
      case_scope <- "Rajasthan / Southern"
    }
  }
  vars_to_check <- intersect(c(vars$miss_vars_all, vars$miss_vars_enrolled), names(df))
  if (!length(vars_to_check) || !any(keep, na.rm = TRUE)) {
    return(data.frame(case_scope = case_scope, n_rows = sum(keep, na.rm = TRUE), note =
  ↪ "No matching rows or missingness variables available.", stringsAsFactors =
  ↪ FALSE))
  }
  case_df <- df[keep, vars_to_check, drop = FALSE]
  any_missing <- !stats::complete.cases(case_df)
  data.frame(
    case_scope = case_scope,
    n_rows = nrow(case_df),
    n_rows_with_any_missing = sum(any_missing),
    n_rows_with_partial_missing = sum(any_missing & rowSums(is.na(case_df)) <
  ↪ ncol(case_df)),
    n_missing_cells = sum(is.na(case_df)),
    n_observed_cells_in_rows_with_missing = sum(!is.na(case_df[any_missing, , drop =
  ↪ FALSE])),
    interpretation = "Current analog of the legacy Rajasthan/Southern View() check: rows
  ↪ with partial cost-variable missingness preserve the legacy concern that a child
  ↪ was not necessarily excluded wholesale when one cost field was missing.",
    stringsAsFactors = FALSE
  )
}

summarize_missingness_chisq <- function(df, probit_vars) {
  df <- as.data.frame(df, stringsAsFactors = FALSE)
  if (!nrow(df) || !length(probit_vars)) return(data.frame())
  probit_vars <- intersect(probit_vars, names(df))
  if (!length(probit_vars)) return(data.frame())
  any_na <- !stats::complete.cases(df[probit_vars])
  safe_test <- function(group_col, label) {
    if (!group_col %in% names(df)) {
      return(data.frame(test = label, status = "not_available", reason = paste("Missing
  ↪ column", group_col), stringsAsFactors = FALSE))
    }
  }
  group <- as.character(df[[group_col]])
  keep <- !is.na(group) & nzchar(group)

```

```

    if (length(unique(group[keep])) < 2L || length(unique(any_na[keep])) < 2L) {
      return(data.frame(test = label, status = "not_estimated", reason = "Insufficient
        ↪ variation for chi-square test.", stringsAsFactors = FALSE))
    }
    tab <- table(any_na = any_na[keep], group = group[keep])
    fit <- suppressWarnings(stats::chisq.test(tab))
    data.frame(
      test = label,
      status = "estimated",
      statistic = unname(fit$statistic),
      parameter = unname(fit$parameter),
      p.value = unname(fit$p.value),
      n = sum(tab),
      method = fit$method,
      stringsAsFactors = FALSE
    )
  }
  safe_bind_rows(list(
    safe_test("state_0708", "any_probit_model_na_by_state"),
    safe_test("region_0708", "any_probit_model_na_by_region")
  ))
}

summarize_missingness_regions <- function(df, probit_vars, vars, top_n = 20L) {
  empty <- data.frame()
  if (!"state_0708" %in% names(df)) {
    return(list(cost = empty, distance = empty, father_education = empty))
  }
  if (!length(probit_vars)) probit_vars <- names(df)
  temp <- df
  region_cols <- intersect(c("state_0708", "region_0708"), names(temp))
  # Legacy Chunk 8 inspected both state and region. Some cleaned selection
  # inputs no longer expose region_0708, so keep the diagnostic alive at the
  # state level instead of writing empty CSVs.
  if (!"region_0708" %in% region_cols) {
    temp$region_0708 <- "all_regions_available_in_state_only_input"
    region_cols <- c("state_0708", "region_0708")
  }

  temp$any_na_row <- !stats::complete.cases(temp[intersect(probit_vars, names(temp))])
  for (nm in c("dmean_num_ENROLLMENT_COST", "DIST_FROM_NEAREST_PRIMARY_CLASS",
    ↪ "father_educ")) {
    temp[[paste0("miss_", nm)]] <- if (nm %in% names(temp)) as.integer(is.na(temp[[nm]]))
    ↪ else NA_integer_
  }
  temp$is_urban <- if ("SECTOR" %in% names(temp)) as.integer(temp$SECTOR == "Urban") else
  ↪ NA_integer_
  temp$is_female <- if ("SEX" %in% names(temp)) as.integer(temp$SEX == "Female") else
  ↪ NA_integer_
  temp$is_hindu <- if ("RELIGION" %in% names(temp)) as.integer(temp$RELIGION == "Hindu")
  ↪ else NA_integer_
  temp$is_muslim <- if ("RELIGION" %in% names(temp)) as.integer(temp$RELIGION ==
    ↪ "Muslim") else NA_integer_

```



```

temp$is_st_sc_obc <- if ("SOCIAL_GROUP" %in% names(temp)) as.integer(temp$SOCIAL_GROUP
↪ %in% c("Scheduled Tribe", "Scheduled Caste", "Other Backward Class")) else
↪ NA_integer_

measure_cols <- c("any_na_row", "miss_dmean_num_ENROLLMENT_COST",
↪ "miss_DIST_FROM_NEAREST_PRIMARY_CLASS", "miss_father_educ", "is_urban", "is_female",
↪ "is_hindu", "is_muslim", "is_st_sc_obc")
grouped <- stats::aggregate(
  temp[measure_cols],
  temp[region_cols],
  function(x) mean(as.numeric(x), na.rm = TRUE)
)
n <- stats::aggregate(temp$any_na_row, temp[region_cols], length)
names(n)[names(n) == "x"] <- "n"
out <- merge(n, grouped, by = region_cols, all = TRUE)
names(out) <- sub("^any_na_row$", "pct_any_na", names(out))
out$region_diagnostic_level <- if ("region_0708" %in% names(df)) "state_region" else
↪ "state_only_fallback"
pct_cols <- setdiff(names(out), c(region_cols, "n", "region_diagnostic_level"))
out[pct_cols] <- lapply(out[pct_cols], function(x) round(100 * x, 2))
rank_one <- function(col) {
  if (!col %in% names(out)) return(empty)
  out[order(-out[[col]]), , drop = FALSE][seq_len(min(top_n, nrow(out))), , drop =
↪ FALSE]
}
list(
  cost = rank_one("miss_dmean_num_ENROLLMENT_COST"),
  distance = rank_one("miss_DIST_FROM_NEAREST_PRIMARY_CLASS"),
  father_education = rank_one("miss_father_educ")
)
}

missingness_correlation_matrix <- function(df, miss_vars, group_vars, cts_vars) {
  df <- as.data.frame(df, stringsAsFactors = FALSE)
  miss_vars <- intersect(miss_vars, names(df))
  if (!length(miss_vars) || !nrow(df)) return(matrix(numeric(), nrow = 0L, ncol = 0L))
  miss_df <- as.data.frame(lapply(df[miss_vars], function(x) as.integer(is.na(x))),
↪ check.names = FALSE)
  names(miss_df) <- paste0("miss_", names(miss_df))
  cts <- intersect(cts_vars, names(df))
  cts_df <- if (length(cts)) as.data.frame(lapply(df[cts], numeric_like), check.names =
↪ FALSE) else data.frame()
  grp <- intersect(group_vars, names(df))
  grp_df <- if (length(grp)) {
    grp_data <- as.data.frame(lapply(df[grp], function(x) {
      x <- as.character(x)
      x[is.na(x) | !nzchar(x)] <- "missing"
      factor(x)
    }), check.names = FALSE)
    # Legacy Chunk 8 used model.matrix-style expansions of sector, sex,
    # religion, social group, state, and related grouping variables before
    # computing missingness correlations. Small diagnostic subsets, especially
    # enrolled-only slices used in tests or regional filters, may contain only a
    # single observed level for one or more grouping variables. model.matrix()
    # cannot apply contrasts to one-level factors, so drop constant grouping

```



```

# variables while preserving all grouping variables with real variation.
varying <- vapply(grp_data, function(x) length(unique(as.character(x))) > 1L,
  ↪ logical(1))
grp_data <- grp_data[varying]
if (length(grp_data)) {
  stats::model.matrix(~ . - 1, data = grp_data)
} else {
  matrix(numeric(), nrow = nrow(df), ncol = 0L)
}
} else matrix(numeric(), nrow = nrow(df), ncol = 0L)
safe_pairwise_cor(cbind(miss_df, cts_df, grp_df))
}

```

```

binomial_fit_issues <- function(fit, warnings = character(), tolerance =
  ↪ sqrt(.Machine$double.eps)) {
  issues <- unique(as.character(warnings[nzchar(warnings)]))
  probabilities <- stats::fitted(fit)
  boundary_fit <- any(
    is.finite(probabilities) &
    (probabilities <= tolerance | probabilities >= 1 - tolerance)
  )
  if (boundary_fit) {
    issues <- c(issues, "fitted probabilities are numerically near 0 or 1")
  }
  if (!isTRUE(fit$converged)) {
    issues <- c(issues, "binomial model did not converge")
  }
  unique(issues)
}

```

```

#' check missing logit parallel
#'
#' Legacy Chunk 8 used mclapply/parLapply after defining the same one-logit-per-
#' missing-variable problem. Targets already parallelizes at the target level,
#' so this implementation keeps deterministic in-process lapply while preserving
#' the model specification and BH adjustment.
check_missing_logit_parallel <- function(df, miss_vars, covars, method_p = "BH") {
  df <- as.data.frame(df, stringsAsFactors = FALSE)
  miss_vars <- intersect(miss_vars, names(df))
  covars <- intersect(covars, names(df))
  if (!length(miss_vars) || !length(covars)) return(data.frame())
  rhs <- paste(covars, collapse = " + ")
  fit_one <- function(m) {
    y <- is.na(df[[m]])
    if (length(unique(y)) < 2L) return(data.frame())
    f <- stats::as.formula(paste0("is.na(`", m, "`) ~ ", rhs))
    tryCatch({
      fit_warnings <- character()
      fit <- withCallingHandlers(
        stats::glm(f, data = df, family = stats::binomial),
        warning = function(w) {
          fit_warnings <- c(fit_warnings, conditionMessage(w))
          invokeRestart("muffleWarning")
        }
      )
    })
  }

```

```

    )
    fit_issues <- binomial_fit_issues(fit, fit_warnings)
    pseudoR2 <- 1 - fit$deviance / fit$null.deviance
    out <- broom::tidy(fit)
    out$missing_var <- m
    out$pseudoR2 <- pseudoR2
    out$nobs <- stats::nobs(fit)
    out$status <- if (length(fit_issues)) "estimated_with_warning" else "estimated"
    out$reason <- if (length(fit_issues)) paste(fit_issues, collapse = "; ") else
    ↪ NA_character_
    out
  }, error = function(e) {
    data.frame(term = NA_character_, estimate = NA_real_, std.error = NA_real_,
    ↪ statistic = NA_real_, p.value = NA_real_, missing_var = m, pseudoR2 = NA_real_,
    ↪ nobs = length(y), status = "failed", reason = conditionMessage(e),
    ↪ stringsAsFactors = FALSE)
  })
}
out <- safe_bind_rows(lapply(miss_vars, fit_one))
adjust_missingness_pvalues_bh(out, method_p = method_p)
}

run_missingness_logits <- function(df, miss_vars, covars) {
  check_missing_logit_parallel(df, miss_vars, covars)
}

adjust_missingness_pvalues_bh <- function(df, method_p = "BH") {
  df <- as.data.frame(df, stringsAsFactors = FALSE)
  if (!nrow(df) || !"p.value" %in% names(df)) return(df)
  df$p_adj <- NA_real_
  idx <- !is.na(df$p.value) & df$term != "(Intercept)"
  df$p_adj[idx] <- stats::p.adjust(df$p.value[idx], method = method_p)
  df
}

summarize_missingness_logits <- function(df) {
  df <- as.data.frame(df, stringsAsFactors = FALSE)
  if (!nrow(df) || !all(c("missing_var", "p_adj", "pseudoR2") %in% names(df))) {
    return(data.frame(missing_var = character(), n_sig = integer(), pseudoR2 = numeric(),
    ↪ stringsAsFactors = FALSE))
  }
  split_df <- split(df, df$missing_var)
  out <- safe_bind_rows(lapply(split_df, function(x) {
    pmax <- suppressWarnings(max(x$pseudoR2, na.rm = TRUE))
    if (!is.finite(pmax)) pmax <- NA_real_
    data.frame(
      missing_var = x$missing_var[[1]],
      n_sig = sum(x$p_adj < 0.05 & x$term != "(Intercept)", na.rm = TRUE),
      pseudoR2 = pmax,
      stringsAsFactors = FALSE
    )
  })))
  out[order(-out$pseudoR2), , drop = FALSE]
}

```

```

missingness_correlation_pairs <- function(mat, top_n = 50L) {
  mat <- as.matrix(mat)
  if (!length(mat) || nrow(mat) < 2L || ncol(mat) < 2L) {
    return(data.frame(var1 = character(), var2 = character(), correlation = numeric(),
      ↪ abs_correlation = numeric(), stringsAsFactors = FALSE))
  }
  mat[lower.tri(mat, diag = TRUE)] <- NA_real_
  idx <- which(is.finite(mat), arr.ind = TRUE)
  if (!nrow(idx)) {
    return(data.frame(var1 = character(), var2 = character(), correlation = numeric(),
      ↪ abs_correlation = numeric(), stringsAsFactors = FALSE))
  }
  out <- data.frame(
    var1 = rownames(mat)[idx[, "row"]],
    var2 = colnames(mat)[idx[, "col"]],
    correlation = mat[idx],
    stringsAsFactors = FALSE
  )
  out$abs_correlation <- abs(out$correlation)
  out <- out[order(-out$abs_correlation), , drop = FALSE]
  utils::head(out, top_n)
}

save_missingness_correlation_heatmap <- function(mat, path, max_vars = 35L, title =
  ↪ "Missingness correlation matrix") {
  mat <- as.matrix(mat)
  dir.create(dirname(path), recursive = TRUE, showWarnings = FALSE)
  if (!length(mat) || nrow(mat) < 2L || ncol(mat) < 2L) {
    grDevices::png(path, width = 900, height = 500, res = 120)
    graphics::plot.new()
    graphics::text(0.5, 0.5, "No correlation matrix available")
    grDevices::dev.off()
    return(normalizePath(path, mustWork = FALSE))
  }
  score <- apply(abs(mat), 1L, function(x) max(x[is.finite(x) & x < 1], na.rm = TRUE))
  score[!is.finite(score)] <- 0
  keep <- names(sort(score, decreasing = TRUE))[seq_len(min(max_vars, length(score)))]
  mat <- mat[keep, keep, drop = FALSE]
  grDevices::png(path, width = 1400, height = 1200, res = 140)
  old <- graphics::par(no.readonly = TRUE)
  on.exit({ graphics::par(old); grDevices::dev.off() }, add = TRUE)
  graphics::par(mar = c(11, 11, 4, 2))
  graphics::image(
    x = seq_len(ncol(mat)),
    y = seq_len(nrow(mat)),
    z = t(mat[nrow(mat):1, , drop = FALSE]),
    axes = FALSE,
    xlab = "",
    ylab = "",
    main = title,
    zlim = c(-1, 1)
  )
  graphics::axis(1, at = seq_len(ncol(mat)), labels = colnames(mat), las = 2, cex.axis =
  ↪ 0.55)

```

```

    graphics::axis(2, at = seq_len(nrow(mat)), labels = rev(rownames(mat)), las = 2,
    ↪ cex.axis = 0.55)
    graphics::box()
    normalizePath(path, mustWork = FALSE)
  }

save_missingness_logit_plot <- function(summary_df, path, title = "Structured missingness
    ↪ logit screen", top_n = 25L) {
  dir.create(dirname(path), recursive = TRUE, showWarnings = FALSE)
  df <- as.data.frame(summary_df, stringsAsFactors = FALSE)
  grDevices::png(path, width = 1100, height = 850, res = 130)
  old <- graphics::par(no.readonly = TRUE)
  on.exit({ graphics::par(old); grDevices::dev.off() }, add = TRUE)
  if (!nrow(df) || !all(c("missing_var", "pseudoR2") %in% names(df))) {
    graphics::plot.new()
    graphics::text(0.5, 0.5, "No missingness-logit summary available")
    return(normalizePath(path, mustWork = FALSE))
  }
  df <- df[is.finite(df$pseudoR2), , drop = FALSE]
  if (!nrow(df)) {
    graphics::plot.new()
    graphics::text(0.5, 0.5, "No finite pseudo-R2 values available")
    return(normalizePath(path, mustWork = FALSE))
  }
  df <- df[order(df$pseudoR2, decreasing = TRUE), , drop = FALSE]
  df <- utils::head(df, top_n)
  graphics::par(mar = c(5, 13, 4, 2))
  graphics::barplot(
    rev(df$pseudoR2),
    names.arg = rev(df$missing_var),
    horiz = TRUE,
    las = 1,
    xlab = "Pseudo-R2 (predictability of missingness)",
    main = title
  )
  normalizePath(path, mustWork = FALSE)
}

save_missingness_diagnostics <- function(diagnostics, dir =
    ↪ "outputs/diagnostics/extended/missingness") {
  dir.create(dir, recursive = TRUE, showWarnings = FALSE)
  if (!inherits(diagnostics, "emi_missingness_diagnostics")) diagnostics <-
    ↪ list(missing_counts = as.data.frame(diagnostics))
  paths <- c(
    missing_counts = write_diagnostic_csv(diagnostics$missing_counts %||% data.frame(),
    ↪ file.path(dir, "missingness_counts.csv")),
    regional_cost = write_diagnostic_csv(diagnostics$regional_cost %||% data.frame(),
    ↪ file.path(dir, "regional_missingness_cost.csv")),
    regional_distance = write_diagnostic_csv(diagnostics$regional_distance %||%
    ↪ data.frame(), file.path(dir, "regional_missingness_distance.csv")),
    regional_father = write_diagnostic_csv(diagnostics$regional_father_education %||%
    ↪ data.frame(), file.path(dir, "regional_missingness_father_education.csv")),
    logit_all = write_diagnostic_csv(diagnostics$logit_all %||% data.frame(),
    ↪ file.path(dir, "missingness_logits_all.csv")),
  )
}

```

```

logit_enrolled = write_diagnostic_csv(diagnostics$logit_enrolled %||% data.frame(),
  ↪ file.path(dir, "missingness_logits_enrolled.csv")),
logit_summary = write_diagnostic_csv(diagnostics$logit_summary %||% data.frame(),
  ↪ file.path(dir, "missingness_logit_summary.csv")),
case_study = write_diagnostic_csv(diagnostics$case_study %||% data.frame(),
  ↪ file.path(dir, "missingness_rajasthan_southern_case_study.csv")),
chi_square = write_diagnostic_csv(diagnostics$chi_square %||% data.frame(),
  ↪ file.path(dir, "missingness_chi_square_tests.csv")),
schooling_benefits = write_diagnostic_csv(
  diagnostics$schooling_benefit_missingness %||% data.frame(),
  file.path(dir, "missingness_schooling_benefits_enrolled.csv")
),
notes = write_diagnostic_csv(diagnostics$notes %||% data.frame(), file.path(dir,
  ↪ "missingness_legacy_notes.csv"))
)
if (length(diagnostics$corr_all)) {
  paths <- c(
    paths,
    corr_all = write_diagnostic_matrix(diagnostics$corr_all, file.path(dir,
      ↪ "missingness_correlation_all.csv")),
    corr_all_pairs =
      ↪ write_diagnostic_csv(missingness_correlation_pairs(diagnostics$corr_all),
      ↪ file.path(dir, "missingness_correlation_all_top_pairs.csv")),
    corr_all_heatmap = save_missingness_correlation_heatmap(diagnostics$corr_all,
      ↪ file.path(dir, "missingness_correlation_all.png"), title = "Missingness
      ↪ correlations: all observations")
  )
}
if (length(diagnostics$corr_enrolled)) {
  paths <- c(
    paths,
    corr_enrolled = write_diagnostic_matrix(diagnostics$corr_enrolled, file.path(dir,
      ↪ "missingness_correlation_enrolled.csv")),
    corr_enrolled_pairs =
      ↪ write_diagnostic_csv(missingness_correlation_pairs(diagnostics$corr_enrolled),
      ↪ file.path(dir, "missingness_correlation_enrolled_top_pairs.csv")),
    corr_enrolled_heatmap =
      ↪ save_missingness_correlation_heatmap(diagnostics$corr_enrolled, file.path(dir,
      ↪ "missingness_correlation_enrolled.png"), title = "Missingness correlations:
      ↪ enrolled observations")
  )
}
if (nrow(diagnostics$logit_summary %||% data.frame())) {
  paths <- c(
    paths,
    logit_pseudo_r2_plot = save_missingness_logit_plot(
      diagnostics$logit_summary,
      file.path(dir, "missingness_logit_pseudo_r2.png"),
      title = "How structured is missingness?"
    )
  )
}
output_manifest(paths)
}

```

Survey-weighted enrollment probit

Source: R/selection/estimate_selection_probit.R

```
selection_probit_variables <- function(selection_data) {
  controls <- c("AGE", "age", "SEX", "HH_SIZE", "RELIGION", "SOCIAL_GROUP", "SECTOR")
  exclusion_all_kids <- c("DIST_FROM_NEAREST_PRIMARY_CLASS", "father_educ")
  exclusion_district <- c(
    "dmean_num_IS_EDU_FREE", "dmean_num_TUTION_FEE_WAIVED",
    "dmean_num_RECD_SCHOLARSHIP_STIPEND", "dmean_num_RECD_TXT_BOOKS",
    "dmean_num_RECD_STATIONERY", "dmean_num_MID_DAY_MEAL_ETC_RECD",
    "dmean_num_ENROLLMENT_COST"
  )
  intersect(c(controls, exclusion_all_kids, exclusion_district), names(selection_data))
}

#' project the selection sample to the model's actual dependency surface
#'
#' Keeping unrelated Block-5 schooling attributes in the fitted survey design makes
#' the serialized model change whenever those attributes change, even though they
#' are absent from the probit formula. A narrow model frame lets {targets} skip the
#' expensive AME branch when upstream changes do not affect estimation.
project_selection_model_data <- function(selection_data) {
  selection_data <- safe_df(selection_data)
  design <- c("enrolled", "FSU_SL_NO", "weight", "STATE", "STRATUM", "SUB_STRATUM_NO")
  keep <- unique(c(design, selection_probit_variables(selection_data)))
  selection_data[intersect(keep, names(selection_data))]
}

#' estimate selection probit
#'
estimate_selection_probit <- function(selection_data, cfg) {
  if (!"enrolled" %in% names(selection_data) || all(is.na(selection_data$enrolled))) {
    return(list(status = "out_of_active_pipeline", reason = "No enrolled variable."))
  }
  covars <- selection_probit_variables(selection_data)
  if (!length(covars)) {
    return(list(status = "out_of_active_pipeline", reason = "No probit covariates."))
  }
  f_probit <- stats::reformulate(covars, response = "enrolled")
  if (identical(cfg$mode, "final") && requireNamespace("survey", quietly = TRUE)) {
    old_options <- options(survey.lonely.psu = "average")
    on.exit(options(old_options), add = TRUE)

    design <- build_survey_design_selection(selection_data)
    if (!is.null(design)) {
      out <- fit_selection_probit(design, f_probit)
      return(stabilize_selection_model_formula(out, f_probit))
    }
  }
  out <- stats::glm(f_probit, data = selection_data, family = stats::binomial(link =
    ↪ "probit"))
  stabilize_selection_model_formula(out, f_probit)
}

#' store a durable selection-model formula
```

```

#'
#' Programmatically fitted models otherwise retain a call to the local symbol
#' `f_probit`. Packages which reconstruct model data from the saved call cannot
#' resolve that symbol after the model is serialized by targets. Embed the
#' formula object in the call and retain the existing audit attribute.
stabilize_selection_model_formula <- function(model, formula) {
  if (!inherits(formula, "formula")) {
    stop("Selection model formula must inherit from formula.", call. = FALSE)
  }
  if (!is.null(model$call)) model$call$formula <- formula
  attr(model, "selection_probit_formula") <- formula
  model
}

#' build survey design selection
#'
build_survey_design_selection <- function(selection_df) {
  psu <- first_col(selection_df, c("FSU_SL_NO", "fsu", "PSU", "psu"))
  weight <- first_col(selection_df, c("weight", "WEIGHT", "Multiplier", "multiplier"))
  strata_cols <- intersect(c("STATE", "STRATUM", "SUB_STRATUM_NO"), names(selection_df))
  if (is.null(psu) || is.null(weight) || !length(strata_cols)) return(NULL)
  selection_df$.survey_strata <- interaction(selection_df[strata_cols], drop = TRUE)
  survey::svydesign(
    ids = stats::as.formula(paste0("~", psu)),
    strata = ~.survey_strata,
    weights = stats::as.formula(paste0("~", weight)),
    data = selection_df,
    nest = TRUE
  )
}

#' fit selection probit
#'
fit_selection_probit <- function(selection_design, f_probit) {
  survey::svyglm(f_probit, design = selection_design, family = stats::quasibinomial(link
    ↵ = "probit"))
}

```

Average marginal effects via automatic differentiation

Source: R/selection/compute_average_marginal_effects.R

```

# Use automatic differentiation for SE delta-method gradients, as in the legacy Rmd.

#' compute average marginal effects
#'
compute_average_marginal_effects <- function(selection_model, cfg = list()) {
  if (!inherits(selection_model, "glm")) {
    return(ame_out_of_pipeline(
      selection_model$status %||% "out_of_active_pipeline",
      selection_model$reason %||% "Selection model not estimated in smoke mode"
    ))
  }
}

# `selection_model` is often a survey::svyglm object loaded from the targets
# store. Survey lonely-PSU handling is an R option, not a durable model

```

```

# attribute, so it may not be set in the downstream AME target even though it
# was set when the model was estimated. Set it explicitly around all AME
# computations to make reruns deterministic and warning/error-free.
old_lonely_psu <- getOption("survey.lonely.psu")
old_domain_lonely <- getOption("survey.adjust.domain.lonely")
options(survey.lonely.psu = "average", survey.adjust.domain.lonely = TRUE)
on.exit({
  if (is.null(old_lonely_psu)) {
    options(survey.lonely.psu = NULL)
  } else {
    options(survey.lonely.psu = old_lonely_psu)
  }
  if (is.null(old_domain_lonely)) {
    options(survey.adjust.domain.lonely = NULL)
  } else {
    options(survey.adjust.domain.lonely = old_domain_lonely)
  }
}, add = TRUE)

#' run expression while muffling survey lonely-PSU warnings
#'
#' survey emits lonely-PSU warnings even when survey.lonely.psu is set to
#' average and the adjustment is intentional. Muffle only those handled
#' warnings inside the AME target so strict final builds remain warning-clean.
muffle_survey_lonely_psu_warnings <- function(expr) {
  withCallingHandlers(
    expr,
    warning = function(w) {
      msg <- conditionMessage(w)
      lonely_psu_warning <- grepl(
        "has only one PSU at stage",
        msg,
        fixed = TRUE
      )
      if (lonely_psu_warning) invokeRestart("muffleWarning")
    }
  )
}

if (!isTRUE(cfg$run_full_ame)) {
  return(format_ame_results(data.frame(
    term = names(stats::coef(selection_model)),
    estimate = unname(stats::coef(selection_model)),
    std.error = NA_real_,
    statistic = NA_real_,
    p.value = NA_real_,
    s.value = NA_real_,
    conf.low = NA_real_,
    conf.high = NA_real_,
    method = "coefficient_fallback",
    status = "estimated",
    reason = "Draft config uses coefficient fallback instead of full AME computation"
  )))
}

```



```

if (!requireNamespace("marginaleffects", quietly = TRUE)) {
  return(ame_out_of_pipeline(
    "out_of_active_pipeline",
    "Package marginaleffects is required for full AME computation"
  ))
}

out <- muffle_survey_lonely_psu_warnings(
  tryCatch(
    compute_ames_autodiff(selection_model),
    error = function(e) {
      fallback <- compute_ames_probit_analytic(selection_model)
      fallback$method <- "delta_method_analytic_probit"
      fallback$reason <- NA_character_
      fallback
    }
  )
)
format_ame_results(out)
}

ame_out_of_pipeline <- function(status, reason) {
  data.frame(
    term = NA_character_,
    estimate = NA_real_,
    std.error = NA_real_,
    statistic = NA_real_,
    p.value = NA_real_,
    s.value = NA_real_,
    conf.low = NA_real_,
    conf.high = NA_real_,
    method = "not_run",
    status = status,
    reason = reason
  )
}

#' extract model-frame data and AME weights
#'
#' marginaleffects warns, correctly, when weighted/survey models are summarized
#' without an explicit `wts` argument. Build a newdata frame from the exact
#' model estimation sample and attach a synthetic weight column whenever the
#' fitted model exposes usable weights.
ame_model_weight <- function(model_data, model_weights = NULL) {
  weight_cols <- intersect(c("weight", "WEIGHT", "Multiplier", "multiplier",
    ↪ "(weights)"), names(model_data))
  if (length(weight_cols)) {
    return(suppressWarnings(as.numeric(model_data[[weight_cols[[1]]]])))
  }

  if (!is.null(model_weights) && length(model_weights) == nrow(model_data)) {
    return(suppressWarnings(as.numeric(model_weights)))
  }
}

```

```

  NULL
}

#' build marginaleffects newdata and weights
#'
#' @return List with `data`, `wts`, and `has_explicit_weights`.
ame_model_data_and_weights <- function(model) {
  model_data <- as.data.frame(stats::model.frame(model))
  model_weights <- tryCatch(stats::weights(model), error = function(e) NULL)
  weight <- ame_model_weight(model_data, model_weights)

  if (!is.null(weight) && length(weight) == nrow(model_data) && any(is.finite(weight) &
    ↪ weight != 1)) {
    model_data$.ame_weight <- weight
    return(list(data = model_data, wts = ".ame_weight", has_explicit_weights = TRUE))
  }

  list(data = model_data, wts = FALSE, has_explicit_weights = FALSE)
}

#' run marginaleffects while muffling only a redundant explicit-weight warning
#'
run_avg_slopes <- function(model, model_data, wts, numderiv = NULL) {
  args <- list(
    model = model,
    newdata = model_data,
    wts = wts,
    vcov = TRUE,
    type = "response"
  )
  if (!is.null(numderiv)) args$numderiv <- numderiv

  withCallingHandlers(
    do.call(marginaleffects::avg_slopes, args),
    warning = function(w) {
      msg <- conditionMessage(w)
      explicit_weight_warning <- grepl(
        "normally good practice to specify weights using the `wts` argument",
        msg,
        fixed = TRUE
      )
      if (explicit_weight_warning && !identical(wts, FALSE)) {
        invokeRestart("muffleWarning")
      }
    }
  )
}

#' compute ames autodiff
#'
compute_ames_autodiff <- function(model) {
  # Use the exact estimation sample retained by glm/svyglm. Passing the full
  # pre-model selection_data can have extra rows omitted during model fitting,
  # which makes marginaleffects recycle weights/covariates silently.
  amed <- ame_model_data_and_weights(model)

```

```

run_avg_slopes(model, amed$data, amed$wts)
}

#' compute analytic probit AMEs
#'
#' @return Data frame of observed-data average marginal effects.
compute_ames_probit_analytic <- function(model) {
  coefs <- stats::coef(model)
  coef_table <- as.data.frame(summary(model)$coefficients)
  terms <- setdiff(names(coefs), "(Intercept)")
  if (!length(terms)) return(ame_out_of_pipeline("out_of_active_pipeline", "Selection
    ↪ model has no non-intercept terms.))

  newdata <- stats::model.frame(model)
  weights <- if ("weight" %in% names(newdata)) num(newdata$weight) else rep(1,
    ↪ nrow(newdata))
  eta <- stats::predict(model, type = "link")
  mean_phi <- wmean(stats::dnorm(eta), weights)
  factor_levels <- model$xlevels %||% list()

  rows <- lapply(terms, function(term) {
    factor_var <- names(factor_levels)[vapply(names(factor_levels), function(v)
    ↪ startsWith(term, v), logical(1))]]
    estimate <- NA_real_
    if (length(factor_var)) {
      var <- factor_var[[which.max(nchar(factor_var))]]
      level <- sub(paste0("^", var), "", term)
      if (level %in% factor_levels[[var]]) {
        hi <- newdata
        lo <- newdata
        hi[[var]] <- factor(level, levels = factor_levels[[var]])
        lo[[var]] <- factor(factor_levels[[var]][[1]], levels = factor_levels[[var]])
        estimate <- wmean(stats::predict(model, newdata = hi, type = "response") -
    ↪ stats::predict(model, newdata = lo, type = "response"), weights)
      }
    } else {
      estimate <- unname(coefs[[term]]) * mean_phi
    }
    p_col <- intersect(c("Pr(>|z|)", "Pr(>|t|)"), names(coef_table))
    p <- if (term %in% rownames(coef_table) && length(p_col)) coef_table[term,
    ↪ p_col[[1]]] else NA_real_
    se_col <- intersect(c("Std. Error", "std.error"), names(coef_table))
    se_beta <- if (term %in% rownames(coef_table) && length(se_col)) coef_table[term,
    ↪ se_col[[1]]] else NA_real_
    se <- if (is.finite(se_beta)) abs(se_beta * mean_phi) else NA_real_
    z <- if (is.finite(se) && se > 0) estimate / se else NA_real_
    p_out <- if (is.finite(z)) 2 * stats::pnorm(abs(z), lower.tail = FALSE) else p
    data.frame(
      term = term,
      estimate = estimate,
      std.error = se,
      statistic = z,
      p.value = p_out,
      s.value = if (is.finite(p_out) && p_out > 0) -log2(p_out) else NA_real_,

```

```

      conf.low = if (is.finite(se)) estimate - 1.96 * se else NA_real_,
      conf.high = if (is.finite(se)) estimate + 1.96 * se else NA_real_,
      method = "delta_method_analytic_probit",
      status = "estimated",
      reason = NA_character_,
      stringsAsFactors = FALSE
    )
  })
  do.call(rbind, rows)
}

is_marginaeffects_object <- function(x) {
  inherits(x, c("marginaeffects", "comparisons", "avg_comparisons", "slopes",
    ↪ "avg_slopes", "predictions"))
}

copy_first_ame_column <- function(out, target, candidates) {
  if (!target %in% names(out)) out[[target]] <- NA
  for (candidate in candidates) {
    if (!candidate %in% names(out)) next
    missing_target <- is.na(out[[target]])
    if (any(missing_target)) out[[target]][missing_target] <-
      ↪ out[[candidate]][missing_target]
  }
  out
}

normalize_ame_columns <- function(out) {
  # marginaeffects has used both dotted and snake_case names across versions
  # and model classes. Normalize here before selecting the public/audit schema
  # so uncertainty columns are not silently dropped downstream.
  aliases <- list(
    std.error = c("std_error", "std.err", "Std. Error"),
    statistic = c("z", "z.value", "z_value", "t", "t.value", "t_value"),
    p.value = c("p_value", "p", "Pr(>|z|)", "Pr(>|t|)"),
    s.value = c("s_value"),
    conf.low = c("conf_low", "conf.low", "2.5 %"),
    conf.high = c("conf_high", "conf.high", "97.5 %")
  )
  for (target in names(aliases)) {
    out <- copy_first_ame_column(out, target, aliases[[target]])
  }
  out
}

ame_label_lookup <- function() {
  data.frame(
    term = c(
      "AGE", "SEX", "HH_SIZE",
      "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION",
      "SOCIAL_GROUP", "SOCIAL_GROUP", "SOCIAL_GROUP",
      "SECTOR",
      "DIST_FROM_NEAREST_PRIMARY_CLASS", "DIST_FROM_NEAREST_PRIMARY_CLASS",
      ↪ "DIST_FROM_NEAREST_PRIMARY_CLASS", "DIST_FROM_NEAREST_PRIMARY_CLASS",

```

```

"father_educ", "father_educ", "father_educ", "father_educ", "father_educ",
  ↪ "father_educ", "father_educ",
"dmean_num_IS_EDU_FREE", "dmean_num_TUTION_FEE_WAIVED",
  ↪ "dmean_num_RECD_SCHOLARSHIP_STIPEND",
"dmean_num_RECD_TXT_BOOKS", "dmean_num_RECD_STATIONERY",
  ↪ "dmean_num_MID_DAY_MEAL_ETC_RECD", "dmean_num_ENROLLMENT_COST"
),
contrast = c(
  "dY/dX", "Female - Male", "dY/dX",
  "Muslim - Hindu", "Christian - Hindu", "Sikh - Hindu", "Jain - Hindu", "Buddhist -
  ↪ Hindu", "Zoroastrian - Hindu", "Other - Hindu",
  "Scheduled Tribe - Other", "Scheduled Caste - Other", "Other Backward Class -
  ↪ Other",
  "Urban - Rural",
  "1km <= d <2kms - d<1km", "2kms<= d <3kms - d<1km", "3kms <= d <5kms - d<1km",
  ↪ "d>=5kms - d<1km",
  "Literate, no school - Illiterate", "Literate, school < primary - Illiterate",
  ↪ "Primary - Illiterate", "Upper primary - Illiterate", "Secondary - Illiterate",
  ↪ "Higher secondary - Illiterate", "Postsecondary+ - Illiterate",
  "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX"
),
Term = c(
  "Age (years)", "Female (ref: Male)", "Household size",
  "Religion: Muslim (ref: Hindu)", "Religion: Christian", "Religion: Sikh",
  ↪ "Religion: Jain", "Religion: Buddhist", "Religion: Zoroastrian", "Religion:
  ↪ Other",
  "Social group: Scheduled Tribe (ref: Other)", "Social group: Scheduled Caste",
  ↪ "Social group: Other Backward Class",
  "Urban (ref: Rural)",
  "Distance 1-2km (ref: <1km)", "Distance 2-3km", "Distance 3-5km", "Distance > 5km",
  "Father's educ.: Literate, no school (ref: Illiterate)", "Father's educ.: Literate,
  ↪ school < primary", "Father's educ.: Primary", "Father's educ.: Upper primary",
  ↪ "Father's educ.: Secondary", "Father's educ.: Higher secondary", "Father's
  ↪ educ.: Postsecondary+",
  "Educ. free available (ref: No)", "Tuition waiver received", "Scholarship/Stipend
  ↪ received", "Textbook(s) received", "Stationery received", "Mid-day meal, etc.
  ↪ received", "Enrollment cost (Rs.)"
),
stringsAsFactors = FALSE
)
}

ame_label_order <- function() ame_label_lookup()$Term

ame_modelsummary_label <- function(x) {
  x <- as.character(x)
  x <- gsub("\u00d7", ":", x, fixed = TRUE)
  x <- gsub("\s+:\s+", ":", x, perl = TRUE)
  x <- gsub("\\(ref:\s*", "(ref: ", x, perl = TRUE)
  x
}

infer_ame_contrast <- function(out) {
  if ("contrast" %in% names(out)) return(as.character(out$contrast))
  contrast <- rep("dY/dX", nrow(out))

```

```

    if (!"term" %in% names(out)) return(contrast)
    if ("comparison" %in% names(out)) contrast <- as.character(out$comparison)
    contrast[is.na(contrast) | !nzchar(contrast)] <- "dY/dX"
    contrast
  }

ame_factor_base_terms <- function() {
  unique(ame_label_lookup()$term[ame_label_lookup()$contrast != "dY/dX"])
}

normalize_encoded_factor_ame_terms <- function(out) {
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  if (!"term" %in% names(out)) return(out)
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(out)

  lookup <- ame_label_lookup()
  factor_terms <- ame_factor_base_terms()
  factor_terms <- factor_terms[order(nchar(factor_terms), decreasing = TRUE)]

  for (i in seq_len(nrow(out))) {
    current_term <- as.character(out$term[[i]])
    current_contrast <- as.character(out$contrast[[i]])
    if (!nzchar(current_term) || (!is.na(current_contrast) && current_contrast !=
      ↪ "dY/dX")) next

    direct <- lookup$term == current_term & lookup$contrast == current_contrast
    if (any(direct)) next

    for (base_term in factor_terms) {
      if (!startsWith(current_term, base_term)) next
      level <- trimws(sub(paste0("^", base_term), "", current_term))
      if (!nzchar(level)) next
      candidates <- lookup[lookup$term == base_term, , drop = FALSE]
      contrast_hit <- candidates$contrast[startsWith(candidates$contrast, paste0(level, "
      ↪ - "))]
      if (length(contrast_hit)) {
        out$term[[i]] <- base_term
        out$contrast[[i]] <- contrast_hit[[1]]
        break
      }
    }
  }
  out
}

attach_ame_labels <- function(out) {
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  if (!"term" %in% names(out) && "variable" %in% names(out)) out$term <- out$variable
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(out)
  out <- normalize_encoded_factor_ame_terms(out)
  lookup <- ame_label_lookup()
  labeled <- merge(out, lookup, by = c("term", "contrast"), all.x = TRUE, sort = FALSE)
  labeled$Term[is.na(labeled$Term)] <- labeled$term[is.na(labeled$Term)]
  labeled$Term <- factor(labeled$Term, levels = unique(c(ame_label_order(),
  ↪ labeled$Term)))

```

```

labeled <- labeled[order(labeled$Term), , drop = FALSE]
labeled$Term <- as.character(labeled$Term)
labeled
}

modelsummary_margineffects_object <- function(out, native_ame) {
  if (!is_margineffects_object(native_ame)) return(NULL)
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  required <- c("Term", "estimate", "std.error", "statistic", "p.value", "conf.low",
    ↪ "conf.high")
  if (!all(required %in% names(out))) return(native_ame)

  ms <- out[, required, drop = FALSE]
  ms$term <- ame_modelsummary_label(ms$Term)
  ms$contrast <- ""
  ms$Term <- NULL
  # Preserve the margineffects class and non-structural attributes so
  # modelsummary can use its native margineffects/tidy pathway. We still
  # hand modelsummary the labeled, ordered rows that are validated for the
  # public table, because passing the raw object caused modelsummary to reorder
  # contrasts and display labels against the wrong estimates.
  native_attributes <- attributes(native_ame)
  structural_attributes <- c("names", "row.names", "class")
  for (nm in setdiff(names(native_attributes), structural_attributes)) {
    attr(ms, nm) <- native_attributes[[nm]]
  }
  class(ms) <- unique(c(class(native_ame), class(ms)))
  ms
}

#' format ame results
#'
format_ame_results <- function(ame_results) {
  native_ame <- ame_results
  out <- tibble::as_tibble(ame_results)
  if (!"term" %in% names(out) && "variable" %in% names(out)) out$term <- out$variable
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(as.data.frame(out))
  out <- normalize_ame_columns(out)
  if (!"method" %in% names(out)) out$method <- "autodiff"
  if (!"status" %in% names(out)) out$status <- "estimated"
  if (!"reason" %in% names(out)) out$reason <- NA_character_
  labeled_terms <- unique(ame_label_lookup()$term)
  use_labeled_schema <- any(out$term %in% labeled_terms) ||
    any(grepl("^dmean_num_", out$term %||% character(), perl = TRUE)) ||
    any((out$contrast %||% "dY/dX") != "dY/dX", na.rm = TRUE)

  if (isTRUE(use_labeled_schema)) {
    out <- attach_ame_labels(out)
    required <- c(
      "Term", "term", "contrast", "estimate", "std.error", "statistic", "p.value",
      ↪ "s.value",
      "conf.low", "conf.high", "method", "status", "reason"
    )
  } else {

```

```

    required <- c(
      "term", "estimate", "std.error", "statistic", "p.value",
      "s.value", "conf.low", "conf.high", "method", "status", "reason"
    )
  }

  for (nm in setdiff(required, names(out))) out[[nm]] <- NA
  if ("p.value" %in% names(out) && "s.value" %in% names(out)) {
    missing_s <- is.na(out$s.value) & is.finite(out$p.value) & out$p.value > 0
    out$s.value[missing_s] <- -log2(out$p.value[missing_s])
  }
  out <- out[, required, drop = FALSE]
  if (is_marginaffects_object(native_ame)) {
    attr(out, "marginaffects_object") <- modelsummary_marginaffects_object(out,
      ↪ native_ame)
  }
  out
}

```

District harmonization map and many-to-many QA

Source: R/districts/build_district_join_map.R

```

#' Required reviewed district-crosswalk columns
#'
#' @return The name columns needed to attach 2001, 2007-08, 2017-18, and 2020
#' source data to one harmonized district row.
district_join_map_required_columns <- function() {
  c("state_01", "district_01", "state_07", "district_07", "state_17", "district_17",
    ↪ "state_20", "district_20")
}

#' Prepare the reviewed district join map
#'
#' The active pipeline does not infer the harmonization map at build time. It
#' consumes the reviewed metadata crosswalk, gives each row a stable internal
#' identifier, and exposes empty diagnostic attributes until source attachment
#' produces a row-level match ledger in the planned district-matching rewrite.
prepare_district_join_map <- function(crosswalk) {
  out <- safe_df(crosswalk)
  missing <- setdiff(district_join_map_required_columns(), names(out))
  if (length(missing)) {
    stop(
      "District harmonization crosswalk is missing required columns: ",
      paste(missing, collapse = ", "),
      call. = FALSE
    )
  }
  if (!nrow(out)) stop("District harmonization crosswalk contains no rows.", call. =
    ↪ FALSE)

  out$tracker_row <- seq_len(nrow(out))
  out$source <- "harmonization_crosswalk"
  out$match_status <- "reviewed_crosswalk_row"
  out$possible_false_positive <- FALSE
  out$many_to_many <- FALSE

```



```

  attr(out, "unmatched_rows") <- out[0, , drop = FALSE]
  attr(out, "possible_false_positives") <- out[0, , drop = FALSE]
  attr(out, "many_to_many_cases") <- out[0, , drop = FALSE]
  out
}

#' Flag many-to-many matches in a row-level match ledger
#'
#' This helper is retained for matching diagnostics. It is not used to infer the
#' reviewed crosswalk itself.
flag_many_to_many_matches <- function(join_map, source_col = NULL, tracker_col = NULL) {
  join_map <- safe_df(join_map)
  if (!nrow(join_map)) {
    join_map$many_to_many <- logical()
    join_map$many_to_many_type <- character()
    return(join_map)
  }
  source_col <- source_col %||% first_col(join_map, c(".source_row", "source_row",
    ↪ "source_key", "source_id"))
  tracker_col <- tracker_col %||% first_col(join_map, c(".tracker_row", "tracker_row",
    ↪ "tracker_key", "district_panel_id"))
  if (is.null(source_col) || is.null(tracker_col)) {
    join_map$many_to_many <- FALSE
    join_map$many_to_many_type <- "not_evaluable_missing_keys"
    return(join_map)
  }
  source_key <- canon(join_map[[source_col]])
  tracker_key <- canon(join_map[[tracker_col]])
  source_n <- ave(rep(1L, nrow(join_map)), source_key, FUN = length)
  tracker_n <- ave(rep(1L, nrow(join_map)), tracker_key, FUN = length)
  source_dup <- source_n > 1L & !is.na(source_key) & nzchar(source_key)
  tracker_dup <- tracker_n > 1L & !is.na(tracker_key) & nzchar(tracker_key)
  join_map$n_source_matches <- as.integer(source_n)
  join_map$n_tracker_matches <- as.integer(tracker_n)
  join_map$many_to_many <- source_dup & tracker_dup
  join_map$many_to_many_type <- ifelse(
    source_dup & tracker_dup, "many_to_many",
    ifelse(source_dup, "one_source_to_many_tracker",
      ifelse(tracker_dup, "many_source_to_one_tracker", "one_to_one")))
  )
  join_map
}

```

Contiguity-based spatial weights

Source: R/diagnostics/diagnose_spatial_weights.R

```

#' build spatial weights
#'
#' Build the rook-contiguity spatial weights used by the legacy Rmd.
#'
#' The legacy chunk first removed rows with missing analysis values, then called
#' `poly2nb(queen = FALSE)`, `nb2mat(style = "B")`, and `nb2listw(style = "W")`.
#' The target stores all three objects so downstream diagnostics can reproduce
#' both the binary adjacency matrix and the row-standardized listw object.
#'

```

```

#' @return A list containing neighbor, matrix, and listw objects, or an explicit
#' inactive status when geometry is unavailable.
build_spatial_weights <- function(district_panel, cfg) {
  if (!inherits(district_panel, "sf")) {
    return(list(status = "out_of_active_pipeline", reason = "Requires sf geometry."))
  }
  need_pkg("spdep", "spatial weights")
  need_pkg("sf", "spatial weights")

  row_index <- spatial_weight_final_panel_rows(district_panel)
  coverage <- length(row_index) / max(nrow(district_panel), 1L)
  if (!is.finite(coverage) || coverage < 0.75) {
    return(list(
      status = "out_of_active_pipeline",
      reason = paste0(
        "Requires validated non-empty geometry for at least 75% of district-panel rows;
        ↪ current coverage is ",
        round(100 * coverage, 1), "%."
      ))
    )
  }

  weights <- build_spatial_weights_for_rows(district_panel, row_index, queen = FALSE,
  ↪ snap = spatial_snap_setting(cfg))
  weights$coverage <- coverage
  weights$panel_scope <- "current_final_matched_panel_non_empty_geometry"
  weights
}

#' rows used by final-panel spatial diagnostics
#'
#' Spatial diagnostics operate on the active `district_panel` target, not on the
#' legacy exploratory geometry object. The final-panel scope is therefore all
#' rows in the current matched panel with non-empty geometry.
spatial_weight_final_panel_rows <- function(district_panel) {
  if (!inherits(district_panel, "sf")) return(integer())
  geom <- sf::st_geometry(district_panel)
  which(!sf::st_is_empty(geom))
}

spatial_snap_setting <- function(cfg) {
  cfg <- cfg %||% list()
  spatial_cfg <- cfg$spatial_weights %||% list()
  value <- spatial_cfg$snap %||% cfg$spatial_snap %||% NULL
  if (is.null(value) || !length(value) || is.na(value[[1]])) return(NULL)
  value <- suppressWarnings(as.numeric(value[[1]]))
  if (!is.finite(value) || value < 0) stop("Spatial snap must be a finite non-negative
  ↪ number.", call. = FALSE)
  value
}

poly2nb_with_optional_snap <- function(panel, queen, snap = NULL) {
  if (is.null(snap)) return(spdep::poly2nb(panel, queen = queen))
  spdep::poly2nb(panel, queen = queen, snap = snap)
}

```

```

#' build spatial weights for selected district-panel rows
#'
#' @return A list with nb, binary matrix, and row-standardized listw objects.
build_spatial_weights_for_rows <- function(district_panel, rows, queen = FALSE, snap =
  ↪ NULL) {
  if (!inherits(district_panel, "sf")) {
    return(list(status = "out_of_active_pipeline", reason = "Requires sf geometry."))
  }
  need_pkg("spdep", "spatial weights")
  need_pkg("sf", "spatial weights")

  rows <- as.integer(rows)
  rows <- rows[is.finite(rows) & rows >= 1L & rows <= nrow(district_panel)]
  if (!length(rows)) {
    return(list(status = "out_of_active_pipeline", reason = "No rows with usable
      ↪ geometry."))
  }
  geom <- sf::st_geometry(district_panel)
  rows <- rows[!sf::st_is_empty(geom[rows])]
  if (!length(rows)) {
    return(list(status = "out_of_active_pipeline", reason = "Selected rows have empty
      ↪ geometry."))
  }

  spatial_warnings <- character()
  capture_expected_spatial_warning <- function(expr) {
    withCallingHandlers(
      expr,
      warning = function(w) {
        msg <- conditionMessage(w)
        expected <- grepl("no neighbours|sub-graphs", msg, ignore.case = TRUE)
        if (expected) {
          spatial_warnings <- unique(c(spatial_warnings, msg))
          invokeRestart("muffleWarning")
        }
      }
    )
  }

  panel <- district_panel[rows, , drop = FALSE]
  nb <- capture_expected_spatial_warning(poly2nb_with_optional_snap(panel, queen = queen,
  ↪ snap = snap))
  W <- capture_expected_spatial_warning(spdep::nb2mat(nb, style = "B", zero.policy =
  ↪ TRUE))
  listw <- capture_expected_spatial_warning(spdep::nb2listw(nb, style = "W", zero.policy
  ↪ = TRUE))

  neighbor_counts <- spdep::card(nb)
  n_components <- spdep::n.comp.nb(nb)$nc %||% NA_integer_
  snap_used <- if (is.null(snap)) NA_real_ else as.numeric(snap)
  panel_data <- as.data.frame(sf::st_drop_geometry(panel), stringsAsFactors = FALSE)
  identifier_cols <- intersect(
    c("district_panel_id", "state_20", "district_20", "state_std", "district_std"),
    names(panel_data)
  )
}

```

```

neighbor_ledger <- data.frame(
  row_index = rows,
  n_neighbors = neighbor_counts,
  panel_data[identifier_cols],
  check.names = FALSE,
  stringsAsFactors = FALSE
)

out <- list(
  status = "constructed",
  contiguity = if (isTRUE(queen)) "queen" else "rook",
  style = "W",
  matrix_style = "B",
  zero_policy = TRUE,
  row_index = rows,
  nb = nb,
  W = W,
  listw = listw,
  neighbor_counts = neighbor_counts,
  neighbor_ledger = neighbor_ledger,
  n = length(rows),
  n_islands = sum(neighbor_counts == 0L),
  mean_neighbors = mean(neighbor_counts),
  n_subgraphs = n_components,
  snap = snap_used,
  snap_investigation_needed = sum(neighbor_counts == 0L) > 0L || (!is.na(n_components)
    ⇨ && n_components > 1L),
  warnings = spatial_warnings
)
class(out) <- c("emi_spatial_weights", class(out))
out
}

#' diagnose spatial weights
#'
diagnose_spatial_weights <- function(district_panel, spatial_weights, cfg) {
  if (!inherits(district_panel, "sf")) {
    return(data.frame(
      diagnostic = "spatial_weights",
      status = "out_of_active_pipeline",
      reason = "Requires sf geometry.",
      stringsAsFactors = FALSE
    ))
  }
}

comparison <- compare_rook_queen_contiguity(district_panel, snap =
  ⇨ spatial_snap_setting(cfg))
base <- data.frame(
  diagnostic = "spatial_weights",
  status = spatial_weights$status %||% "constructed",
  contiguity = spatial_weights$contiguity %||% NA_character_,
  style = spatial_weights$style %||% NA_character_,
  matrix_style = spatial_weights$matrix_style %||% NA_character_,
  n = spatial_weights$n %||% NA_real_,
  n_islands = spatial_weights$n_islands %||% NA_real_,

```

```

mean_neighbors = spatial_weights$mean_neighbors %||% NA_real_,
n_subgraphs = spatial_weights$n_subgraphs %||% NA_integer_,
snap = spatial_weights$snap %||% NA_real_,
snap_investigation_needed = spatial_weights$snap_investigation_needed %||% NA,
panel_scope = spatial_weights$panel_scope %||%
  ↪ "current_final_matched_panel_non_empty_geometry",
warnings = paste(spatial_weights$warnings %||% attr(spatial_weights,
  ↪ "spatial_warnings") %||% character(), collapse = "; "),
stringsAsFactors = FALSE
)
attr(base, "rook_queen_comparison") <- add_spatial_weight_reference(comparison)
attr(base, "legacy_reference") <- spatial_weight_reference()
attr(base, "neighbor_counts") <- summarize_neighbor_counts(spatial_weights)
attr(base, "islands") <- summarize_islands(spatial_weights)
attr(base, "connectivity") <- summarize_spatial_connectivity(spatial_weights)
base
}

#' compare rook and queen contiguity
#'
#' The legacy Rmd commented out an sfExtras rook-vs-queen check and recorded
#' nearly identical mean neighbor counts (rook 4.780165, queen 4.783471) plus
#' similar run times. To avoid an extra sfExtras dependency, this project uses
#' the same `spdep::poly2nb()` implementation used by the final rook weights and
#' records elapsed time and average neighbor counts for both choices.
compare_rook_queen_contiguity <- function(district_panel, snap = NULL) {
  if (!inherits(district_panel, "sf")) return(tibble::tibble())
  need_pkg("spdep", "rook/queen contiguity comparison")
  need_pkg("sf", "rook/queen contiguity comparison")
  panel <- district_panel[spatial_weight_final_panel_rows(district_panel), , drop =
  ↪ FALSE]
  if (!nrow(panel)) return(tibble::tibble())

  one <- function(queen) {
    warnings <- character()
    elapsed <- system.time({
      nb <- withCallingHandlers(
        poly2nb_with_optional_snap(panel, queen = queen, snap = snap),
        warning = function(w) {
          msg <- conditionMessage(w)
          if (grepl("no neighbours|sub-graphs", msg, ignore.case = TRUE)) {
            warnings <- unique(c(warnings, msg))
            invokeRestart("muffleWarning")
          }
        }
      )
    })
  }
  }][["elapsed"]]
cards <- spdep::card(nb)
tibble::tibble(
  contiguity = if (isTRUE(queen)) "queen" else "rook",
  n = length(nb),
  mean_neighbors = mean(cards),
  n_islands = sum(cards == 0L),
  n_subgraphs = spdep::n.comp.nb(nb)$nc %||% NA_integer_,
  snap = if (is.null(snap)) NA_real_ else as.numeric(snap),

```

```

    panel_scope = "current_final_matched_panel_non_empty_geometry",
    elapsed_seconds = unname(elapsed),
    warnings = paste(warnings, collapse = "; ")
  )
}

safe_bind_rows(list(one(FALSE), one(TRUE)))
}

spatial_weight_reference <- function() {
  data.frame(
    contiguity = c("rook", "queen"),
    legacy_method = c(
      "sfExtras::st_rook() timing comment; final weights use spdep::poly2nb(queen =
      ↪ FALSE)",
      "sfExtras::st_queen() timing comment; benchmark uses spdep::poly2nb(queen = TRUE)"
    ),
    legacy_mean_neighbors = c(4.780165, 4.783471),
    legacy_elapsed_note = c("legacy comment recorded similar run time to queen", "legacy
    ↪ comment recorded similar run time to rook"),
    interpretation = "Current means are intentionally computed on the active final
    ↪ matched panel with non-empty geometry; they may differ from legacy exploratory
    ↪ sfExtras objects, but they now match the panel used for current spatial
    ↪ diagnostics.",
    stringsAsFactors = FALSE
  )
}

add_spatial_weight_reference <- function(comparison) {
  comparison <- as.data.frame(comparison, stringsAsFactors = FALSE)
  if (!nrow(comparison)) return(comparison)
  ref <- spatial_weight_reference()[c("contiguity", "legacy_mean_neighbors")]
  out <- merge(comparison, ref, by = "contiguity", all.x = TRUE, sort = FALSE)
  out$mean_neighbor_delta_from_legacy <- out$mean_neighbors - out$legacy_mean_neighbors
  out$pct_delta_from_legacy <- 100 * out$mean_neighbor_delta_from_legacy /
  ↪ out$legacy_mean_neighbors
  out
}

#' summarize islands
#'
summarize_islands <- function(spatial_weights) {
  ledger <- summarize_neighbor_counts(spatial_weights)
  if (!nrow(ledger)) return(ledger)
  ledger[ledger$n_neighbors == 0L, , drop = FALSE]
}

#' summarize neighbor counts
#'
summarize_neighbor_counts <- function(spatial_weights) {
  if (!inherits(spatial_weights, "emi_spatial_weights")) return(tibble::tibble())
  ledger <- spatial_weights$neighbor_ledger %||% data.frame(
    row_index = spatial_weights$row_index,
    n_neighbors = spatial_weights$neighbor_counts
  )
}

```

```

)
tibble::as_tibble(ledger)
}

#' Summarize graph connectivity and whether `poly2nb()` snap needs review
summarize_spatial_connectivity <- function(spatial_weights) {
  if (!inherits(spatial_weights, "emi_spatial_weights")) return(tibble::tibble())
  tibble::tibble(
    n = spatial_weights$n,
    n_islands = spatial_weights$n_islands,
    n_subgraphs = spatial_weights$n_subgraphs,
    snap = spatial_weights$snap,
    snap_investigation_needed = isTRUE(spatial_weights$snap_investigation_needed),
    recommendation = if (isTRUE(spatial_weights$snap_investigation_needed)) {
      "Inspect geometry validity and sensitivity to a documented poly2nb snap value
      ↪ before interpreting spatial results."
    } else {
      "No island/subgraph-driven snap investigation is currently indicated."
    }
  )
}

#' save spatial weight diagnostics
#'
save_spatial_weight_diagnostics <- function(diagnostics, dir =
  ↪ "outputs/diagnostics/extended/spatial") {
  dir.create(dir, recursive = TRUE, showWarnings = FALSE)
  paths <- c(
    summary = write_diagnostic_csv(as.data.frame(diagnostics), file.path(dir,
      ↪ "spatial_weights_summary.csv")),
    rook_queen_comparison = write_diagnostic_csv(attr(diagnostics,
      ↪ "rook_queen_comparison") %||% data.frame(), file.path(dir,
      ↪ "rook_queen_contiguity_comparison.csv")),
    legacy_reference = write_diagnostic_csv(attr(diagnostics, "legacy_reference") %||%
      ↪ data.frame(), file.path(dir, "spatial_weights_legacy_reference.csv")),
    neighbor_counts = write_diagnostic_csv(attr(diagnostics, "neighbor_counts") %||%
      ↪ data.frame(), file.path(dir, "spatial_weights_neighbor_counts.csv")),
    islands = write_diagnostic_csv(attr(diagnostics, "islands") %||% data.frame(),
      ↪ file.path(dir, "spatial_weights_islands.csv")),
    connectivity = write_diagnostic_csv(attr(diagnostics, "connectivity") %||%
      ↪ data.frame(), file.path(dir, "spatial_weights_connectivity.csv"))
  )
  output_manifest(paths)
}

```

Reusable IV specification definitions

Source: R/iv/build_iv_formulas.R

```

make_iv_formula <- function(dep, endog, instruments, controls = NULL, fixed_effects =
  ↪ NULL) {
  stats::as.formula(paste(
    dep, '~', paste(c(endog, controls, fixed_effects), collapse = ' + '), '|',
    paste(c(instruments, controls, fixed_effects), collapse = ' + ')
  ))
}

```

```

}

legacy_2007_iv_controls <- function() {
  c(
    'consumption_0708', 'gini_cons_0708', 'pct_urban', 'avg_hh_size',
    'dependency_ratio', 'pct_fem_head', 'pct_hindu', 'pct_muslim',
    'pct_st', 'pct_sc', 'pct_obc', 'pct_small_land', 'pct_medium_land',
    'pct_large_land', 'pct_head_lit_to_primary', 'pct_head_secondary_plus'
  )
}

preferred_iv_variables <- function() {
  list(
    treatment = "emi_exposure_all_children_0708",
    instrument = "ling_distance_nonzero_mean"
  )
}

# Historical specification retained only for the optional legacy-geography comparison.
build_legacy_iv_formulas <- function() {
  controls <- legacy_2007_iv_controls()
  list(
    consumption = make_iv_formula(
      "consumption_pct_change", "EMIE", "wavg_ling_degrees", controls
    ),
    gini = make_iv_formula(
      "gini_change", "EMIE", "wavg_ling_degrees", controls
    )
  )
}

# Preferred public formulas: all-child EMI exposure, the full-distribution
# nonzero linguistic-distance scalar, predetermined Census controls, and state FE.
build_revised_iv_formulas <- function() {
  spec <- preferred_iv_variables()
  controls <- census_2001_main_controls()
  state_fe <- 'factor(state_code_2001)'
  list(
    consumption = make_iv_formula(
      'real_log_consumption_change', spec$treatment, spec$instrument,
      controls = controls, fixed_effects = state_fe
    ),
    consumption_ancova = make_iv_formula(
      'log_real_consumption_1718', spec$treatment, spec$instrument,
      controls = c('log_real_consumption_0708', controls), fixed_effects = state_fe
    ),
    consumption_nominal = make_iv_formula(
      'log_consumption_difference', spec$treatment, spec$instrument,
      controls = controls, fixed_effects = state_fe
    ),
    consumption_legacy_controls = make_iv_formula(
      'log_consumption_difference', spec$treatment, spec$instrument,
      controls = legacy_2007_iv_controls(), fixed_effects = state_fe
    )
  )
}

```



```
)
}
```

Multicollinearity and spatial autocorrelation diagnostics

Source: R/diagnostics/diagnose_multicollinearity.R

```
#' Structural-regressor formula for multicollinearity diagnostics
#'
#' For `ivreg` models, multicollinearity is diagnosed on the stage-two
#' structural regressors, not on the instrument matrix. The `ivreg` package
#' exposes this component through its formula/model-matrix methods.
multicollinearity_formula <- function(model) {
  if (inherits(model, "ivreg")) {
    return(stats::formula(model, component = "regressors"))
  }
  stats::formula(model)
}

#' Structural-regressor design matrix
multicollinearity_design_matrix <- function(model) {
  iv_structural_model_matrix(model)
}

#' Normalize `car::vif()` output to a stable data-frame contract
normalize_vif_output <- function(x, model_scope) {
  if (is.null(x) || !length(x)) return(data.frame())
  if (is.matrix(x)) {
    out <- data.frame(
      term = rownames(x),
      gvif = as.numeric(x[, "GVIF"]),
      df = as.integer(x[, "Df"]),
      gvif_scaled = as.numeric(x[, "GVIF^(1/(2*Df))"]),
      stringsAsFactors = FALSE
    )
    out$vif <- ifelse(out$df == 1L, out$gvif, NA_real_)
  } else {
    out <- data.frame(
      term = names(x),
      gvif = as.numeric(x),
      df = 1L,
      gvif_scaled = sqrt(as.numeric(x)),
      vif = as.numeric(x),
      stringsAsFactors = FALSE
    )
  }
  out$model_scope <- model_scope
  out$status <- "estimated"
  out$reason <- NA_character_
  out[c("term", "model_scope", "df", "vif", "gvif", "gvif_scaled", "status", "reason")]
}

#' Compute term-aware VIF/GVIF diagnostics
compute_vif_if_applicable <- function(model) {
  scope <- if (inherits(model, "ivreg")) "ivreg_structural_regressors" else
  ↪ "model_regressors"
```

```

if (inherits(model, "ivreg") && !requireNamespace("ivreg", quietly = TRUE)) {
  return(data.frame(
    term = NA_character_, model_scope = scope, df = NA_integer_,
    vif = NA_real_, gvif = NA_real_, gvif_scaled = NA_real_,
    status = "unavailable", reason = "Package 'ivreg' is not installed.",
    stringsAsFactors = FALSE
  ))
}
if (!requireNamespace("car", quietly = TRUE)) {
  return(data.frame(
    term = NA_character_, model_scope = scope, df = NA_integer_,
    vif = NA_real_, gvif = NA_real_, gvif_scaled = NA_real_,
    status = "unavailable", reason = "Package 'car' is not installed.",
    stringsAsFactors = FALSE
  ))
}
tryCatch(
  normalize_vif_output(car::vif(model), scope),
  error = function(e) data.frame(
    term = NA_character_, model_scope = scope, df = NA_integer_,
    vif = NA_real_, gvif = NA_real_, gvif_scaled = NA_real_,
    status = "unavailable", reason = conditionMessage(e),
    stringsAsFactors = FALSE
  )
)
}

#' Scale-free structural-regressor condition number
#'
#' The raw condition number is dominated by the intercept and arbitrary units.
#' Following the IV diagnostics guidance, diagnose numerical collinearity on
#' centered, unit-variance non-intercept structural regressors instead.
standardized_design_condition_number <- function(X) {
  X <- as.matrix(X)
  if (!nrow(X) || !ncol(X)) return(NA_real_)
  keep <- colnames(X) != "(Intercept)"
  X <- X[, keep, drop = FALSE]
  if (!ncol(X)) return(NA_real_)
  sds <- apply(X, 2L, stats::sd)
  X <- X[, is.finite(sds) & sds > 0, drop = FALSE]
  if (!ncol(X)) return(NA_real_)
  X <- scale(X)
  qr_x <- qr(X)
  if (!qr_x$rank) return(NA_real_)
  estimable <- qr_x$pivot[seq_len(qr_x$rank)]
  kappa(X[, estimable, drop = FALSE], exact = TRUE)
}

multicollinearity_model_names <- function(iv_models) {
  models <- if (is.list(iv_models) && !inherits(iv_models, c("lm", "ivreg"))) iv_models
  ↪ else list(iv_models)
  model_names <- names(models)
  if (is.null(model_names)) model_names <- rep("", length(models))
  missing_names <- !nzchar(model_names)

```

```

model_names[missing_names] <- paste0("model_", which(missing_names))
list(models = models, names = model_names)
}

#' Diagnose one model's structural-regressor design
single_model_multicollinearity <- function(model, model_name) {
  scope <- if (inherits(model, "ivreg")) "ivreg_structural_regressors" else
  ↪ "model_regressors"
  design <- tryCatch({
    X <- multicollinearity_design_matrix(model)
    data.frame(
      model = model_name,
      diagnostic = "design_matrix",
      term = NA_character_,
      model_scope = scope,
      n = nrow(X),
      rank = qr(X)$rank,
      columns = ncol(X),
      kappa = standardized_design_condition_number(X),
      df = NA_integer_,
      vif = NA_real_,
      gvif = NA_real_,
      gvif_scaled = NA_real_,
      status = "estimated",
      reason = NA_character_,
      stringsAsFactors = FALSE
    )
  }, error = function(e) {
    data.frame(
      model = model_name,
      diagnostic = "design_matrix",
      term = NA_character_, model_scope = scope,
      n = NA_integer_, rank = NA_integer_, columns = NA_integer_, kappa = NA_real_,
      df = NA_integer_, vif = NA_real_, gvif = NA_real_, gvif_scaled = NA_real_,
      status = "out_of_active_pipeline", reason = conditionMessage(e),
      stringsAsFactors = FALSE
    )
  })

  vif <- compute_vif_if_applicable(model)
  if (!nrow(vif)) return(design)
  vif$model <- model_name
  vif$diagnostic <- "vif"
  vif$n <- NA_integer_
  vif$rank <- NA_integer_
  vif$columns <- NA_integer_
  vif$kappa <- NA_real_
  vif <- vif[names(design)]
  rbind(design, vif)
}

#' Diagnose multicollinearity
#'
#' @return A data frame containing a design-matrix summary row and term-level
#' VIF/GVIF rows for every supplied model. Factor terms use GVIF and the

```

```

#'    dimension-adjusted  $\widehat{GVIF^{(1/(2*Df))}}$ ; IV models use structural regressors.
diagnose_multicollinearity <- function(district_panel, iv_models, cfg) {
  named <- multicollinearity_model_names(iv_models)
  safe_bind_rows(Map(single_model_multicollinearity, named$models, named$names))
}

#' Save public multicollinearity diagnostics
save_multicollinearity_diagnostics <- function(diagnostics, path =
  ↪ "outputs/diagnostics/public/multicollinearity_diagnostics.csv") {
  write_diagnostic_csv(diagnostics, path)
}

```