

Coding Sample: Escaping Inequality in India

Selected Excerpts

Rishav Roy

What follows are excerpts from the 138-page code appendix to my independent paper and writing sample, “Escaping Inequality in India: The Role of English-Medium Instruction.”

The next 25 pages illustrate data ingestion and harmonization, complex survey econometrics, IV estimation, and careful diagnostics. Each excerpt is a part of the code chunks listed below:

14-16. Scalable **import, cleaning, and merging** of multi-file census and geospatial data.

- Chunk 4: Import key data files

31-32. **QA and identifier-disambiguation checks** for district-level crosswalk.

- Chunk 6: Match districts: Construct district tracking dfs

50-52. **Testing for systematic missingness** using the Benjamini–Hochberg procedure and parallelized logistic regressions.

- Chunk 8: Participation model: Are NAs randomly distributed

55-57. **Survey-weighted probit estimation**, average marginal effects derivation, and setup for sample selection correction under a complex survey design.

- Chunk 9: Participation model: Estimation and IMR
- Chunk 10: Participation model: Derive AME

80-88. Reusable **record-linkage functions** using cascading fuzzy matches and unmatched/false-match diagnostics to harmonize district identifiers across years and data sources.

- Chunk 16: Match districts: Test joining methods
- Chunk 17: Match districts: Define joining functions

105-107. Reusable **IV specification and estimation pipeline** with clustered inference, plus contiguity-based spatial weights and setup for spatially lagged models.

- Chunk 24: Geospatial: Neighbor list construction
- Chunk 25: 2SLS: Controls and 2SLS formula
- Chunk 26: 2SLS: Baseline 2SLS models

111-112. **Robustness checks and diagnostics** for multicollinearity, spatial autocorrelation.

- Chunk 28: Multicollinearity check
- Chunk 29: Geospatial: Spatial autocorrelation tests

Scalable import, cleaning, and merging of Census/geospatial data

Source: R/io/read_inputs.R

```
## read nss 2007 education
##
read_nss_2007_education <- function(paths) {
  read_manifest_group(paths, "nss_2007_education")
}

## read nss 2007 consumption
##
read_nss_2007_consumption <- function(paths) {
  read_manifest_group(paths, "nss_2007_consumption")
}

## read nss 2017 education
##
read_nss_2017_education <- function(paths) {
  read_manifest_group(paths, "nss_2017_education")
}

## read census 2001 mother tongue
##
read_census_2001_mother_tongue <- function(paths) {
  read_manifest_group(paths, "census_2001_mother_tongue")
}

## read district boundaries 2020
##
read_district_boundaries_2020 <- function(paths) {
  rows <- require_manifest_files(paths, "district_boundaries_2020")
  shp <- rows[tolower(rows$file_type) == "shp", , drop = FALSE]
  if (!nrow(shp)) stop("The district-boundary manifest includes shapefile sidecars but no
    ↪ .shp row.", call. = FALSE)
  need_pkg("sf", "district boundaries")
  sf::st_read(shp$absolute_path[[1]], quiet = TRUE)
}

## read district change sources
##
read_district_change_sources <- function(paths) {
  read_manifest_group(paths, "district_changes")
}

## list ilo figure paths
##
list_ilo_figure_paths <- function(paths) {
  rows <- require_manifest_files(paths, "ilo_figures")
  stats::setNames(rows$absolute_path, rows$file_id)
}

## read one manifest row
##
## @return Reader output for raw data files, or a validated path for sidecars/assets.
read_by_manifest_row <- function(row) {
```

```

path <- row$absolute_path[[1]]
file_type <- tolower(row$file_type[[1]])
switch(
  file_type,
  sav = read_sav_short(path),
  csv = read_csv_short(path),
  xls = read_excel_short(path),
  xlsx = read_excel_short(path),
  ods = read_ods_short(path),
  shp = {
    need_pkg("sf", "shapefiles")
    sf::st_read(path, quiet = TRUE)
  },
  shx = path,
  dbf = path,
  prj = path,
  png = path,
  stop("No reader implemented for ", file_type, call. = FALSE)
)
}

#' read all manifest rows for a source
#'
#' @return Named list of reader outputs keyed by manifest file_id.
read_manifest_group <- function(paths, source_id) {
  rows <- require_manifest_files(paths, source_id)
  stats::setNames(lapply(seq_len(nrow(rows)), function(i) read_by_manifest_row(rows[i,
↵ ])), rows$file_id)
}

```

QA and identifier-disambiguation checks

Source: R/districts/build_district_tracker.R

```

#' build district tracker
#'
build_district_tracker <- function(raw_district_changes) {
  combine_district_tracker_sources(raw_district_changes)
}

#' parse alluvial district changes
#'
parse_alluvial_district_changes <- function(x) {
  x
}

#' parse india district tracker
#'
parse_india_district_tracker <- function(x) {
  x
}

#' parse carveouts renamings
#'
parse_carveouts_renamings <- function(x) {
  x
}

```

```

}

#' parse new districts created
#'
parse_new_districts_created <- function(x) {
  x
}

#' parse name changes
#'
parse_name_changes <- function(x) {
  x
}

#' parse district splits
#'
parse_district_splits <- function(x) {
  x
}

#' combine district tracker sources
#'
combine_district_tracker_sources <- function(raw_district_changes) {
  out <- safe_bind_rows(lapply(names(raw_district_changes), function(name) {
    x <- safe_df(raw_district_changes[[name]])
    x$source_file_id <- name
    x
  })))
  if (nrow(out)) out$.row_in_source <- ave(seq_len(nrow(out)), out$source_file_id, FUN =
    ↪ seq_along)
  out
}

#' standardize tracker years
#'
standardize_tracker_years <- function(tracker) {
  tracker
}

#' standardize tracker names
#'
standardize_tracker_names <- function(tracker) {
  tracker
}

```

Parallelized missingness diagnostics with BH correction

Source: R/selection/diagnose_missingness.R

```

#' diagnose missingness
#'
diagnose_missingness <- function(selection_data, cfg) {
  data.frame(
    missing_var = names(selection_data)[vapply(selection_data, function(x) any(is.na(x)),
    ↪ logical(1))],
    stringsAsFactors = FALSE
  )
}

```

```

)
}

#' check missing logit parallel
#'
check_missing_logit_parallel <- function(df, miss_vars, covars, method_p = "BH") {
  rhs <- paste(covars, collapse = " + ")
  fit_one <- function(m) { f <- stats::as.formula(paste0("is.na(", m, ") ~ ", rhs)); fit
  <- stats::glm(f, data = df, family = stats::binomial); broom::tidy(fit) |>
  dplyr::mutate(missing_var = m, pseudoR2 = 1 - fit$deviance / fit$null.deviance, nobs
  = stats::nobs(fit)) }
  out <- dplyr::bind_rows(lapply(miss_vars, fit_one))
  out |> dplyr::group_by(missing_var) |> dplyr::mutate(p_adj = {adj <- rep(NA_real_,
  dplyr::n()); idx <- term != "(Intercept)"; adj[idx] <- p.adjust(p.value[idx], method
  = method_p); adj}) |> dplyr::ungroup()
}

#' summarize missingness by variable
#'
summarize_missingness_by_variable <- function(df) {
  data.frame(
    missing_var = names(df),
    n_missing = vapply(df, function(x) sum(is.na(x)), integer(1)),
    stringsAsFactors = FALSE
  )
}

#' run missingness logits
#'
run_missingness_logits <- function(df, miss_vars, covars) {
  check_missing_logit_parallel(df, miss_vars, covars)
}

#' adjust missingness pvalues bh
#'
adjust_missingness_pvalues_bh <- function(df) {
  df
}

#' save missingness diagnostics
#'
save_missingness_diagnostics <- function(diagnostics) {
  diagnostics
}

```

Survey-weighted probit and IMR setup

Source: R/selection/estimate_selection_probit.R

```

legacy_probit_variables <- function(selection_data) {
  controls <- c("AGE", "age", "SEX", "HH_SIZE", "RELIGION", "SOCIAL_GROUP", "SECTOR")
  exclusion_all_kids <- c("DIST_FROM_NEAREST_PRIMARY_CLASS", "father_educ")
  exclusion_district <- c(
    "dmean_num_IS_EDU_FREE", "dmean_num_TUTION_FEE_WAIVED",
    "dmean_num_REC_D_SCHOLARSHIP_STIPEND", "dmean_num_REC_D_TXT_BOOKS",
    "dmean_num_REC_D_STATIONERY", "dmean_num_MID_DAY_MEAL_ETC_REC_D",

```

```

    "dmean_num_ENROLLMENT_COST"
  )
  intersect(c(controls, exclusion_all_kids, exclusion_district), names(selection_data))
}

#' estimate selection probit
#'
estimate_selection_probit <- function(selection_data, cfg) {
  if (!"enrolled" %in% names(selection_data) || all(is.na(selection_data$enrolled))) {
    return(list(status = "out_of_active_pipeline", reason = "No enrolled variable."))
  }
  covars <- legacy_probit_variables(selection_data)
  if (!length(covars)) {
    return(list(status = "out_of_active_pipeline", reason = "No probit covariates."))
  }
  f_probit <- stats::reformulate(covars, response = "enrolled")
  if (identical(cfg$mode, "final") && requireNamespace("survey", quietly = TRUE)) {
    design <- build_survey_design_selection(selection_data)
    if (!is.null(design)) {
      out <- fit_selection_probit(design, f_probit)
      attr(out, "legacy_probit_formula") <- f_probit
      return(out)
    }
  }
  out <- stats::glm(f_probit, data = selection_data, family = stats::binomial(link =
↪ "probit"))
  attr(out, "legacy_probit_formula") <- f_probit
  out
}

#' build survey design selection
#'
build_survey_design_selection <- function(selection_df) {
  psu <- first_col(selection_df, c("FSU_SL_NO", "fsu", "PSU", "psu"))
  weight <- first_col(selection_df, c("weight", "WEIGHT", "Multiplier", "multiplier"))
  strata_cols <- intersect(c("STATE", "STRATUM", "SUB_STRATUM_NO"), names(selection_df))
  if (is.null(psu) || is.null(weight) || !length(strata_cols)) return(NULL)
  options(survey.lonely.psu = "average")
  selection_df$.survey_strata <- interaction(selection_df[strata_cols], drop = TRUE)
  survey::svydesign(
    ids = stats::as.formula(paste0("~", psu)),
    strata = ~.survey_strata,
    weights = stats::as.formula(paste0("~", weight)),
    data = selection_df,
    nest = TRUE
  )
}

#' fit selection probit
#'
fit_selection_probit <- function(selection_design, f_probit) {
  survey::svyglm(f_probit, design = selection_design, family = stats::quasibinomial(link
↪ = "probit"))
}

```

```

#' compute inverse mills ratio
#'
compute_inverse_mills_ratio <- function(model, selection_df, f_probit = NULL) {
  if (is.null(f_probit)) f_probit <- attr(model, "legacy_probit_formula") %||%
    ↪ stats::formula(model)
  needed <- all.vars(f_probit)
  needed <- intersect(needed, names(selection_df))
  comp_cases <- stats::complete.cases(selection_df[, needed, drop = FALSE])
  lp <- rep(NA_real_, nrow(selection_df))
  lp[comp_cases] <- stats::predict(model, newdata = selection_df[comp_cases, , drop =
    ↪ FALSE], type = "link")
  phi_lp <- stats::dnorm(lp)
  Phi_lp <- pmin(pmax(stats::pnorm(lp), .Machine$double.eps), 1 - .Machine$double.eps)
  enrolled <- as.character(selection_df$enrolled)
  selection_df$IMR <- ifelse(enrolled == "Yes", phi_lp / Phi_lp, ifelse(enrolled == "No",
    ↪ phi_lp / (1 - Phi_lp), NA_real_))
  selection_df
}

#' tidy selection model
#'
tidy_selection_model <- function(model) {
  broom::tidy(model)
}

```

Average marginal effects via automatic differentiation

Source: R/selection/compute_average_marginal_effects.R

Use automatic differentiation for SE delta-method gradients, as in the legacy Rmd.

```

#' compute average marginal effects
#'
compute_average_marginal_effects <- function(selection_model, selection_data, cfg =
  ↪ list()) {
  if (!inherits(selection_model, "glm")) {
    return(ame_out_of_pipeline(
      selection_model$status %||% "out_of_active_pipeline",
      selection_model$reason %||% "Selection model not estimated in smoke mode"
    ))
  }

  # `selection_model` is often a survey::svyglm object loaded from the targets
  # store. Survey lonely-PSU handling is an R option, not a durable model
  # attribute, so it may not be set in the downstream AME target even though it
  # was set when the model was estimated. Set it explicitly around all AME
  # computations to make reruns deterministic and warning/error-free.
  old_lonely_psu <- getOption("survey.lonely.psu")
  old_domain_lonely <- getOption("survey.adjust.domain.lonely")
  options(survey.lonely.psu = "average", survey.adjust.domain.lonely = TRUE)
  on.exit({
    if (is.null(old_lonely_psu)) {
      options(survey.lonely.psu = NULL)
    } else {
      options(survey.lonely.psu = old_lonely_psu)
    }
  })
}

```

```

    if (is.null(old_domain_lonely)) {
      options(survey.adjust.domain.lonely = NULL)
    } else {
      options(survey.adjust.domain.lonely = old_domain_lonely)
    }
  }, add = TRUE)

#' run expression while muffling survey lonely-PSU warnings
#'
#' survey emits lonely-PSU warnings even when survey.lonely.psu is set to
#' average and the adjustment is intentional. Muffle only those handled
#' warnings inside the AME target so strict final builds remain warning-clean.
muffle_survey_lonely_psu_warnings <- function(expr) {
  withCallingHandlers(
    expr,
    warning = function(w) {
      msg <- conditionMessage(w)
      lonely_psu_warning <- grepl(
        "has only one PSU at stage",
        msg,
        fixed = TRUE
      )
      if (lonely_psu_warning) invokeRestart("muffleWarning")
    }
  )
}

if (!isTRUE(cfg$run_full_ame)) {
  return(format_ame_results(data.frame(
    term = names(stats::coef(selection_model)),
    estimate = unname(stats::coef(selection_model)),
    std.error = NA_real_,
    statistic = NA_real_,
    p.value = NA_real_,
    s.value = NA_real_,
    conf.low = NA_real_,
    conf.high = NA_real_,
    method = "coefficient_fallback",
    status = "estimated",
    reason = "Draft config uses coefficient fallback instead of full AME computation"
  )))
}

if (!requireNamespace("marginaleffects", quietly = TRUE)) {
  return(ame_out_of_pipeline(
    "out_of_active_pipeline",
    "Package marginaleffects is required for full AME computation"
  ))
}

out <- muffle_survey_lonely_psu_warnings(
  tryCatch(
    compute_ames_autodiff(selection_model, selection_data),
    error = function(e) {

```



```

        fallback <- compute_ames_probit_analytic(selection_model, selection_data)
        fallback$method <- "delta_method_analytic_probit"
        fallback$reason <- NA_character_
        fallback
      }
    )
  )
  format_ame_results(out)
}

ame_out_of_pipeline <- function(status, reason) {
  data.frame(
    term = NA_character_,
    estimate = NA_real_,
    std.error = NA_real_,
    statistic = NA_real_,
    p.value = NA_real_,
    s.value = NA_real_,
    conf.low = NA_real_,
    conf.high = NA_real_,
    method = "not_run",
    status = status,
    reason = reason
  )
}

#' extract model-frame data and AME weights
#'
#' marginalEffects warns, correctly, when weighted/survey models are summarized
#' without an explicit `wts` argument. Build a newdata frame from the exact
#' model estimation sample and attach a synthetic weight column whenever the
#' fitted model exposes usable weights.
ame_model_weight <- function(model_data, model_weights = NULL) {
  weight_cols <- intersect(c("weight", "WEIGHT", "Multiplier", "multiplier",
    ↪ "(weights)"), names(model_data))
  if (length(weight_cols)) {
    return(suppressWarnings(as.numeric(model_data[[weight_cols[[1]]]])))
  }

  if (!is.null(model_weights) && length(model_weights) == nrow(model_data)) {
    return(suppressWarnings(as.numeric(model_weights)))
  }

  NULL
}

#' build marginalEffects newdata and weights
#'
#' @return List with `data`, `wts`, and `has_explicit_weights`.
ame_model_data_and_weights <- function(model) {
  model_data <- as.data.frame(stats::model.frame(model))
  model_weights <- tryCatch(stats::weights(model), error = function(e) NULL)
  weight <- ame_model_weight(model_data, model_weights)

```

```

if (!is.null(weight) && length(weight) == nrow(model_data) && any(is.finite(weight) &
  ↪ weight != 1)) {
  model_data$.ame_weight <- weight
  return(list(data = model_data, wts = ".ame_weight", has_explicit_weights = TRUE))
}

list(data = model_data, wts = FALSE, has_explicit_weights = FALSE)
}

#' run marginaleffects while muffling only a redundant explicit-weight warning
#'
run_avg_slopes <- function(model, model_data, wts) {
  withCallingHandlers(
    marginaleffects::avg_slopes(
      model,
      newdata = model_data,
      wts = wts,
      vcov = TRUE,
      type = "response"
    ),
    warning = function(w) {
      msg <- conditionMessage(w)
      explicit_weight_warning <- grepl(
        "normally good practice to specify weights using the `wts` argument",
        msg,
        fixed = TRUE
      )
      if (explicit_weight_warning && !identical(wts, FALSE)) {
        invokeRestart("muffleWarning")
      }
    }
  )
}

#' compute ames autodiff
#'
compute_ames_autodiff <- function(model, newdata) {
  # Use the exact estimation sample retained by glm/svyglm. Passing the full
  # pre-model selection_data can have extra rows omitted during model fitting,
  # which makes marginaleffects recycle weights/covariates silently.
  amed <- ame_model_data_and_weights(model)
  run_avg_slopes(model, amed$data, amed$wts)
}

#' compute ames fast draft
#'
compute_ames_fast_draft <- function(model, newdata, n = 200) {
  amed <- ame_model_data_and_weights(model)
  run_avg_slopes(
    model,
    dplyr::slice_sample(amed$data, n = min(n, nrow(amed$data))),
    amed$wts
  )
}

```

```

#' compute analytic probit AMEs
#'
#' @return Data frame of observed-data average marginal effects.
compute_ames_probit_analytic <- function(model, newdata) {
  coefs <- stats::coef(model)
  coef_table <- as.data.frame(summary(model)$coefficients)
  terms <- setdiff(names(coefs), "(Intercept)")
  if (!length(terms)) return(ame_out_of_pipeline("out_of_active_pipeline", "Selection
    ↪ model has no non-intercept terms.))

  newdata <- stats::model.frame(model)
  weights <- if ("weight" %in% names(newdata)) num(newdata$weight) else rep(1,
    ↪ nrow(newdata))
  eta <- stats::predict(model, type = "link")
  mean_phi <- wmean(stats::dnorm(eta), weights)
  factor_levels <- model$xlevels %||% list()

  rows <- lapply(terms, function(term) {
    factor_var <- names(factor_levels)[vapply(names(factor_levels), function(v)
    ↪ startsWith(term, v), logical(1))]]
    estimate <- NA_real_
    if (length(factor_var)) {
      var <- factor_var[[which.max(nchar(factor_var))]]
      level <- sub(paste0("^", var), "", term)
      if (level %in% factor_levels[[var]]) {
        hi <- newdata
        lo <- newdata
        hi[[var]] <- factor(level, levels = factor_levels[[var]])
        lo[[var]] <- factor(factor_levels[[var]][[1]], levels = factor_levels[[var]])
        estimate <- wmean(stats::predict(model, newdata = hi, type = "response") -
    ↪ stats::predict(model, newdata = lo, type = "response"), weights)
      }
    } else {
      estimate <- unname(coefs[[term]]) * mean_phi
    }
    p_col <- intersect(c("Pr(>|z|)", "Pr(>|t|)"), names(coef_table))
    p <- if (term %in% rownames(coef_table) && length(p_col)) coef_table[term,
    ↪ p_col[[1]] else NA_real_
    se_col <- intersect(c("Std. Error", "std.error"), names(coef_table))
    se_beta <- if (term %in% rownames(coef_table) && length(se_col)) coef_table[term,
    ↪ se_col[[1]] else NA_real_
    se <- if (is.finite(se_beta)) abs(se_beta * mean_phi) else NA_real_
    z <- if (is.finite(se) && se > 0) estimate / se else NA_real_
    p_out <- if (is.finite(z)) 2 * stats::pnorm(abs(z), lower.tail = FALSE) else p
    data.frame(
      term = term,
      estimate = estimate,
      std.error = se,
      statistic = z,
      p.value = p_out,
      s.value = if (is.finite(p_out) && p_out > 0) -log2(p_out) else NA_real_,
      conf.low = if (is.finite(se)) estimate - 1.96 * se else NA_real_,
      conf.high = if (is.finite(se)) estimate + 1.96 * se else NA_real_,
      method = "delta_method_analytic_probit",
      status = "estimated",

```

```

    reason = NA_character_,
    stringsAsFactors = FALSE
  )
})
do.call(rbind, rows)
}

is_marginaeffects_object <- function(x) {
  inherits(x, c("marginaeffects", "comparisons", "avg_comparisons", "slopes",
    ↪ "avg_slopes", "predictions"))
}

copy_first_ame_column <- function(out, target, candidates) {
  if (!target %in% names(out)) out[[target]] <- NA
  for (candidate in candidates) {
    if (!candidate %in% names(out)) next
    missing_target <- is.na(out[[target]])
    if (any(missing_target)) out[[target]][missing_target] <-
      ↪ out[[candidate]][missing_target]
  }
  out
}

normalize_ame_columns <- function(out) {
  # marginaeffects has used both dotted and snake_case names across versions
  # and model classes. Normalize here before selecting the public/audit schema
  # so uncertainty columns are not silently dropped downstream.
  aliases <- list(
    std.error = c("std_error", "std.err", "Std. Error"),
    statistic = c("z", "z.value", "z_value", "t", "t.value", "t_value"),
    p.value = c("p_value", "p", "Pr(>|z|)", "Pr(>|t|)"),
    s.value = c("s_value"),
    conf.low = c("conf_low", "conf.low", "2.5 %"),
    conf.high = c("conf_high", "conf.high", "97.5 %")
  )
  for (target in names(aliases)) {
    out <- copy_first_ame_column(out, target, aliases[[target]])
  }
  out
}

legacy_ame_lookup <- function() {
  data.frame(
    term = c(
      "AGE", "SEX", "HH_SIZE",
      "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION", "RELIGION",
      "SOCIAL_GROUP", "SOCIAL_GROUP", "SOCIAL_GROUP",
      "SECTOR",
      "DIST_FROM_NEAREST_PRIMARY_CLASS", "DIST_FROM_NEAREST_PRIMARY_CLASS",
      ↪ "DIST_FROM_NEAREST_PRIMARY_CLASS", "DIST_FROM_NEAREST_PRIMARY_CLASS",
      "father_educ", "father_educ", "father_educ", "father_educ", "father_educ",
      ↪ "father_educ", "father_educ",
      "dmean_num_IS_EDU_FREE", "dmean_num_TUTION_FEE_WAIVED",
      ↪ "dmean_num_RECD_SCHOLARSHIP_STIPEND",

```

```

      "dmean_num_RECD_TXT_BOOKS", "dmean_num_RECD_STATIONERY",
      ↪ "dmean_num_MID_DAY_MEAL_ETC_RECD", "dmean_num_ENROLLMENT_COST"
    ),
    contrast = c(
      "dY/dX", "Female - Male", "dY/dX",
      "Muslim - Hindu", "Christian - Hindu", "Sikh - Hindu", "Jain - Hindu", "Buddhist -
      ↪ Hindu", "Zoroastrian - Hindu", "Other - Hindu",
      "Scheduled Tribe - Other", "Scheduled Caste - Other", "Other Backward Class -
      ↪ Other",
      "Urban - Rural",
      "1km <= d <2kms - d<1km", "2kms<= d <3kms - d<1km", "3kms <= d <5kms - d<1km",
      ↪ "d>=5kms - d<1km",
      "Literate, no school - Illiterate", "Literate, school < primary - Illiterate",
      ↪ "Primary - Illiterate", "Upper primary - Illiterate", "Secondary - Illiterate",
      ↪ "Higher secondary - Illiterate", "Postsecondary+ - Illiterate",
      "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX", "dY/dX"
    ),
    Term = c(
      "Age (years)", "Female (ref: Male)", "Household size",
      "Religion: Muslim (ref: Hindu)", "Religion: Christian", "Religion: Sikh",
      ↪ "Religion: Jain", "Religion: Buddhist", "Religion: Zoroastrian", "Religion:
      ↪ Other",
      "Social group: Scheduled Tribe (ref: Other)", "Social group: Scheduled Caste",
      ↪ "Social group: Other Backward Class",
      "Urban (ref: Rural)",
      "Distance 1-2km (ref: <1km)", "Distance 2-3km", "Distance 3-5km", "Distance > 5km",
      "Father's educ.: Literate, no school (ref: Illiterate)", "Father's educ.: Literate,
      ↪ school < primary", "Father's educ.: Primary", "Father's educ.: Upper primary",
      ↪ "Father's educ.: Secondary", "Father's educ.: Higher secondary", "Father's
      ↪ educ.: Postsecondary+",
      "Educ. free available (ref: No)", "Tuition waiver received", "Scholarship/Stipend
      ↪ received", "Textbook(s) received", "Stationery received", "Mid-day meal, etc.
      ↪ received", "Enrollment cost (Rs.)"
    ),
    stringsAsFactors = FALSE
  )
}

```

```

legacy_ame_label_order <- function() legacy_ame_lookup()$Term

```

```

legacy_ame_modelsummary_label <- function(x) {
  x <- as.character(x)
  x <- gsub("\u00d7", ":", x, fixed = TRUE)
  x <- gsub("\\s+:", "\\s+", ":", x, perl = TRUE)
  x <- gsub("\\(ref:\\s*", "(ref: ", x, perl = TRUE)
  x
}

```

```

infer_ame_contrast <- function(out) {
  if ("contrast" %in% names(out)) return(as.character(out$contrast))
  contrast <- rep("dY/dX", nrow(out))
  if (!"term" %in% names(out)) return(contrast)
  if ("comparison" %in% names(out)) contrast <- as.character(out$comparison)
  contrast[is.na(contrast) | !nzchar(contrast)] <- "dY/dX"
  contrast
}

```

```

}

ame_factor_base_terms <- function() {
  unique(legacy_ame_lookup()$term[legacy_ame_lookup()$contrast != "dY/dX"])
}

normalize_encoded_factor_ame_terms <- function(out) {
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  if (!"term" %in% names(out)) return(out)
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(out)

  lookup <- legacy_ame_lookup()
  factor_terms <- ame_factor_base_terms()
  factor_terms <- factor_terms[order(nchar(factor_terms), decreasing = TRUE)]

  for (i in seq_len(nrow(out))) {
    current_term <- as.character(out$term[[i]])
    current_contrast <- as.character(out$contrast[[i]])
    if (!nzchar(current_term) || (!is.na(current_contrast) && current_contrast !=
      ↪ "dY/dX")) next

    direct <- lookup$term == current_term & lookup$contrast == current_contrast
    if (any(direct)) next

    for (base_term in factor_terms) {
      if (!startsWith(current_term, base_term)) next
      level <- trimws(sub(paste0("^", base_term), "", current_term))
      if (!nzchar(level)) next
      candidates <- lookup[lookup$term == base_term, , drop = FALSE]
      contrast_hit <- candidates$contrast[startsWith(candidates$contrast, paste0(level, "
      ↪ - "))]
      if (length(contrast_hit)) {
        out$term[[i]] <- base_term
        out$contrast[[i]] <- contrast_hit[[1]]
        break
      }
    }
  }
  out
}

attach_legacy_ame_labels <- function(out) {
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  if (!"term" %in% names(out) && "variable" %in% names(out)) out$term <- out$variable
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(out)
  out <- normalize_encoded_factor_ame_terms(out)
  lookup <- legacy_ame_lookup()
  labeled <- merge(out, lookup, by = c("term", "contrast"), all.x = TRUE, sort = FALSE)
  labeled$Term[is.na(labeled$Term)] <- labeled$term[is.na(labeled$Term)]
  labeled$Term <- factor(labeled$Term, levels = unique(c(legacy_ame_label_order(),
  ↪ labeled$Term)))
  labeled <- labeled[order(labeled$Term), , drop = FALSE]
  labeled$Term <- as.character(labeled$Term)
  labeled
}

```

```

legacy_modelsummary_marginaleffects_object <- function(out, native_ame) {
  if (!is_marginaleffects_object(native_ame)) return(NULL)
  out <- as.data.frame(out, stringsAsFactors = FALSE)
  required <- c("Term", "estimate", "std.error", "statistic", "p.value", "conf.low",
    ↪ "conf.high")
  if (!all(required %in% names(out))) return(native_ame)

  ms <- out[, required, drop = FALSE]
  ms$term <- legacy_ame_modelsummary_label(ms$Term)
  ms$contrast <- ""
  ms$Term <- NULL
  # Preserve the marginaleffects class and non-structural attributes so
  # modelsummary can use its native marginaleffects/tidy pathway. We still
  # hand modelsummary the labeled, ordered rows that are validated for the
  # public table, because passing the raw object caused modelsummary to reorder
  # contrasts and display labels against the wrong estimates.
  native_attributes <- attributes(native_ame)
  structural_attributes <- c("names", "row.names", "class")
  for (nm in setdiff(names(native_attributes), structural_attributes)) {
    attr(ms, nm) <- native_attributes[[nm]]
  }
  class(ms) <- unique(c(class(native_ame), class(ms)))
  ms
}

#' format ame results
#'
format_ame_results <- function(ame_results) {
  native_ame <- ame_results
  out <- tibble::as_tibble(ame_results)
  if (!"term" %in% names(out) && "variable" %in% names(out)) out$term <- out$variable
  if (!"contrast" %in% names(out)) out$contrast <- infer_ame_contrast(as.data.frame(out))
  out <- normalize_ame_columns(out)
  if (!"method" %in% names(out)) out$method <- "autodiff"
  if (!"status" %in% names(out)) out$status <- "estimated"
  if (!"reason" %in% names(out)) out$reason <- NA_character_
  legacy_terms <- unique(legacy_ame_lookup())$term
  use_legacy_schema <- any(out$term %in% legacy_terms) ||
    any(grepl("^dmean_num_", out$term %||% character()), perl = TRUE) ||
    any((out$contrast %||% "dY/dX") != "dY/dX", na.rm = TRUE)

  if (isTRUE(use_legacy_schema)) {
    out <- attach_legacy_ame_labels(out)
    required <- c(
      "Term", "term", "contrast", "estimate", "std.error", "statistic", "p.value",
      ↪ "s.value",
      "conf.low", "conf.high", "method", "status", "reason"
    )
  } else {
    required <- c(
      "term", "estimate", "std.error", "statistic", "p.value",
      "s.value", "conf.low", "conf.high", "method", "status", "reason"
    )
  }
}

```

```

}

for (nm in setdiff(required, names(out))) out[[nm]] <- NA
if ("p.value" %in% names(out) && "s.value" %in% names(out)) {
  missing_s <- is.na(out$s.value) & is.finite(out$p.value) & out$p.value > 0
  out$s.value[missing_s] <- -log2(out$p.value[missing_s])
}
out <- out[, required, drop = FALSE]
if (is_marginaleffects_object(native_ame)) {
  attr(out, "legacy_marginaleffects") <-
    ↪ legacy_modelsummary_marginaleffects_object(out, native_ame)
}
out
}

#' save ame results
#'
save_ame_results <- function(ame_results, path =
  ↪ "outputs/tables/diagnostics/ame_results.csv") {
  readr::write_csv(tibble::as_tibble(ame_results), path); path
}

```

Cascading fuzzy district record linkage

Source: R/districts/fuzzy_join_districts.R

```

# The functions below preserve the cascading fuzzy-match architecture from the legacy
  ↪ Rmd.
# sample-end marker appears near the end of this file after the functions are migrated.

#' evaluate distances
#'
evaluate_distances <- function(pairs, methods, thresholds, col1 = "str1", col2 = "str2")
  ↪ {
  if (length(methods) != length(thresholds)) stop("\"methods\" and \"thresholds\" must
    ↪ have the same length.")
  safe_bind_rows(lapply(seq_along(methods), function(i) {
    distance <- utils::adist(canon(pairs[[col1]]), canon(pairs[[col2]]), ignore.case =
  ↪ TRUE)[, 1]
    data.frame(
      str1 = pairs[[col1]],
      str2 = pairs[[col2]],
      method = methods[i],
      distance = distance,
      threshold = thresholds[i],
      match = distance <= thresholds[i],
      stringsAsFactors = FALSE
    )
  })))
}

#' fuzzy join sequence
#'
fuzzy_join_sequence <- function(df1, df2, dist1, state1, dist2, state2, methods,
  ↪ thresholds, mode = "full") {

```



```

df1_id <- df1 |> dplyr::mutate(.id1 = dplyr::row_number())
df2_id <- df2 |> dplyr::mutate(.id2 = dplyr::row_number())
df1_curr <- df1_id; df2_curr <- df2_id; matched_all <- tibble::tibble()
for (i in seq_along(methods)) {
  full_j <- fuzzyjoin::stringdist_join(df1_curr, df2_curr, by =
  ↪ stats::setNames(c(dist2, state2), c(dist1, state1)), mode = mode, method =
  ↪ methods[i], max_dist = thresholds[i], distance_col = "dist")
  matched_i <- full_j |> dplyr::filter(!is.na(.id1), !is.na(.id2)) |>
  ↪ dplyr::group_by(.id1) |> dplyr::slice_min(dist, n = 1, with_ties = FALSE) |>
  ↪ dplyr::ungroup() |> dplyr::group_by(.id2) |> dplyr::slice_min(dist, n = 1, with_ties
  ↪ = FALSE) |> dplyr::ungroup()
  matched_all <- dplyr::bind_rows(matched_all, matched_i)
  df1_curr <- df1_curr |> dplyr::filter(!(.id1 %in% matched_i$.id1)); df2_curr <-
  ↪ df2_curr |> dplyr::filter(!(.id2 %in% matched_i$.id2))
}
list(joined = matched_all, unmatched_df1 = df1_curr |> dplyr::select(-.id1),
  ↪ unmatched_df2 = df2_curr |> dplyr::select(-.id2))
}

#' unmatched rows
#'
#' @return A data frame of unmatched rows when the join map carries that attribute.
unmatched_rows <- function(join_map, ...) {
  attr(join_map, "unmatched_rows") %||% join_map[0, , drop = FALSE]
}

#' merge dfs into tracker
#'
#' Legacy-compatible cascading district matcher. `df_names` is a named list (or a
#' data frame) whose elements carry source-specific district/state names. For
#' each source we try every requested tracker year in order, keep one-to-one
#' canonical name matches, and then fuzzy-match remaining rows within state.
#'
#' @return A list with `joined_df`, `unmatched_df`, and `flagged_df`, matching
#' the object contract used by the legacy Rmd.
merge_dfs_into_tracker <- function(df_names, tracker = NULL, years_of_interest =
  ↪ c("2001", "2007", "2008", "2017", "2018", "2020"), flag = TRUE, ...) {
  if (is.null(tracker)) return(list(joined_df = safe_df(df_names), unmatched_df =
  ↪ data.frame(), flagged_df = data.frame()))
  tracker <- safe_df(tracker)
  if (!nrow(tracker)) return(list(joined_df = data.frame(), unmatched_df =
  ↪ safe_df(df_names), flagged_df = data.frame()))
  if (!".tracker_row" %in% names(tracker)) tracker$.tracker_row <- seq_len(nrow(tracker))
  xs <- if (inherits(df_names, "data.frame")) list(source = df_names) else
  ↪ as_input_list(df_names)

  joined_all <- list()
  unmatched_all <- list()
  flagged_all <- list()
  for (source_name in names(xs)) {
    source <- safe_df(xs[[source_name]])
    if (!nrow(source)) next
    source$.source_row <- seq_len(nrow(source))
    source$.source_name <- source_name
    source <- add_source_name_keys(source)
  }

```

```

if (!all(c(".source_state_key", ".source_district_key") %in% names(source))) {
  unmatched_all[[source_name]] <- source
  next
}

source$.matched <- FALSE
for (yr in years_of_interest) {
  sfx <- substr(as.character(yr), 3, 4)
  state_col <- paste0("state_", sfx)
  district_col <- paste0("district_", sfx)
  if (!all(c(state_col, district_col) %in% names(tracker))) next
  candidate <- tracker[c(".tracker_row", state_col, district_col)]
  candidate$.tracker_state_key <- canon(candidate[[state_col]])
  candidate$.tracker_district_key <- canon(candidate[[district_col]])
  exact <- merge(
    source[!source$.matched, , drop = FALSE],
    candidate,
    by.x = c(".source_state_key", ".source_district_key"),
    by.y = c(".tracker_state_key", ".tracker_district_key"),
    all.x = FALSE,
    all.y = FALSE
  )
  if (!nrow(exact)) next
  exact <- exact[!duplicated(exact$.source_row) & !duplicated(exact$.tracker_row), ,
  drop = FALSE]
  if (!nrow(exact)) next
  exact$match_year <- as.character(yr)
  exact$match_status <- "exact_name"
  joined_all[[paste(source_name, yr, sep = "_")] <- exact
  source$.matched[source$.source_row %in% exact$.source_row] <- TRUE
}

if (any(!source$.matched)) {
  fuzzy <- fuzzy_match_remaining_to_tracker(source[!source$.matched, , drop = FALSE],
  tracker, years_of_interest)
  if (nrow(fuzzy)) {
    joined_all[[paste0(source_name, "_fuzzy")]] <- fuzzy
    source$.matched[source$.source_row %in% fuzzy$.source_row] <- TRUE
    if (flag) flagged_all[[source_name]] <- fuzzy[fuzzy$possible_false_positive %in%
    TRUE, , drop = FALSE]
  }
}
unmatched_all[[source_name]] <- source[!source$.matched, , drop = FALSE]
}
list(
  joined_df = safe_bind_rows(joined_all),
  unmatched_df = safe_bind_rows(unmatched_all),
  flagged_df = safe_bind_rows(flagged_all)
)
}

add_source_name_keys <- function(source) {
  state <- first_col(source, c("state", "state_01", "state_07", "state_08", "state_17",
  "state_18", "state_20", "state_0708", "state_1718"))

```

```

    district <- first_col(source, c("district", "district_01", "district_07",
    ↪ "district_08", "district_17", "district_18", "district_20", "district_0708",
    ↪ "district_1718", "district_name"))
    if (!is.null(state)) source$.source_state_key <- canon(source[[state]])
    if (!is.null(district)) source$.source_district_key <- canon(source[[district]])
    source
  }

fuzzy_match_remaining_to_tracker <- function(source, tracker, years_of_interest) {
  out <- list()
  for (yr in years_of_interest) {
    sfx <- substr(as.character(yr), 3, 4)
    state_col <- paste0("state_", sfx)
    district_col <- paste0("district_", sfx)
    if (!all(c(state_col, district_col) %in% names(tracker))) next
    candidate <- tracker[c(".tracker_row", state_col, district_col)]
    candidate$.tracker_state_key <- canon(candidate[[state_col]])
    candidate$.tracker_district_key <- canon(candidate[[district_col]])
    for (i in seq_len(nrow(source))) {
      pool <- candidate[candidate$.tracker_state_key == source$.source_state_key[[i]], ,
    ↪ drop = FALSE]
      if (!nrow(pool)) next
      dist <- utils::adist(source$.source_district_key[[i]], pool$.tracker_district_key,
    ↪ ignore.case = TRUE)[1, ]
      best <- which.min(dist)
      if (!length(best) || !is.finite(dist[[best]]) || dist[[best]] > 2) next
      row <- cbind(source[i, , drop = FALSE], pool[best, , drop = FALSE])
      row$match_year <- as.character(yr)
      row$match_status <- "fuzzy_name"
      row$dist <- dist[[best]]
      row$possible_false_positive <- dist[[best]] > 0
      out[[length(out) + 1L]] <- row
    }
  }
  x <- safe_bind_rows(out)
  if (nrow(x)) x <- x[!duplicated(x$.source_row) & !duplicated(x$.tracker_row), , drop =
    ↪ FALSE]
  x
}

#' join year to tracker
#'
join_year_to_tracker <- function(source, tracker, years_of_interest, ...) {
  merge_dfs_into_tracker(source, tracker = tracker, years_of_interest =
    ↪ years_of_interest, ...)
}

#' join all district sources
#'
join_all_district_sources <- function(df_names, tracker, years_of_interest, ...) {
  merge_dfs_into_tracker(df_names, tracker = tracker, years_of_interest =
    ↪ years_of_interest, ...)
}

#' flag many to many matches

```

```

#'
flag_many_to_many_matches <- function(join_map) {
  join_map$many_to_many <- duplicated(join_map$.source_row) |
  ↪ duplicated(join_map$.tracker_row)
  join_map
}

#' score candidate matches
#'
score_candidate_matches <- function(candidates) {
  if (!"dist" %in% names(candidates)) candidates$dist <- 0
  candidates$score <- -num(candidates$dist)
  candidates
}

#' fuzzy join districts
#'
fuzzy_join_districts <- function(district_tracker, district_keys_2001,
  ↪ district_keys_2007, district_keys_2017, district_keys_2020, cfg) {
  tracker <- safe_df(district_tracker)
  if (nrow(tracker) && all(c("state_01", "district_01", "state_07", "district_07",
  ↪ "state_17", "district_17", "state_20", "district_20") %in% names(tracker))) {
    tracker$.tracker_row <- seq_len(nrow(tracker))
    tracker$source <- "legacy_tracker"
    tracker$match_status <- "legacy_tracker_row"
    tracker$possible_false_positive <- FALSE
    tracker$many_to_many <- FALSE
    attr(tracker, "unmatched_rows") <- tracker[0, , drop = FALSE]
    attr(tracker, "possible_false_positives") <- tracker[0, , drop = FALSE]
    attr(tracker, "many_to_many_cases") <- tracker[0, , drop = FALSE]
    return(tracker)
  }

  out <- safe_bind_rows(Map(function(keys, source) {
    if (!nrow(keys)) return(data.frame())
    keys$source <- source
    keys$match_status <- "source_key_unmatched"
    keys$possible_false_positive <- FALSE
    keys$many_to_many <- FALSE
    keys
  }, list(district_keys_2001, district_keys_2007, district_keys_2017,
  ↪ district_keys_2020), c("2001", "2007", "2017", "2020")))
  attr(out, "unmatched_rows") <- out
  attr(out, "possible_false_positives") <- out[out$possible_false_positive, , drop =
  ↪ FALSE]
  attr(out, "many_to_many_cases") <- out[out$many_to_many, , drop = FALSE]
  out
}

```

Contiguity-based spatial weights

Source: R/diagnostics/diagnose_spatial_weights.R

```

#' build spatial weights
#'

```

```

#' @return A spatial weights object, or an explicit inactive status when geometry is
↪ unavailable.
build_spatial_weights <- function(district_panel, cfg) {
  if (!inherits(district_panel, "sf")) {
    return(list(status = "out_of_active_pipeline", reason = "Requires sf geometry."))
  }
  geom <- sf::st_geometry(district_panel)
  coverage <- mean(!sf::st_is_empty(geom))
  if (!is.finite(coverage) || coverage < 0.75) {
    return(list(
      status = "out_of_active_pipeline",
      reason = paste0(
        "Requires validated non-empty geometry for at least 75% of district-panel rows;
↪ current coverage is ",
        round(100 * coverage, 1), "%."))
    ))
  }
  spatial_warnings <- character()
  capture_expected_spatial_warning <- function(expr) {
    withCallingHandlers(
      expr,
      warning = function(w) {
        msg <- conditionMessage(w)
        expected <- grepl("no neighbours|sub-graphs", msg, ignore.case = TRUE)
        if (expected) {
          spatial_warnings <-<- unique(c(spatial_warnings, msg))
          invokeRestart("muffleWarning")
        }
      }
    )
  }
  nb <- capture_expected_spatial_warning(
    spdep::poly2nb(district_panel[!sf::st_is_empty(geom), , drop = FALSE], queen = FALSE)
  )
  out <- capture_expected_spatial_warning(
    spdep::nb2listw(nb, style = "W", zero.policy = TRUE)
  )
  attr(out, "spatial_warnings") <- spatial_warnings
  out
}

#' diagnose spatial weights
#'
diagnose_spatial_weights <- function(district_panel, spatial_weights, cfg) {
  data.frame(
    diagnostic = "spatial_weights",
    status = spatial_weights$status %||% "constructed",
    warnings = paste(attr(spatial_weights, "spatial_warnings") %||% character(), collapse
↪ = "; "),
    stringsAsFactors = FALSE
  )
}

#' compare rook queen contiguity

```

```

#'
compare_rook_queen_contiguity <- function(district_panel) {
  tibble::tibble()
}

#' summarize islands
#'
summarize_islands <- function(spatial_weights) {
  tibble::tibble()
}

#' summarize neighbor counts
#'
summarize_neighbor_counts <- function(spatial_weights) {
  tibble::tibble()
}

#' save spatial weight diagnostics
#'
save_spatial_weight_diagnostics <- function(diagnostics) {
  diagnostics
}

```

Reusable IV specification and estimation

Source: R/iv/build_iv_formulas.R

```

#' make iv formula
#'
make_iv_formula <- function(dep, endog, instruments, controls = NULL, fixed_effects =
  ↪ NULL) {
  stats::as.formula(paste(
    dep,
    "~",
    paste(c(endog, controls, fixed_effects), collapse = " + "),
    "|",
    paste(c(instruments, controls, fixed_effects), collapse = " + ")
  ))
}

legacy_iv_controls <- function() {
  c(
    "consumption_0708", "gini_cons_0708",
    "pct_urban", "avg_hh_size", "dependency_ratio",
    "pct_fem_head",
    "pct_hindu", "pct_muslim",
    "pct_st", "pct_sc", "pct_obc",
    "pct_small_land", "pct_medium_land", "pct_large_land",
    "pct_head_lit_to_primary", "pct_head_secondary_plus"
  )
}

#' build iv formulas
#'
build_iv_formulas <- function(cfg) {
  controls <- legacy_iv_controls()
}

```

```

list(
  consumption = make_iv_formula(
    "consumption_pct_change",
    "EMIE",
    "wavg_ling_degrees",
    controls
  ),
  gini = make_iv_formula(
    "gini_change",
    "EMIE",
    "wavg_ling_degrees",
    controls
  )
)
}

#' build baseline 2sls formula
#'
build_baseline_2sls_formula <- function(...) make_iv_formula(...)

#' build fd 2sls formula
#'
build_fd_2sls_formula <- function(...) {
  list(status = "out_of_active_pipeline", reason = "First-difference 2SLS formula is
    ↪ deferred until the FD redesign is specified.")
}

#' build state fe 2sls formula
#'
build_state_fe_2sls_formula <- function(...) {
  list(status = "out_of_active_pipeline", reason = "State fixed-effect 2SLS formula is
    ↪ deferred until the state-FE redesign is specified.")
}

```

Multicollinearity and spatial autocorrelation diagnostics

Source: R/diagnostics/diagnose_multicollinearity.R

```

#' diagnose multicollinearity
#'
#' @return A data frame with design-matrix condition diagnostics.
diagnose_multicollinearity <- function(district_panel, iv_models, cfg) {
  model <- if (is.list(iv_models) && length(iv_models)) iv_models[[1]] else iv_models
  out <- tryCatch({
    X <- stats::model.matrix(model)
    data.frame(
      test = "kappa",
      n = nrow(X),
      rank = qr(X)$rank,
      columns = ncol(X),
      kappa = kappa(X, exact = TRUE),
      status = "estimated",
      reason = NA_character_,
      stringsAsFactors = FALSE
    )
  }, error = function(e) {

```

```

    data.frame(
      test = "kappa",
      n = NA_integer_,
      rank = NA_integer_,
      columns = NA_integer_,
      kappa = NA_real_,
      status = "out_of_active_pipeline",
      reason = conditionMessage(e),
      stringsAsFactors = FALSE
    )
  })
  out
}

#' compute design matrix rank
#'
compute_design_matrix_rank <- function(formula, data) {
  X <- stats::model.matrix(formula, data = data); tibble::tibble(rank = qr(X)$rank, ncol
  ↪   = ncol(X))
}

#' compute kappa
#'
compute_kappa <- function(formula, data) {
  X <- stats::model.matrix(formula, data = data); kappa(X, exact = TRUE)
}

#' compute vif if applicable
#'
compute_vif_if_applicable <- function(model) {
  if (requireNamespace("car", quietly = TRUE)) car::vif(model) else NULL
}

#' save multicollinearity diagnostics
#'
save_multicollinearity_diagnostics <- function(diagnostics) {
  diagnostics
}

```